


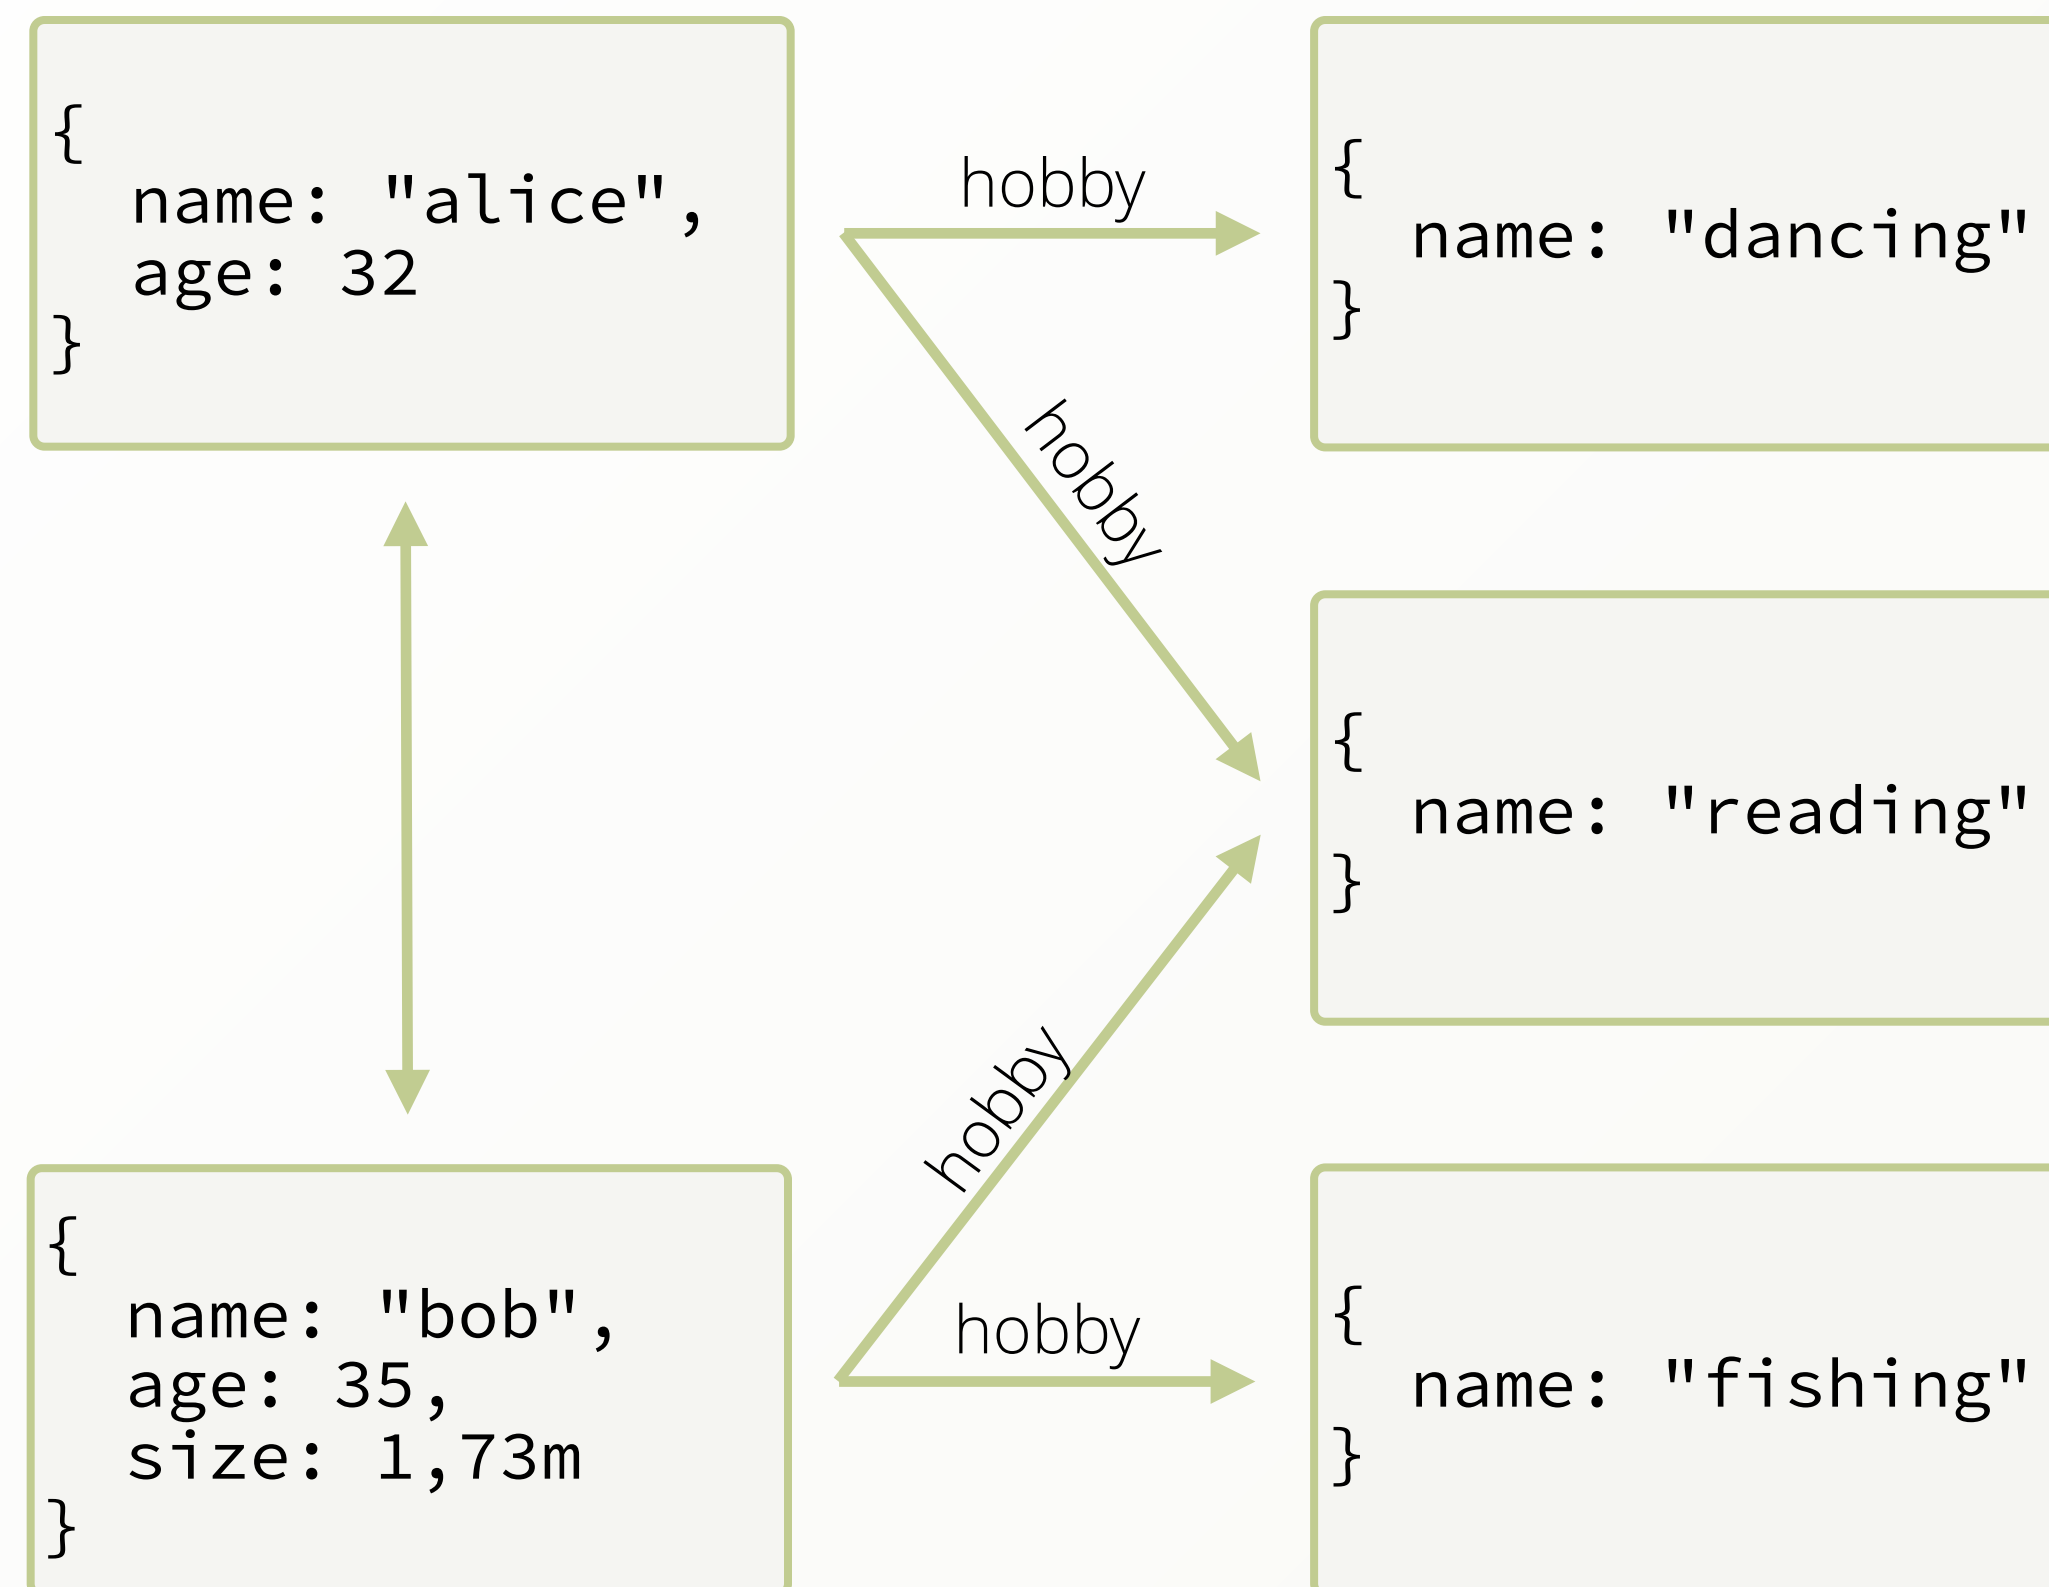


# *Handling Billions Of Edges in a Graph Database*

Michael Hackstein  
@mchacki

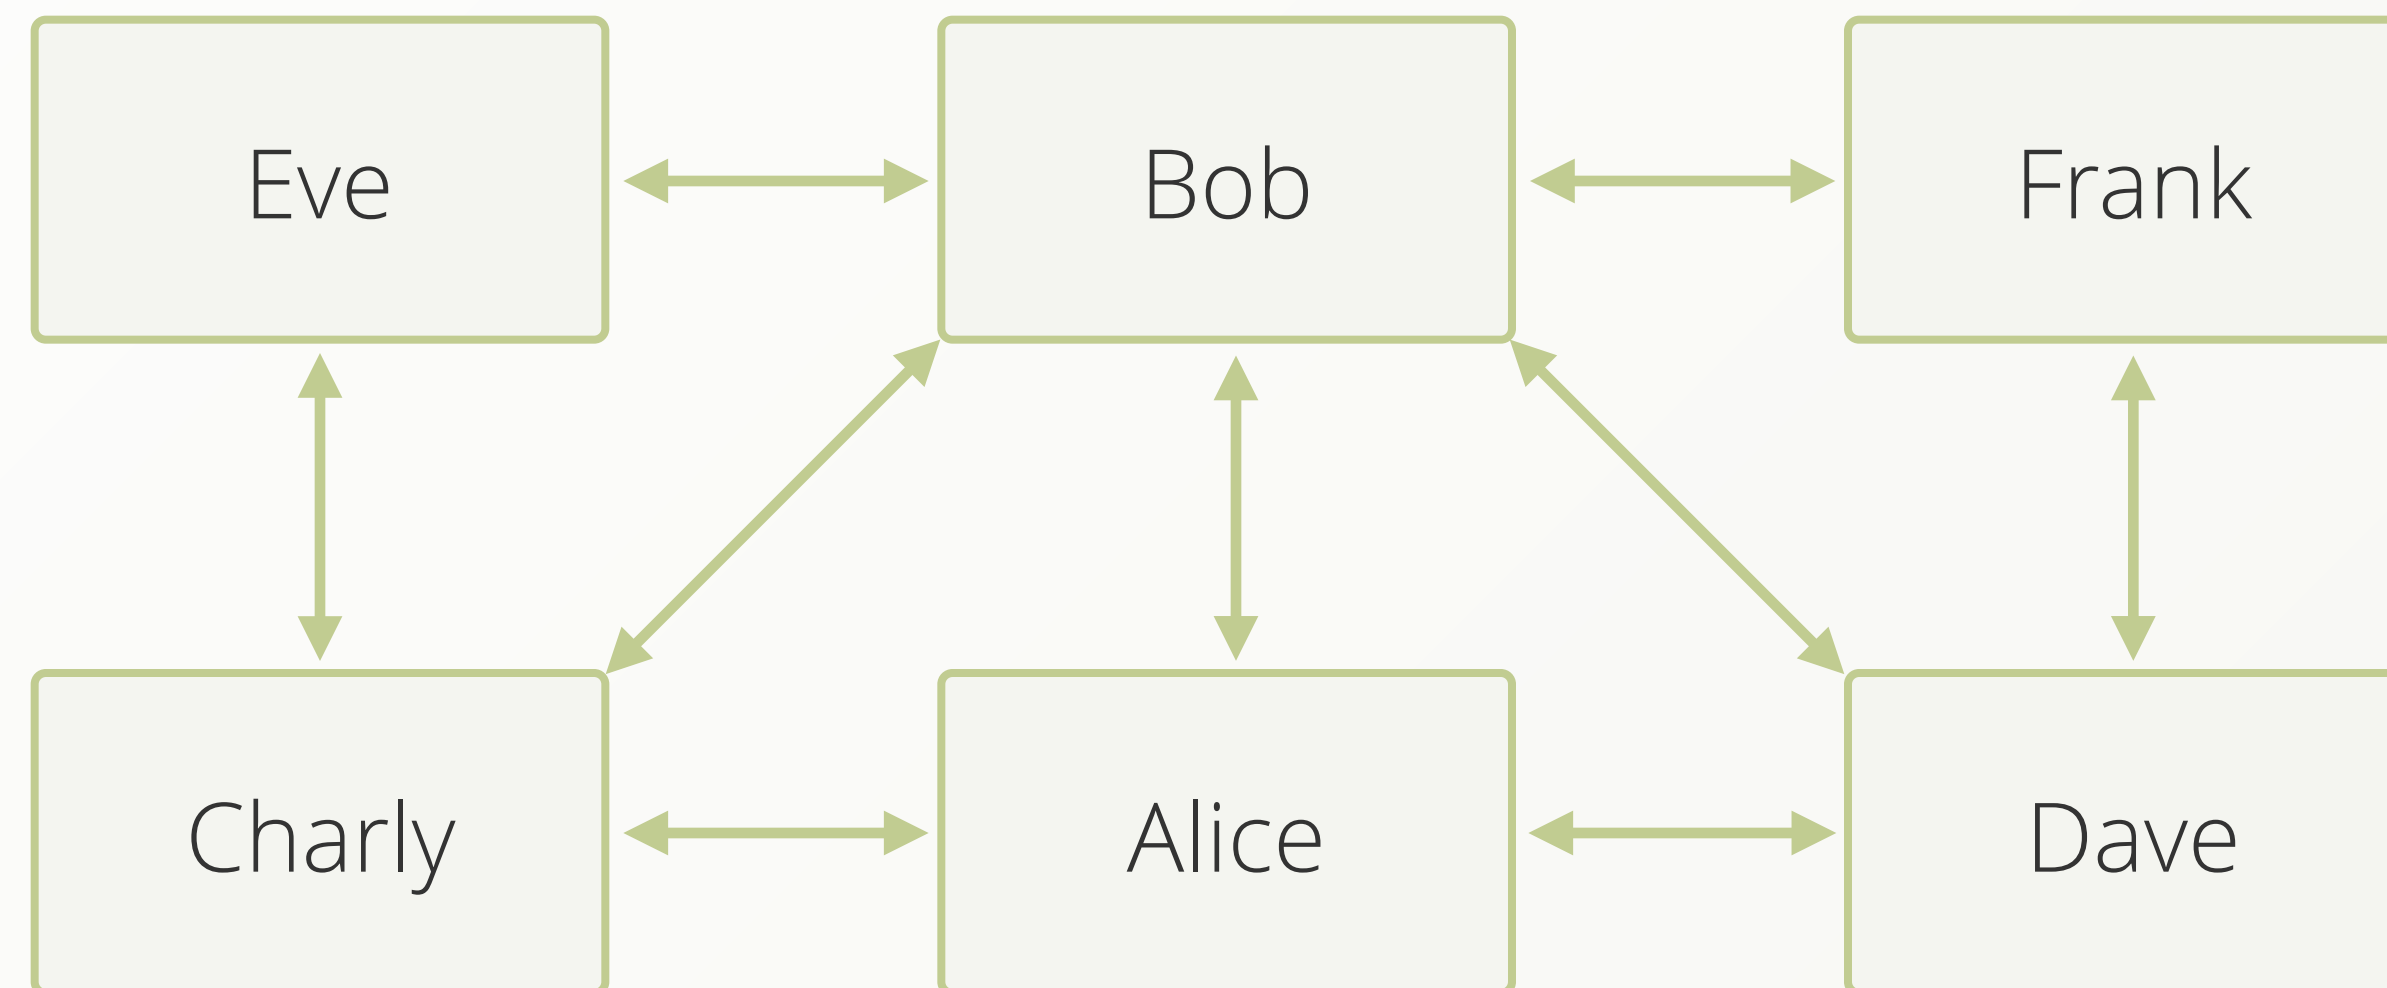


- ▶ Schema-free Objects (Vertices)
- ▶ Relations between them (Edges)
- ▶ Edges have a direction

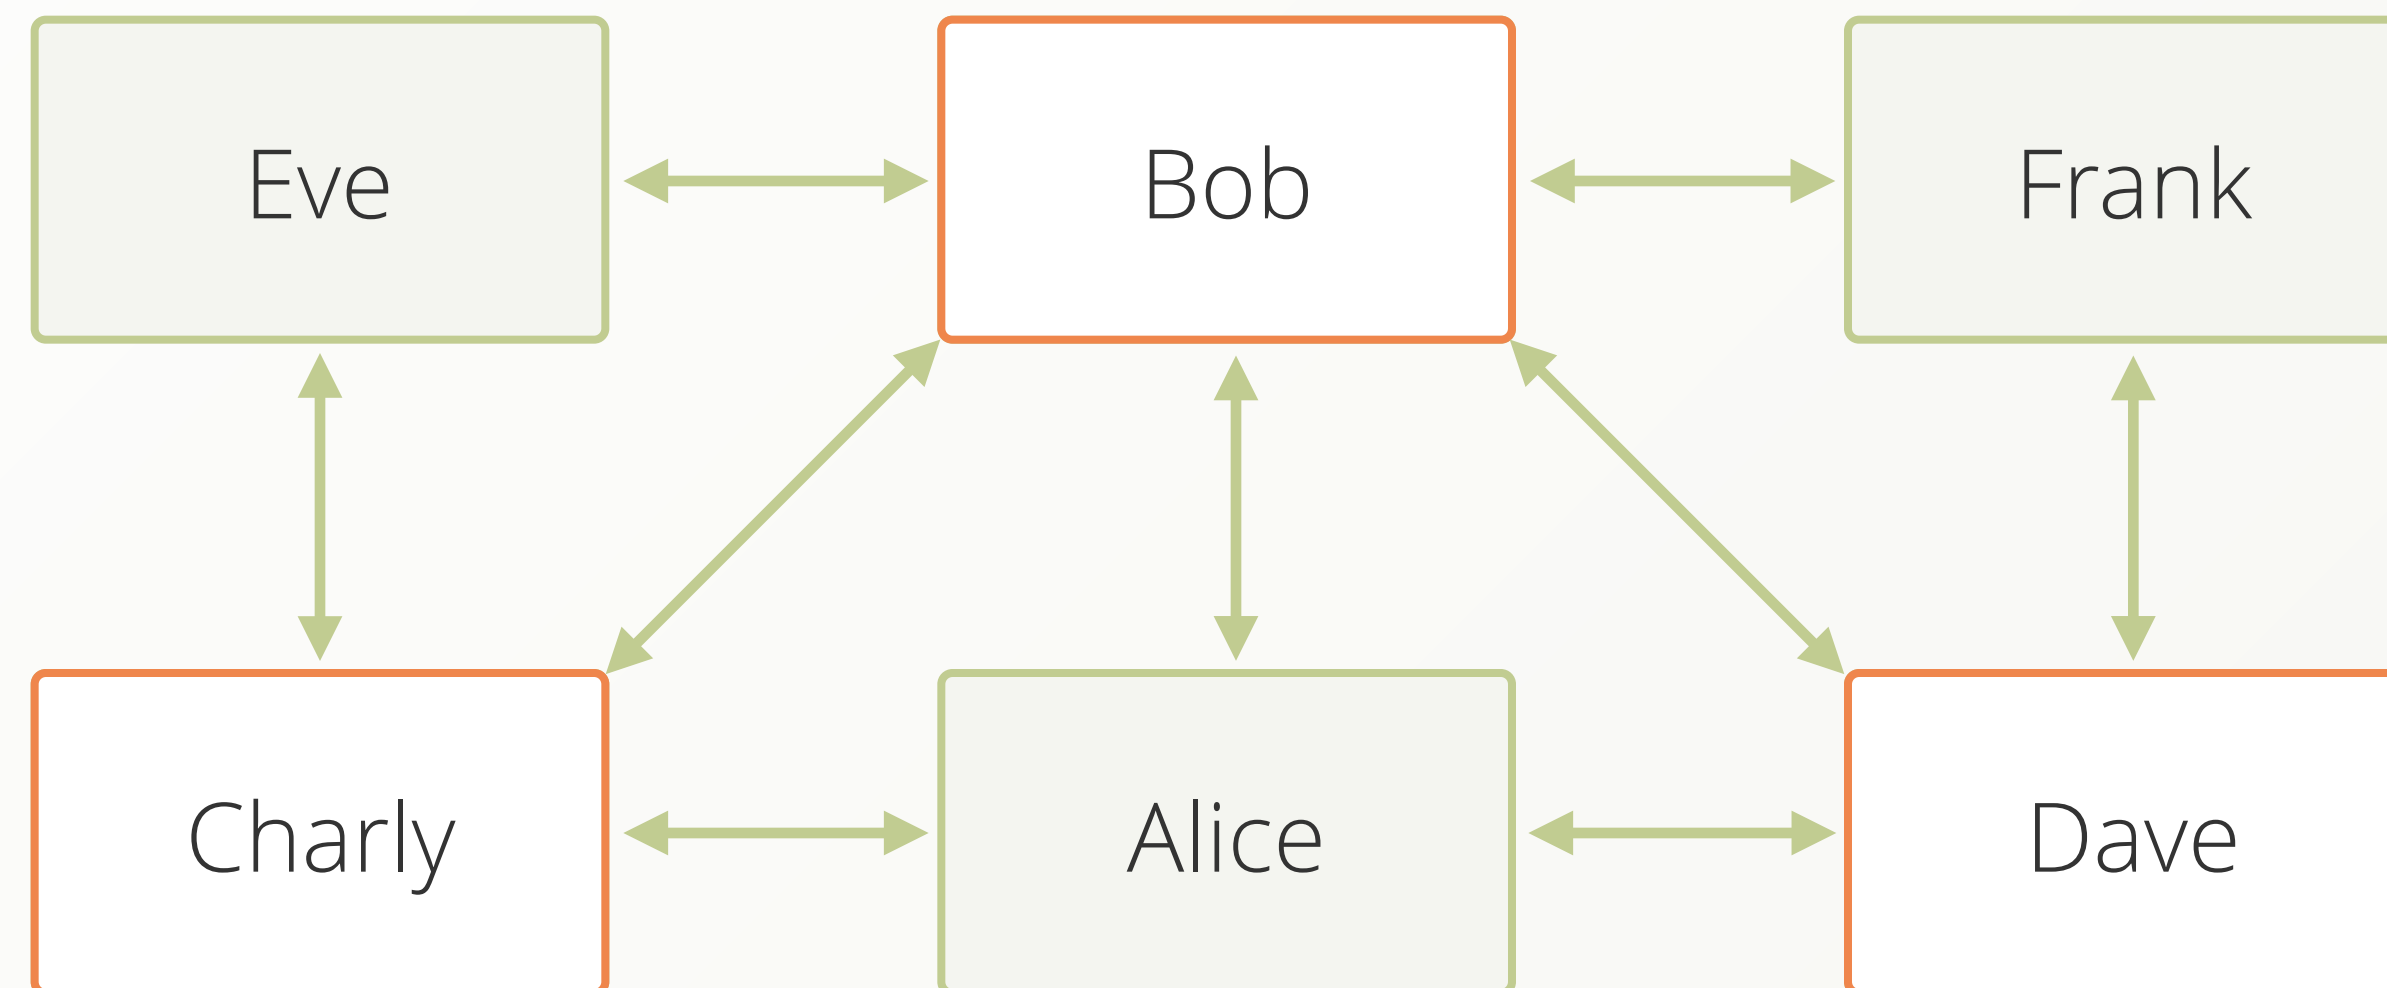


- ▶ Schema-free Objects (Vertices)
- ▶ Relations between them (Edges)
- ▶ Edges have a direction
- ▶ Edges can be queried in both directions
- ▶ Easily query a range of edges (2 to 5)
- ▶ Undefined number of edges (1 to \*)
- ▶ Shortest Path between two vertices


- ▶ Give me all friends of Alice



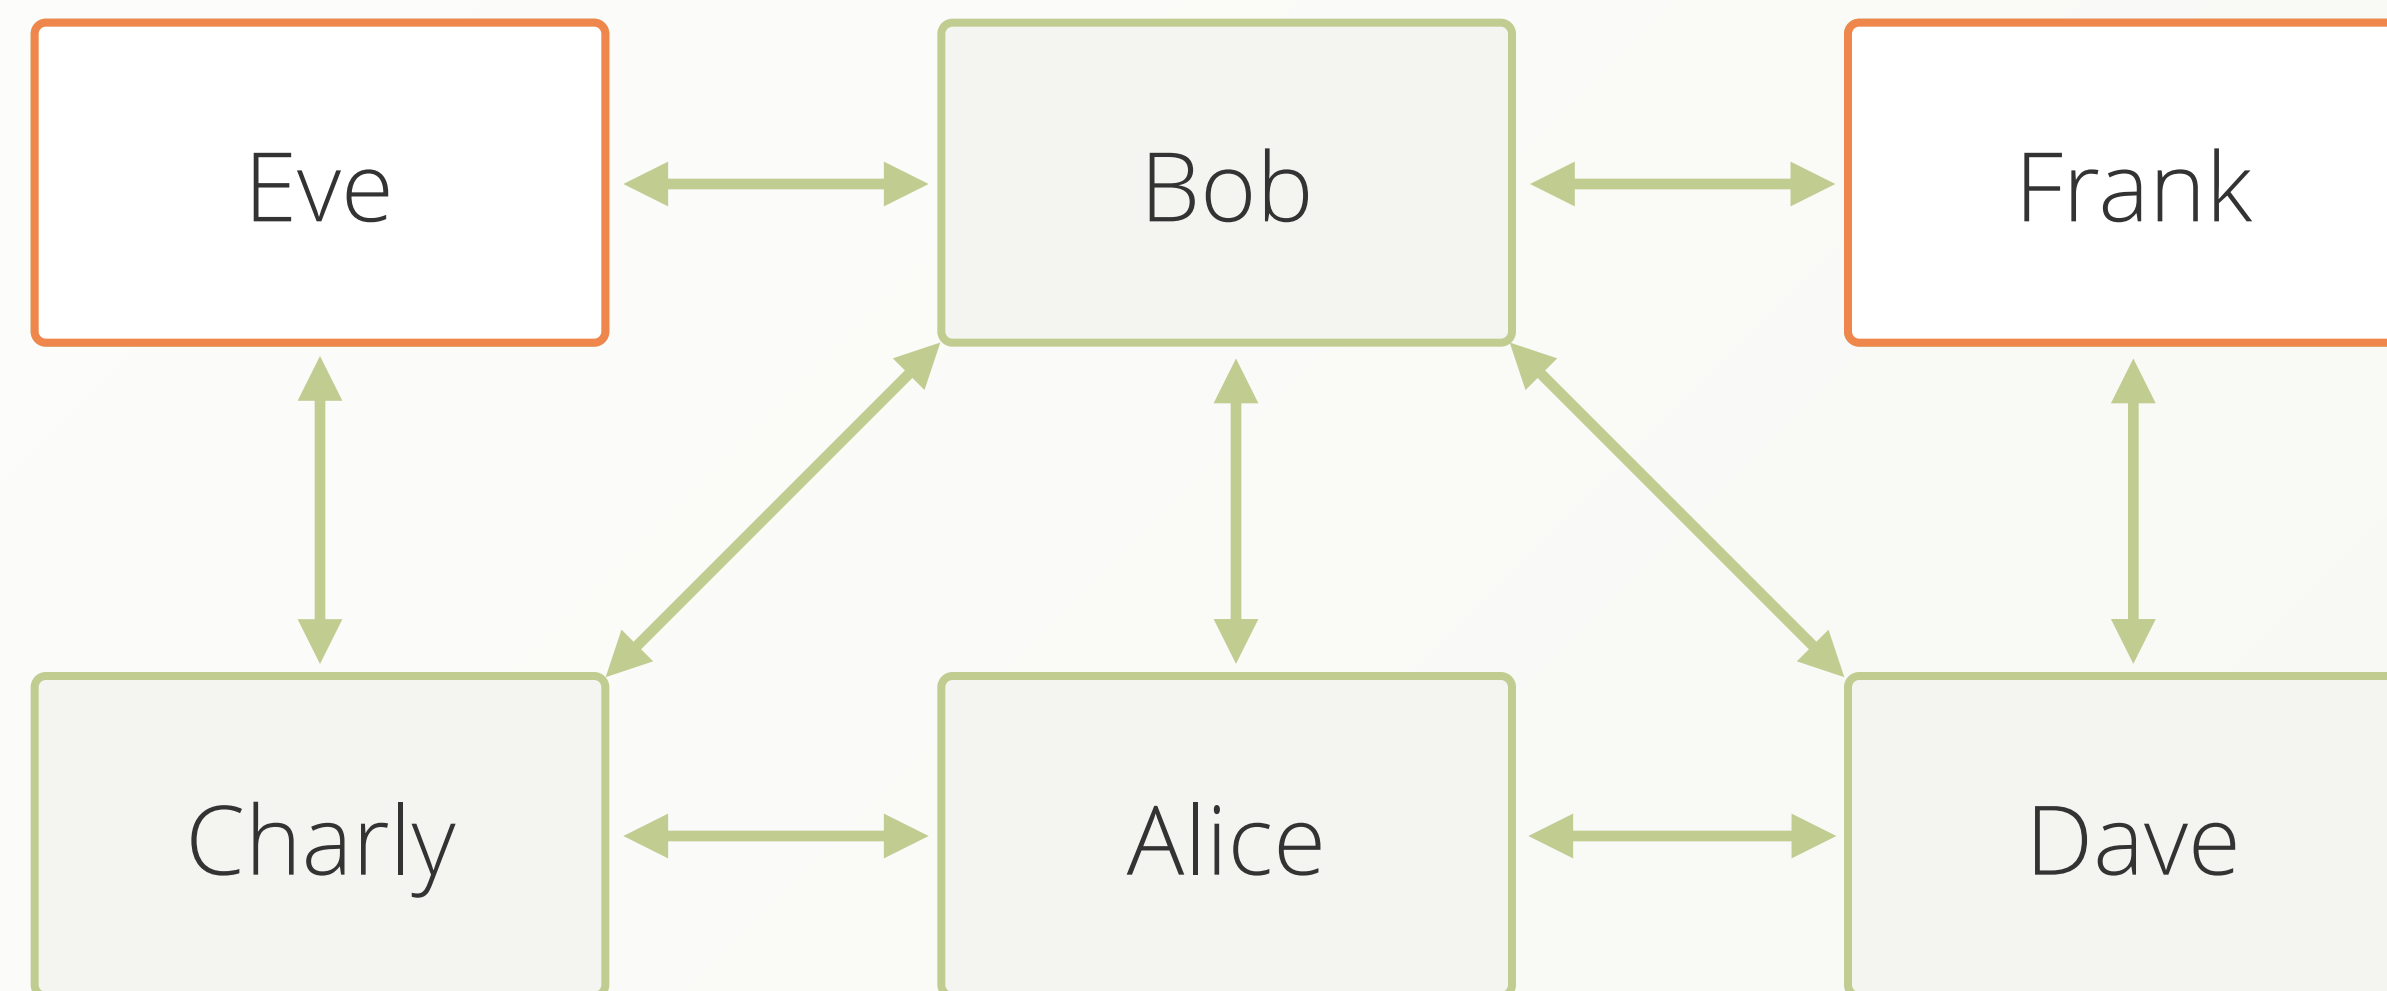
- ▶ Give me all friends of Alice




- ▶ Give me all friends-of-friends of Alice



- ▶ Give me all friends-of-friends of Alice

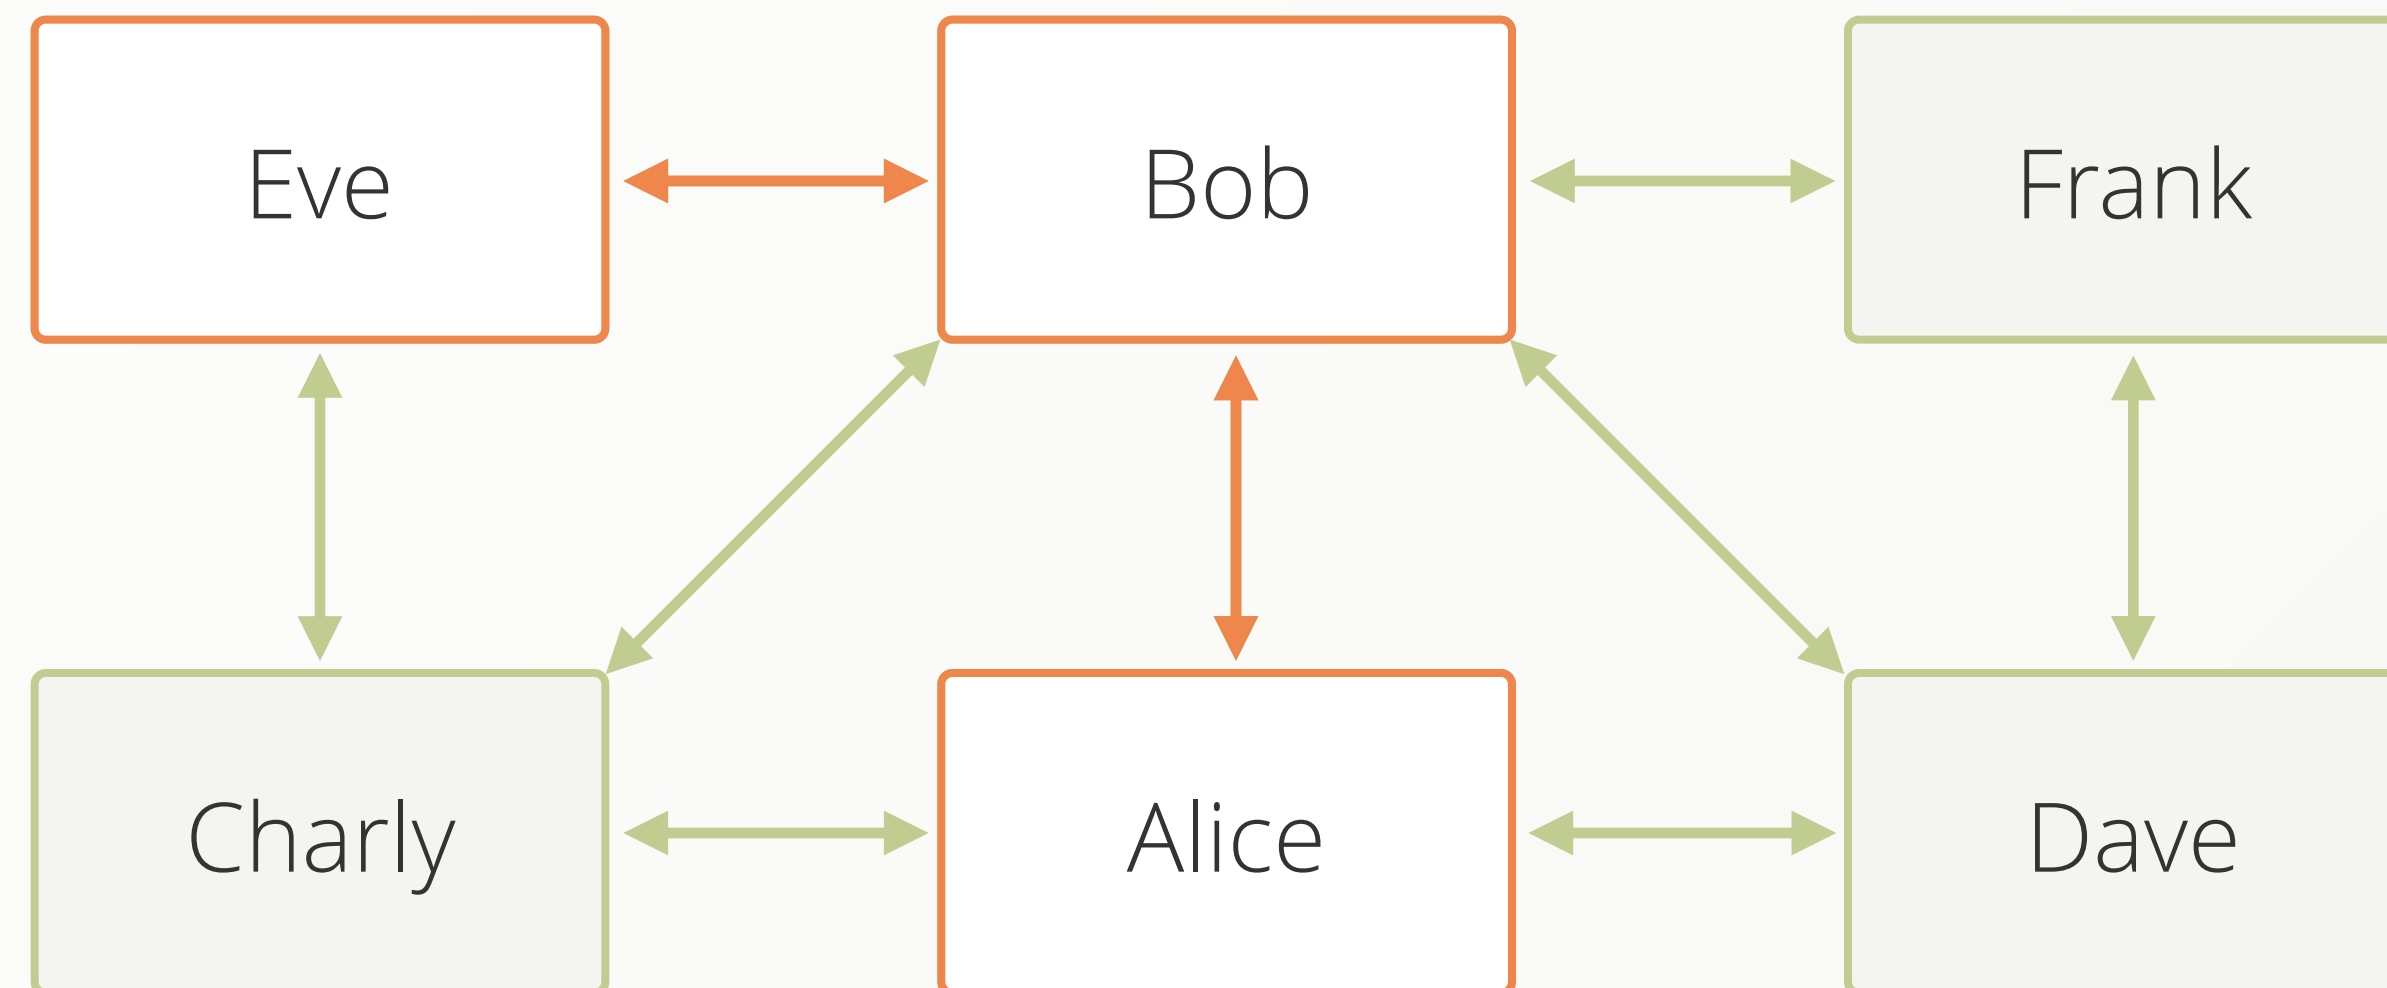


- ▶ What is the linking path between Alice and Eve

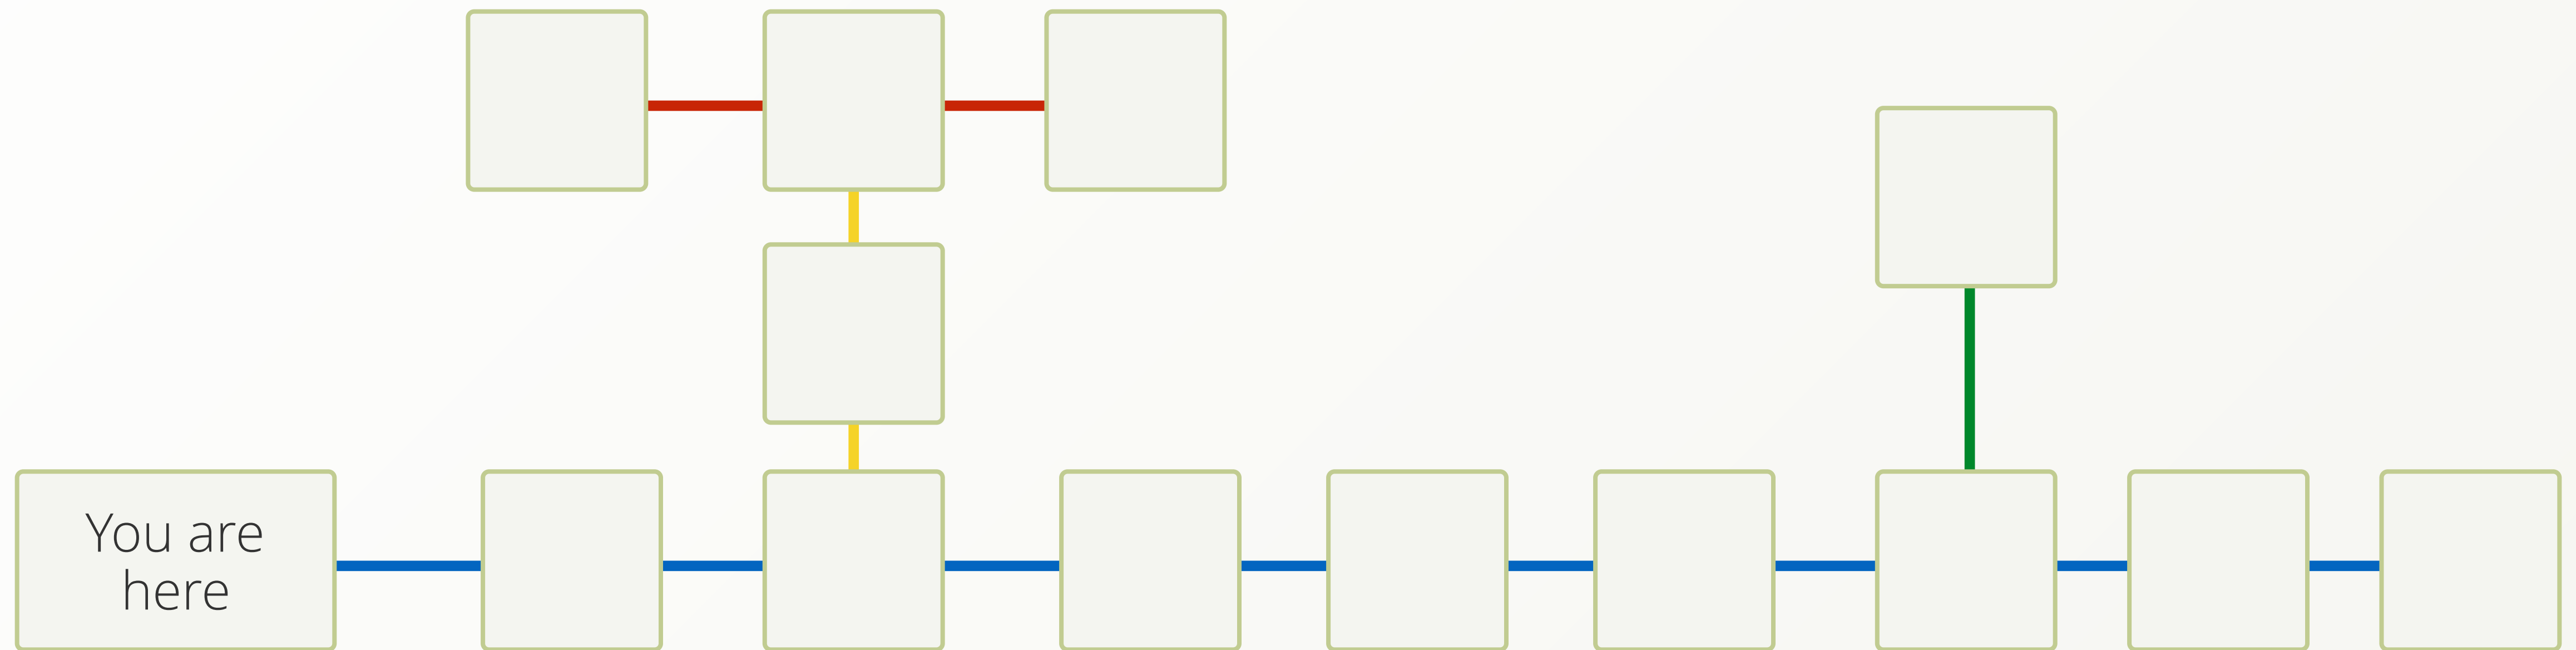




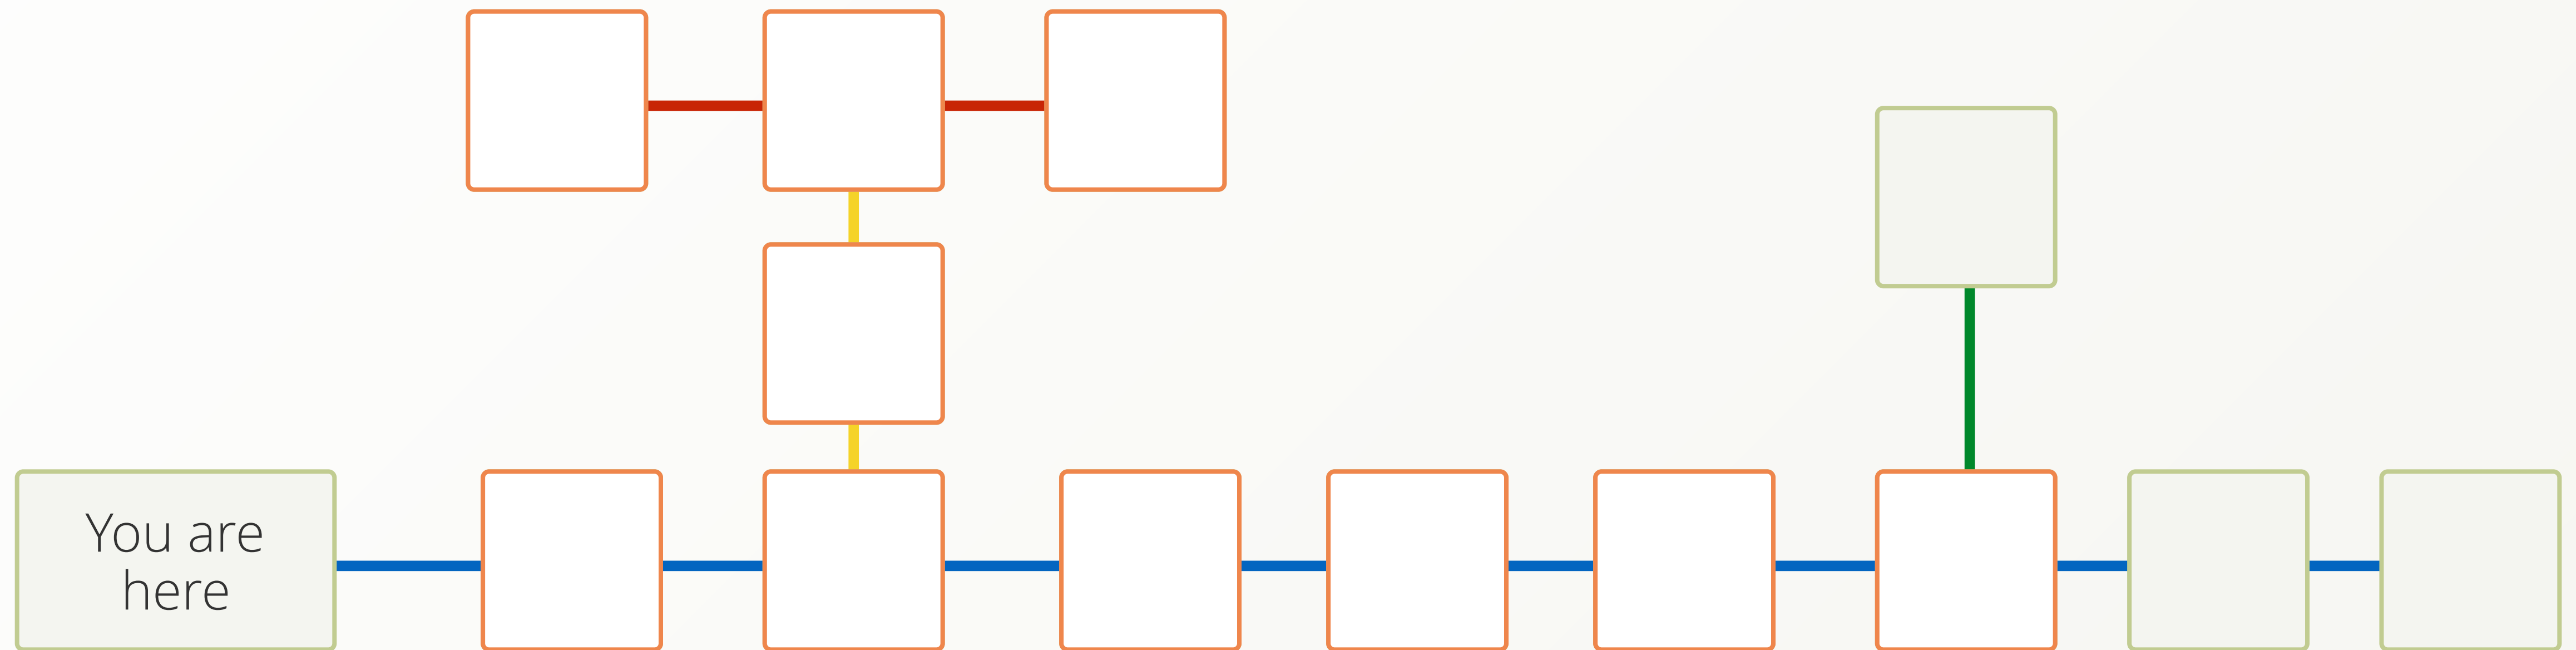
- ▶ What is the linking path between **Alice** and **Eve**




- ▶ Which Trainstations can I reach if I am allowed to drive a distance of at most 6 stations on my ticket




- ▶ Which **Trainstations** can I reach if I am allowed to drive **a distance of at most 6 stations** on my ticket



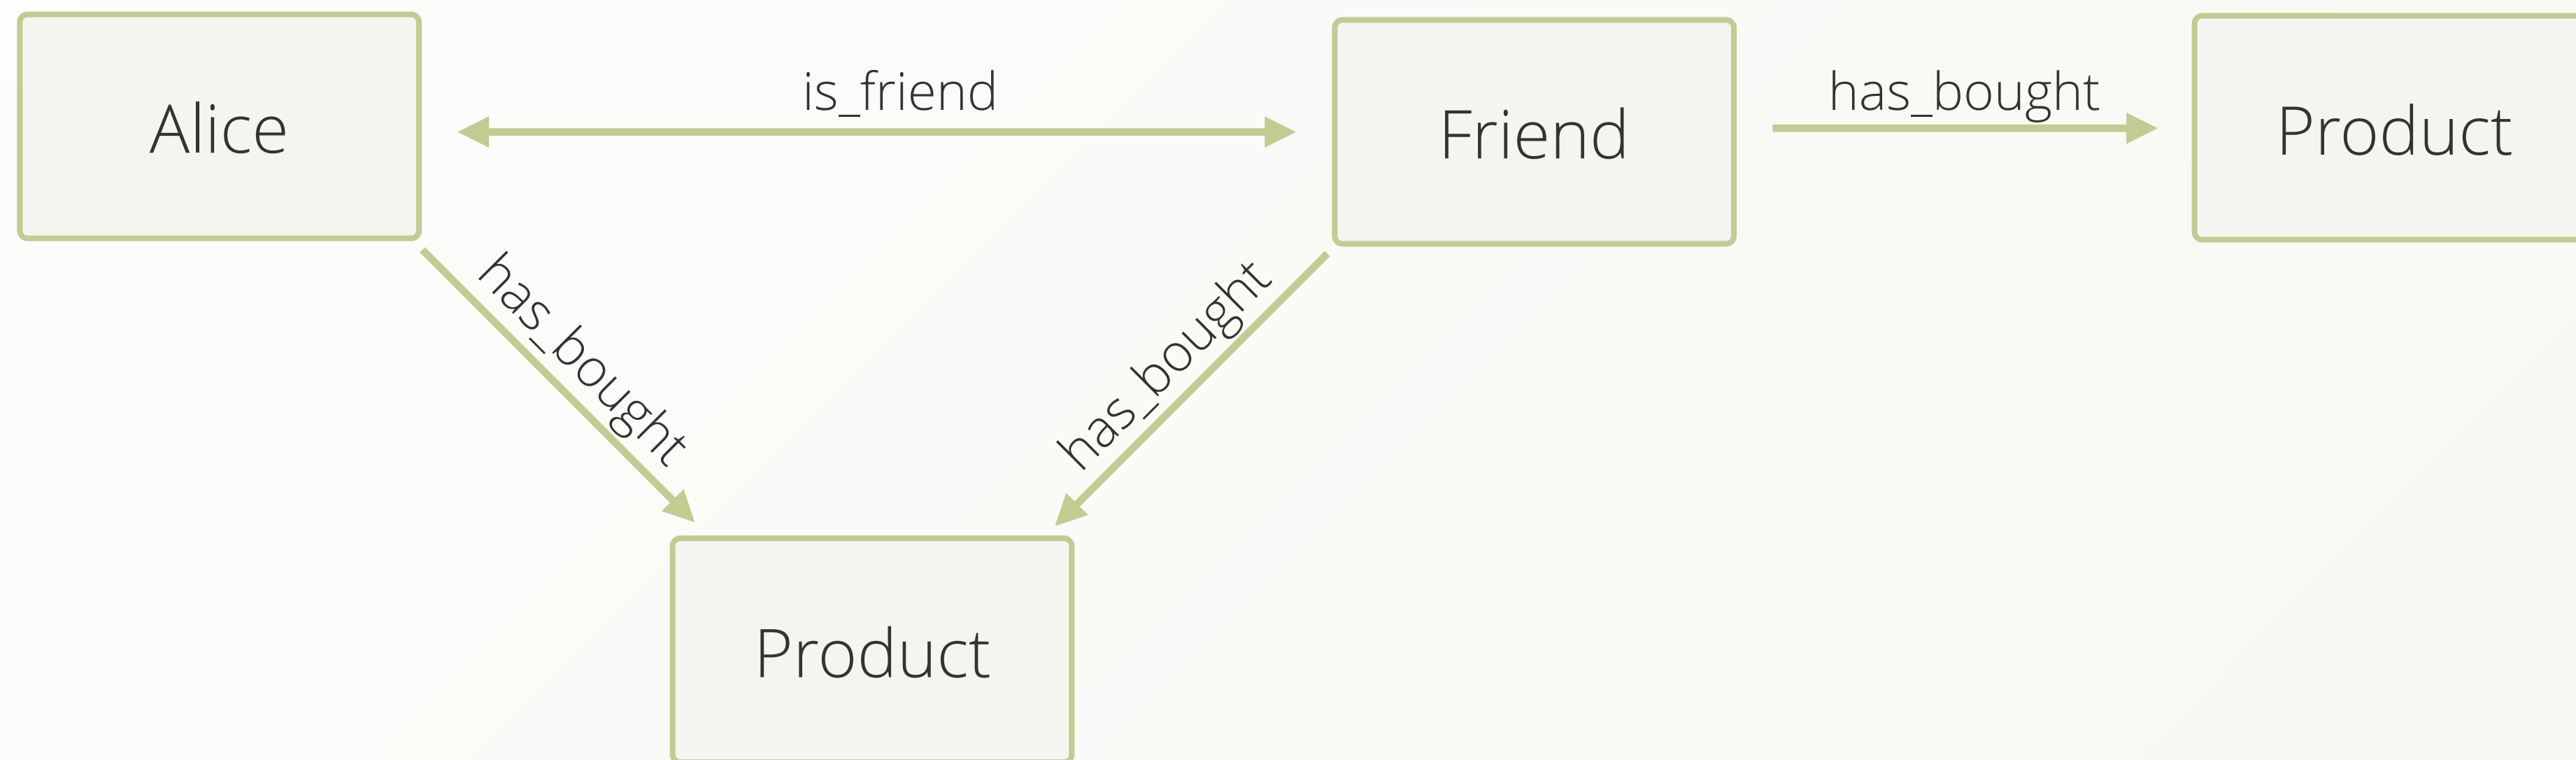
- ▶ Give me all **users** that share two **hobbies** with **Alice**



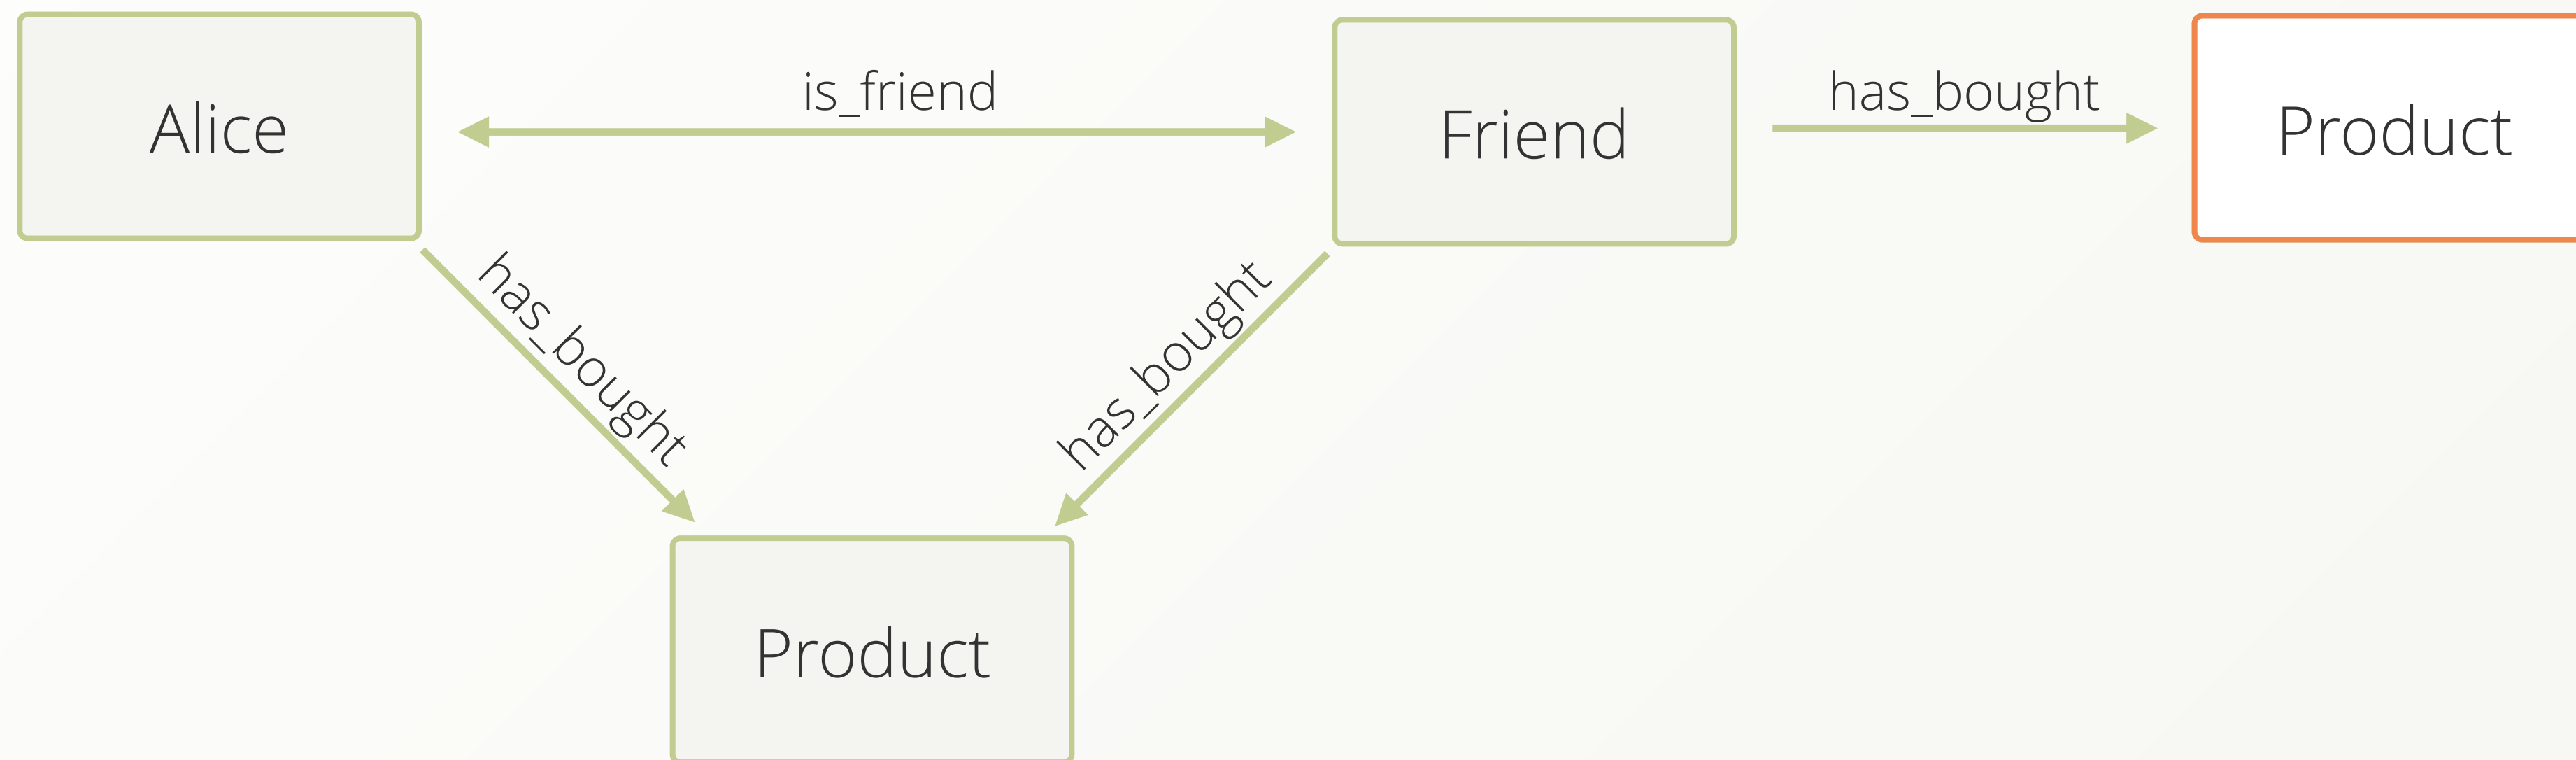
- ▶ Give me all **users** that share two **hobbies** with **Alice**



- ▶ Give me all products that at least one of my friends has bought together with the products I already own, ordered by how many friends have bought it and the products rating, but only 20 of them.



- ▶ Give me all products that at least one of my friends has bought together with the products I already own, ordered by how many friends have bought it and the products rating, but only 20 of them.







- ▶ Give me all users which have an `age` attribute between `21` and `35`.

- ▶ Give me all users which have an **age** attribute between **21** and **35**.
- ▶ Give me the **age** distribution of all **users**


- ▶ Give me all users which have an **age** attribute between **21** and **35**.
- ▶ Give me the **age** distribution of all **users**
- ▶ Group all **users** by their **name**




- ▶ We first pick a start vertex (S)




- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S






- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges




- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)







- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)
- ▶ We apply filters on edges




- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)
- ▶ We apply filters on edges
- ▶ The next vertex (E) is in desired depth. Return the path  $S \rightarrow A \rightarrow E$




- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)
- ▶ We apply filters on edges
- ▶ The next vertex (E) is in desired depth. Return the path  $S \rightarrow A \rightarrow E$
- ▶ Go back to the next unfinished vertex (B)



- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)
- ▶ We apply filters on edges
- ▶ The next vertex (E) is in desired depth. Return the path  $S \rightarrow A \rightarrow E$
- ▶ Go back to the next unfinished vertex (B)
- ▶ We iterate down on (B)



- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)
- ▶ We apply filters on edges
- ▶ The next vertex (E) is in desired depth. Return the path S -> A -> E
- ▶ Go back to the next unfinished vertex (B)
- ▶ We iterate down on (B)
- ▶ We apply filters on edges



- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)
- ▶ We apply filters on edges
- ▶ The next vertex (E) is in desired depth. Return the path S -> A -> E
- ▶ Go back to the next unfinished vertex (B)
- ▶ We iterate down on (B)
- ▶ We apply filters on edges
- ▶ The next vertex (F) is in desired depth. Return the path S -> B -> F


▶ Once:		O
▶ Find the start vertex	Depends on indexes: Hash:	1
▶ For every depth:		
▶ Find all connected edges	Edge-Index or Index-Free:	1
▶ Filter non-matching edges	Linear in edges:	n
▶ Find connected vertices	Depends on indexes: Hash:	n * 1
▶ Filter non-matching vertices	Linear in vertices:	n
	Only one pass:	3n



- ▶ Linear sounds evil?
  - ▶ NOT linear in All Edges  $O(E)$
  - ▶ Only Linear in relevant Edges  $n < E$
- ▶ Traversals solely scale with their result size
- ▶ They are not effected at all by total amount of data
- ▶ BUT: Every depth increases the exponent:  $O((3n)^d)$
- ▶ "7 degrees of separation":  $(3n)^6 < E < (3n)^7$



- ▶ MULTI-MODEL database
  - ▶ Stores Key Value, Documents, and Graphs
  - ▶ All in one core
- ▶ Query language AQL
  - ▶ Document Queries
  - ▶ Graph Queries
  - ▶ Joins
  - ▶ All can be combined in the same statement
- ▶ ACID support including Multi Collection Transactions



```
FOR user IN users  
  RETURN user
```

```
FOR user IN users  
  FILTER user.name == "alice"  
  RETURN user
```




Alice

```
FOR user IN users
  FILTER user.name == "alice"
  FOR product IN OUTBOUND user has_bought
  RETURN product
```




Alice

```
FOR user IN users
  FILTER user.name == "alice"
  FOR product IN OUTBOUND user has_bought
  RETURN product
```





```
FOR user IN users
  FILTER user.name == "alice"
  FOR recommendation, action, path IN 3 ANY user has_bought
    FILTER path.vertices[2].age <= user.age + 5
    AND path.vertices[2].age >= user.age - 5
    FILTER recommendation.price < 25
  LIMIT 10
  RETURN recommendation
```



```
FOR user IN users
  FILTER user.name == "alice"
  FOR recommendation, action, path IN 3 ANY user has_bought
    FILTER path.vertices[2].age <= user.age + 5
    AND path.vertices[2].age >= user.age - 5
    FILTER recommendation.price < 25
  LIMIT 10
  RETURN recommendation
```




- ▶ Many graphs have "Supernodes"
  - ▶ Vertices with many inbound and/or outbound edges
- ▶ Traversing over them is expensive
- ▶ Often you only need a subset of edges





- ▶ Remember Complexity?  $O((3n)^d)$
- ▶ Filtering of non-matching edges is linear for every depth
- ▶ Index all edges based on their vertices and arbitrary other attributes
  - ▶ Find initial set of edges in identical time
  - ▶ Less / No post-filtering required
  - ▶ This decreases the  $n$  significantly



- ▶ We have the rise of big data
  - ▶ Store everything you can
- ▶ Dataset easily grows beyond one machine
- ▶ This includes graph data!

- ▶ Distribute graph on several machines (sharding)
- ▶ How to query it now?
  - ▶ No global view of the graph possible any more
  - ▶ What about edges between servers?
- ▶ In a shared environment network most of the time is the bottleneck
  - ▶ Reduce network hops
- ▶ Vertex-Centric Indexes again help with super-nodes
  - ▶ But: Only on a local machine

First let's do  
the cluster thingy

- ▶ ArangoDB is the first fully certified database including the persistence primitives for DC/OS
- ▶ ArangoDB's cluster resource management is based on Apache Mesos
- ▶ Later this year we will also support Kubernetes



- ▶ ArangoDB can run clusters without it
    - ▶ Setup Requires manual effort (can be scripted)
  - ▶ This works:
    - ▶ Automatic Failover (Follower takes over if leader dies)
    - ▶ Rebalancing of shards
    - ▶ Everything inside of ArangoDB
  - ▶ This is based on Mesos:
    - ▶ Complete self healing
    - ▶ Automatic restart of ArangoDBs (on new machines)
- ➡ We suggest you have someone on call

Demo Time  
DC/OS

Now distribute  
the graph



- ▶ Only parts of the graph on every machine
  - ▶ Neighboring vertices may be on different machines
  - ▶ Even edges could be on other machines than their vertices
- 
- ▶ Queries need to be executed in a distributed way
  - ▶ Result needs to be merged locally

# Random Distribution


---

- ▶ Advantages:

- ▶ every server takes an equal portion of data
- ▶ easy to realize
- ▶ no knowledge about data required
- ▶ always works


- ▶ Disadvantages:

- ▶ Neighbors on different machines
- ▶ Probably edges on other machines than their vertices
- ▶ A lot of network overhead is required for querying



# Random Distribution


- ▶ Advantages:
  - ▶ every server takes an equal portion of data
  - ▶ easy to realize
  - ▶ no knowledge about data required
  - ▶ always works
- ▶ Disadvantages:
  - ▶ Neighbors on different machines
  - ▶ Probably edges on other machines than their vertices
  - ▶ A lot of network overhead is required for querying



# Index-Free Adjacency

---



- ▶ Used by most other graph databases
- ▶ Every vertex maintains two lists of its edges (IN and OUT)
  - ▶ Do not use an index to find edges
  - ▶ How to shard this?



# Index-Free Adjacency

---

- ▶ Used by most other graph databases
- ▶ Every vertex maintains two lists of its edges (IN and OUT)
  - ▶ Do not use an index to find edges
  - ▶ How to shard this?




# Index-Free Adjacency

---

- ▶ Used by most other graph databases
- ▶ Every vertex maintains two lists of its edges (IN and OUT)
  - ▶ Do not use an index to find edges
  - ▶ How to shard this?

—




# Index-Free Adjacency

---

- ▶ Used by most other graph databases
- ▶ Every vertex maintains two lists of its edges (IN and OUT)
  - ▶ Do not use an index to find edges
  - ▶ How to shard this?

????




# Index-Free Adjacency

---

- ▶ Used by most other graph databases
- ▶ Every vertex maintains two lists of its edges (IN and OUT)
  - ▶ Do not use an index to find edges
  - ▶ How to shard this?

????





- ▶ ArangoDB uses a hash-based EdgeIndex ( $O(1)$  - lookup)
  - ▶ The vertex is independent of its edges
  - ▶ It can be stored on a different machine



# Domain Based Distribution

---

- ▶ Many Graphs have a natural distribution
  - ▶ By country/region for People
  - ▶ By tags for Blogs
  - ▶ By category for Products
- ▶ Most edges in same group
- ▶ Rare edges between groups



# Domain Based Distribution

---

- ▶ Many Graphs have a natural distribution
  - ▶ By country/region for People
  - ▶ By tags for Blogs
  - ▶ By category for Products
- ▶ Most edges in same group
- ▶ Rare edges between groups



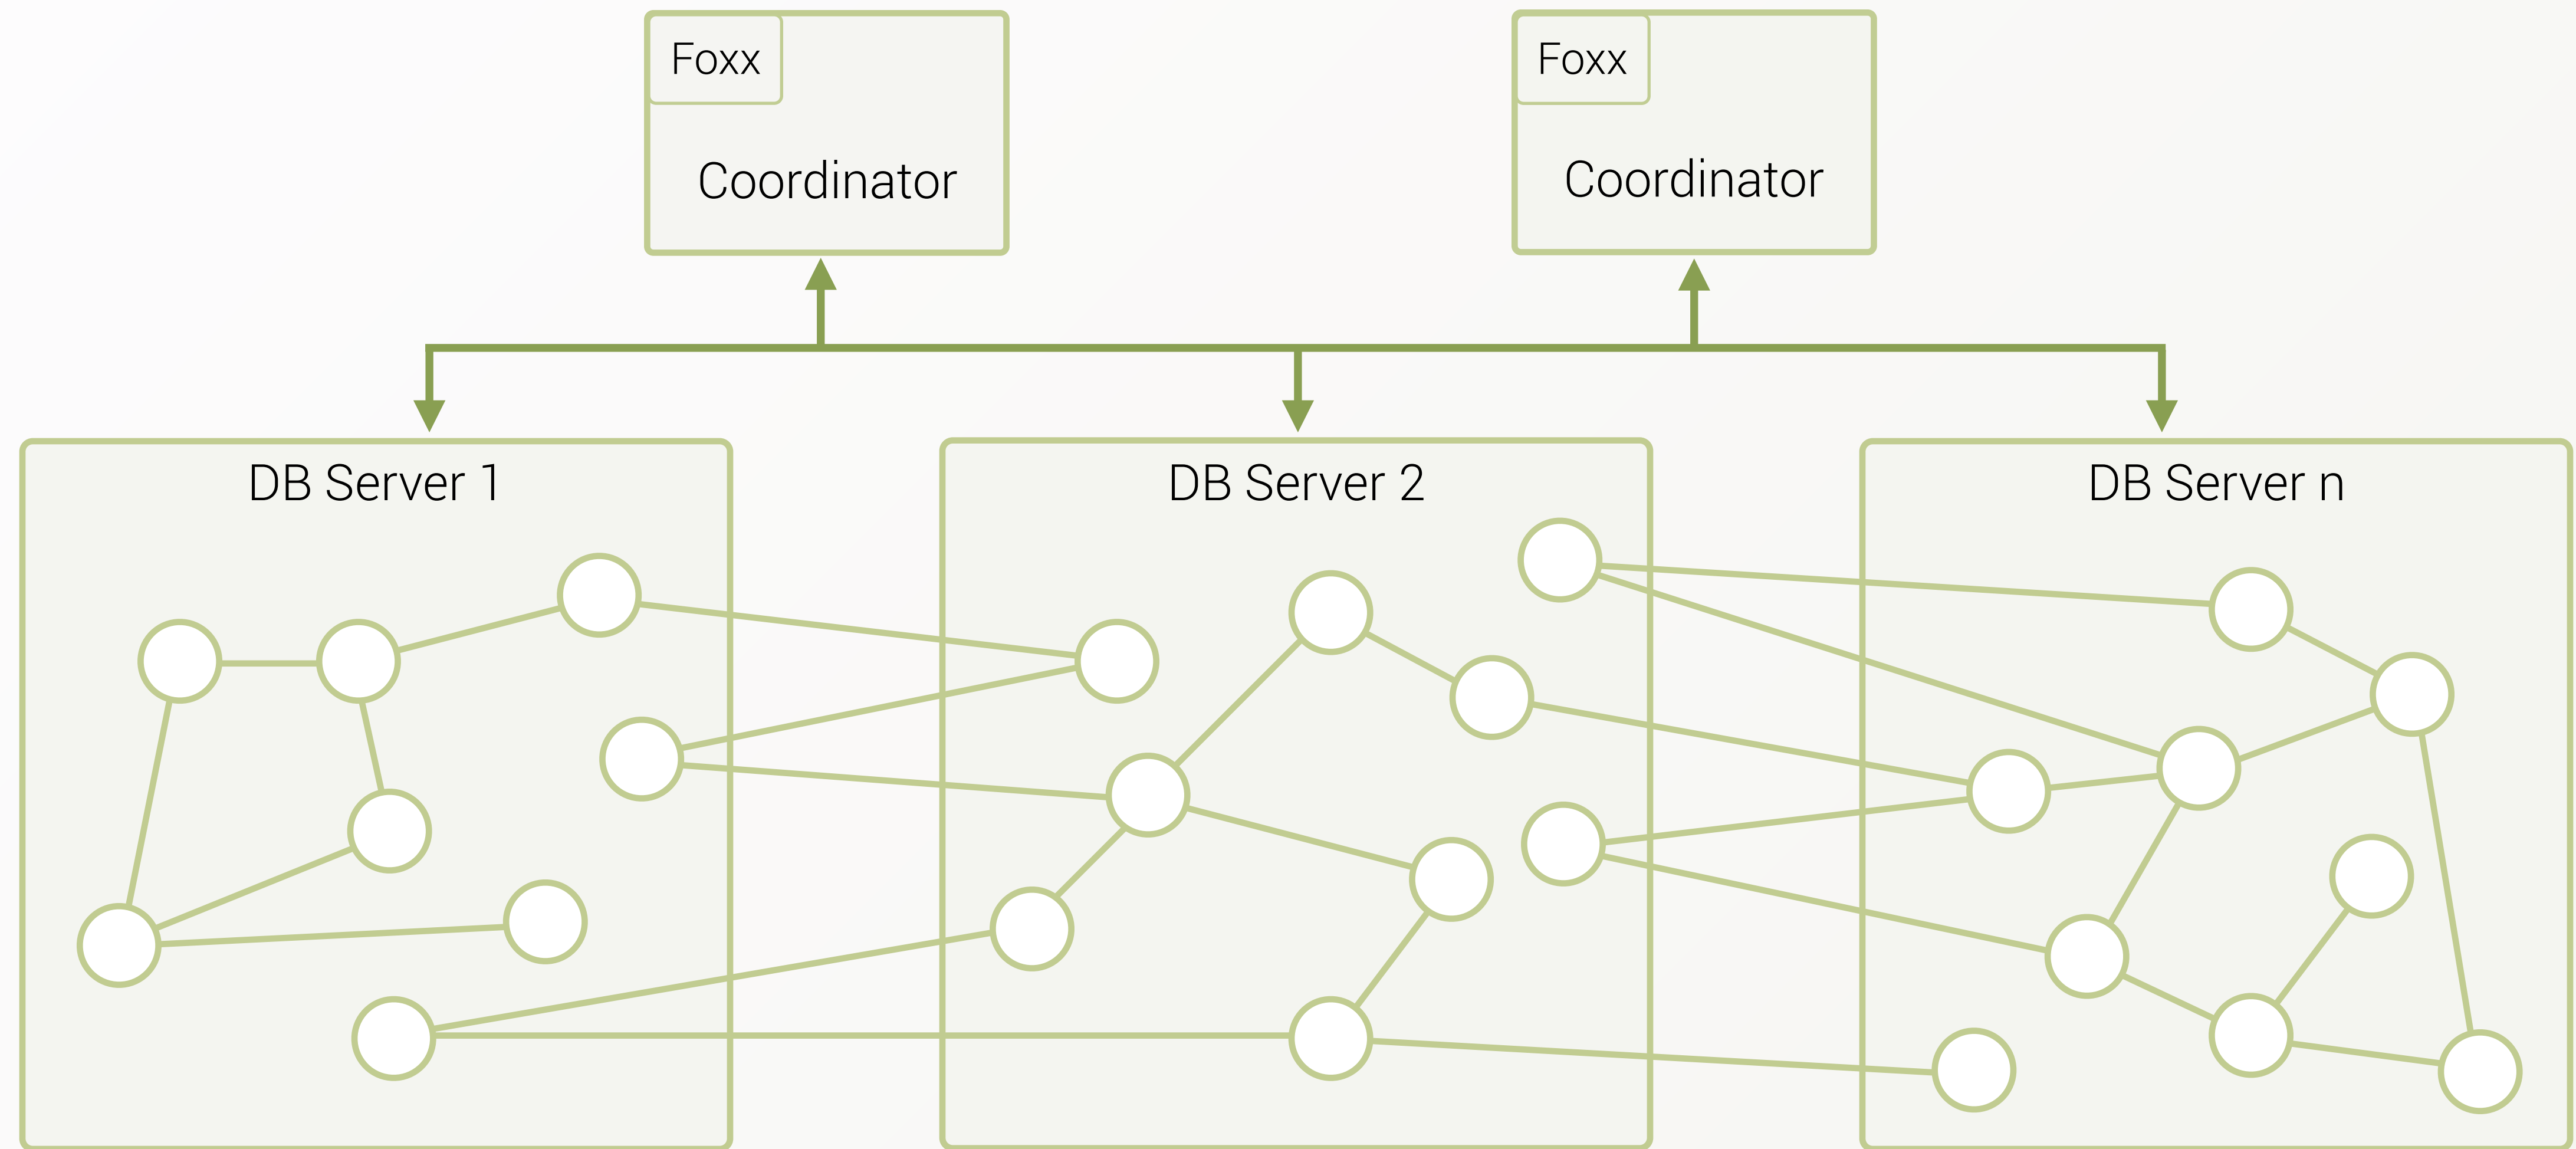
# Domain Based Distribution

---

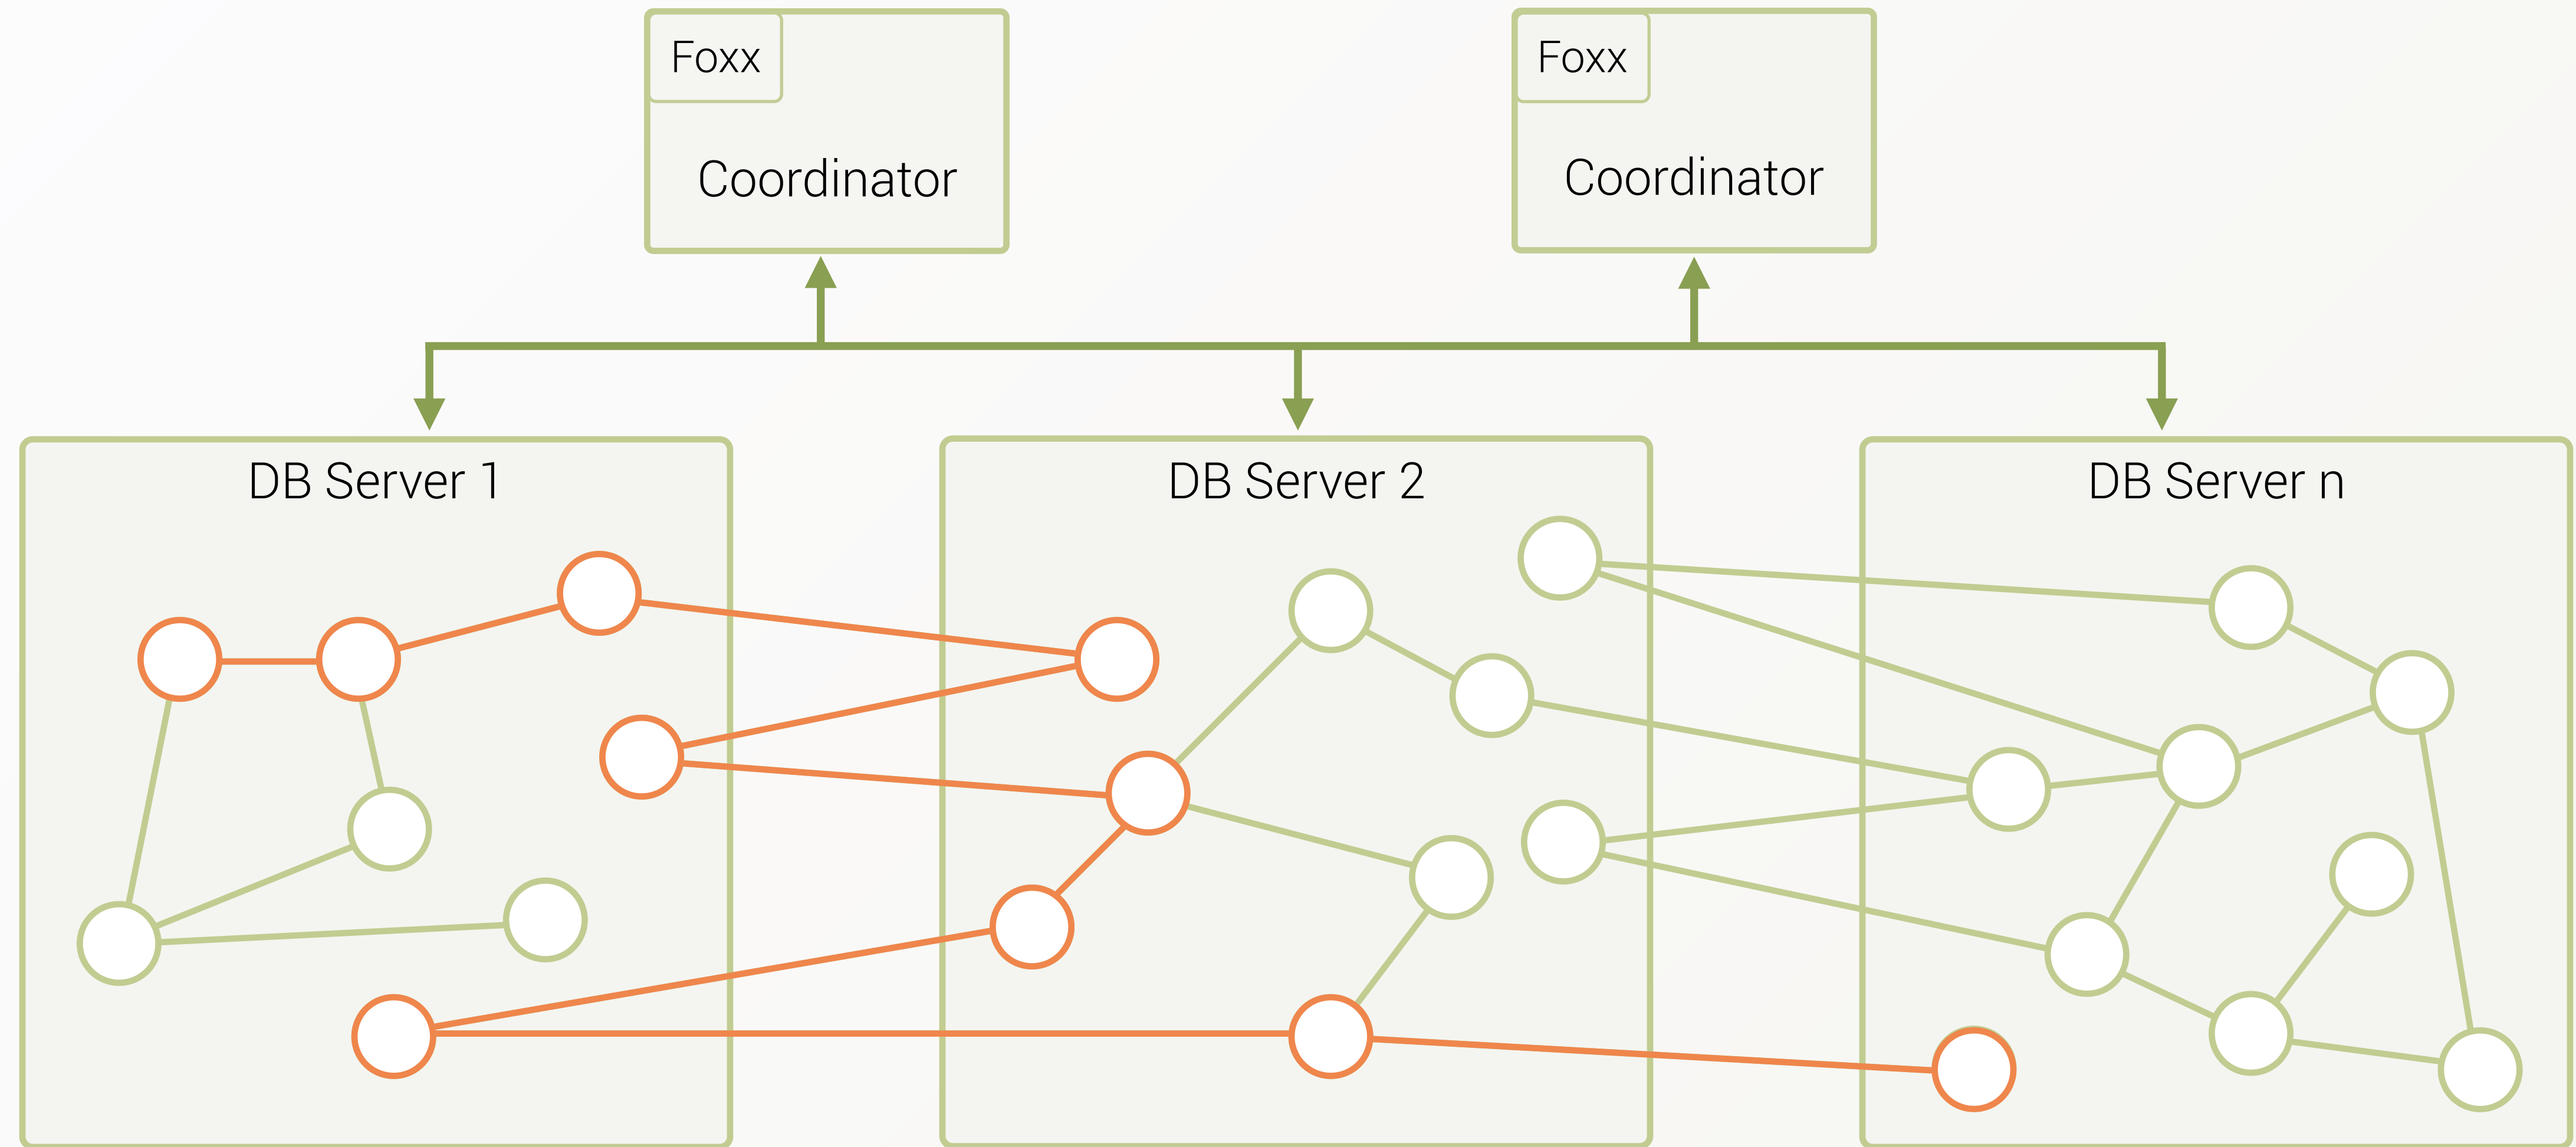
- ▶ Many Graphs have a natural distribution
  - ▶ By country/region for People
  - ▶ By tags for Blogs
  - ▶ By category for Products
- ▶ Most edges in same group
- ▶ Rare edges between groups

ArangoDB Enterprise Edition  
uses Domain Knowledge  
for short-cuts

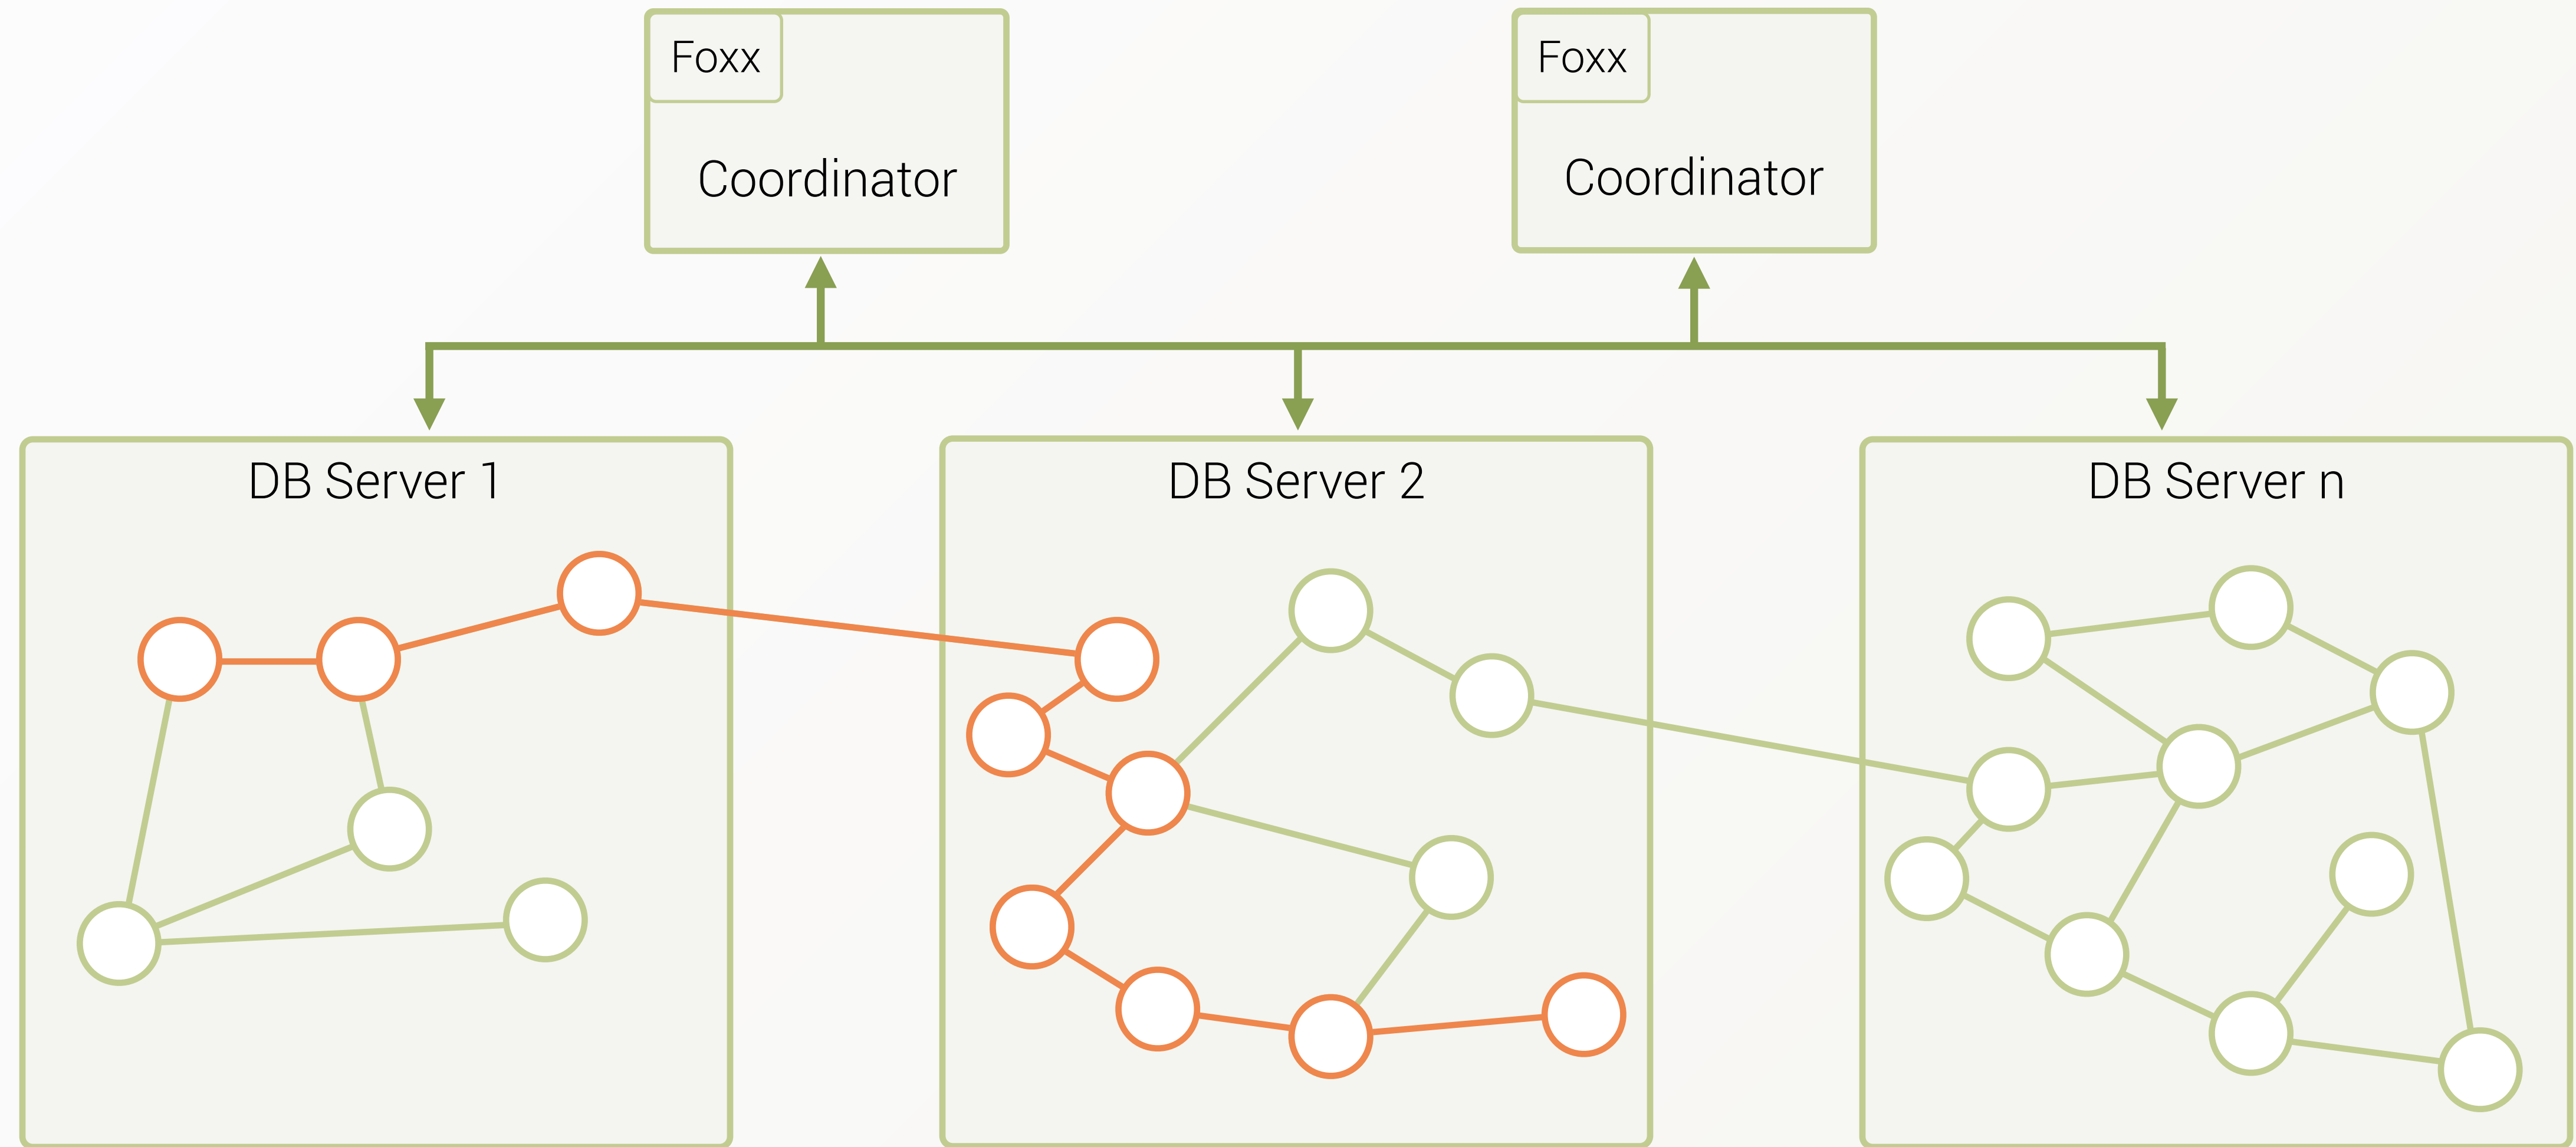
# SmartGraphs - How it works



# SmartGraphs - How it works



# SmartGraphs - How it works





# Sneak Preview SmartGraphs



# NoSQL Performance Test

ArangoDB, Postgres, MongoDB, Neo4j and OrientDB




\*) neighbors and neighbors of neighbors (distinct)

Database versions: ArangoDB 2.7 RC2, OrientDB 2.2 alpha, MongoDB 3.0.6, Neo4J 2.3 M3, PostgreSQL 9.4.4

Weinberger 2015-10-13 (r207)



Database	document write transactions per virtual CPU/second
ArangoDB	1,730
Aerospike	2,500
Cassandra	965
Couchbase	1,375
FoundationDB	750



# Thank You

- ▶ Further questions?
  - ▶ Follow us on twitter: @arangodb
  - ▶ Join our slack: [slack.arangodb.com](https://slack.arangodb.com)
- ▶ Follow me on twitter/github: @mchacki
- ▶ Write me a mail: [michael@arangodb.com](mailto:michael@arangodb.com)