




Handling Billions Of Edges in a Graph Database

Michael Hackstein
@mchacki

Michael Hackstein

- ▶ ArangoDB Core Team

- ▶ Web Frontend
- ▶ Graph visualisation
- ▶ Graph features
- ▶ SmartGraphs



- ▶ Host of cologne.js

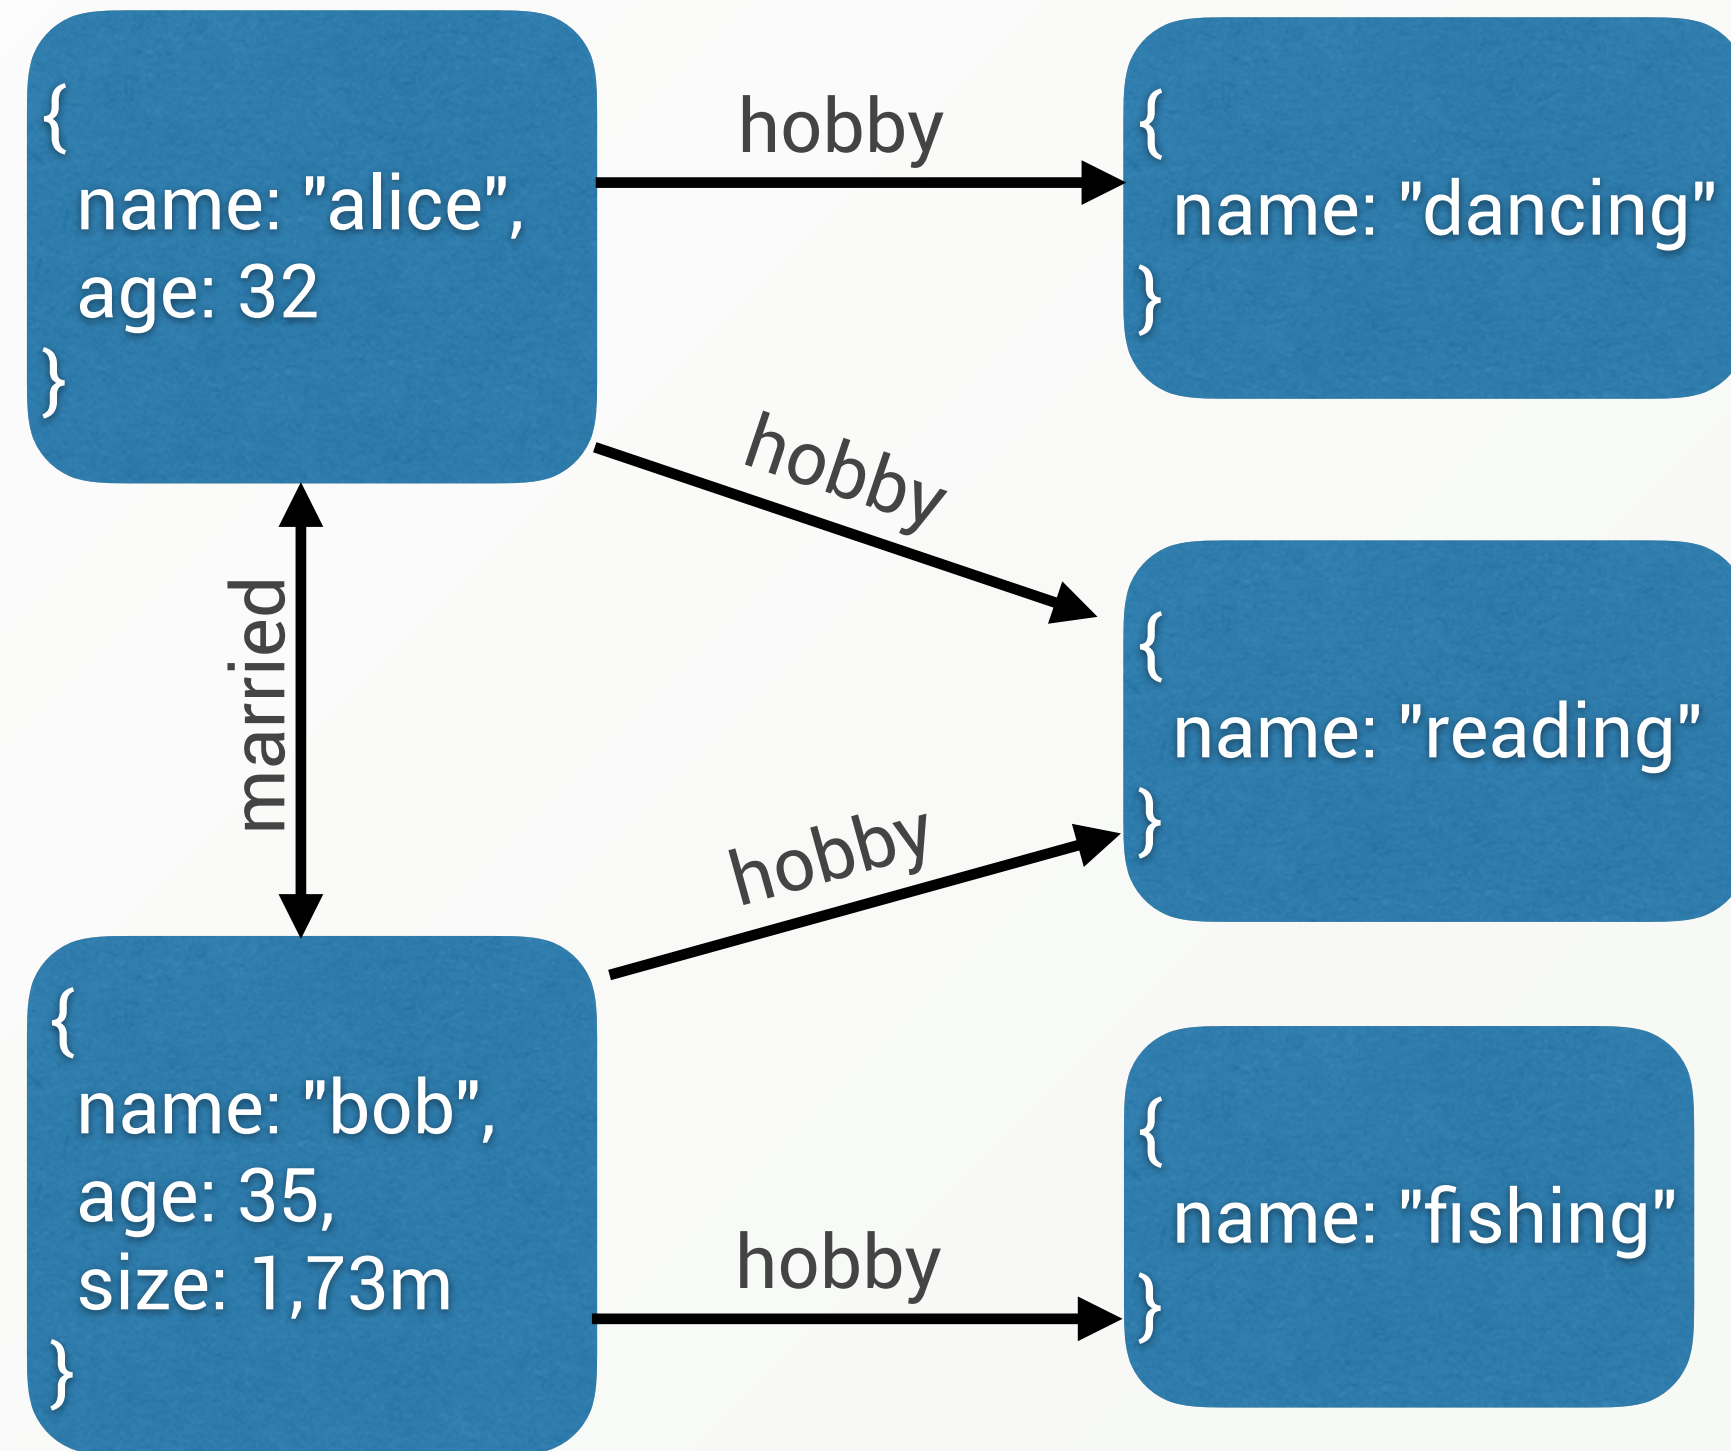


- ▶ Master's Degree
(spec. Databases and
Information Systems)

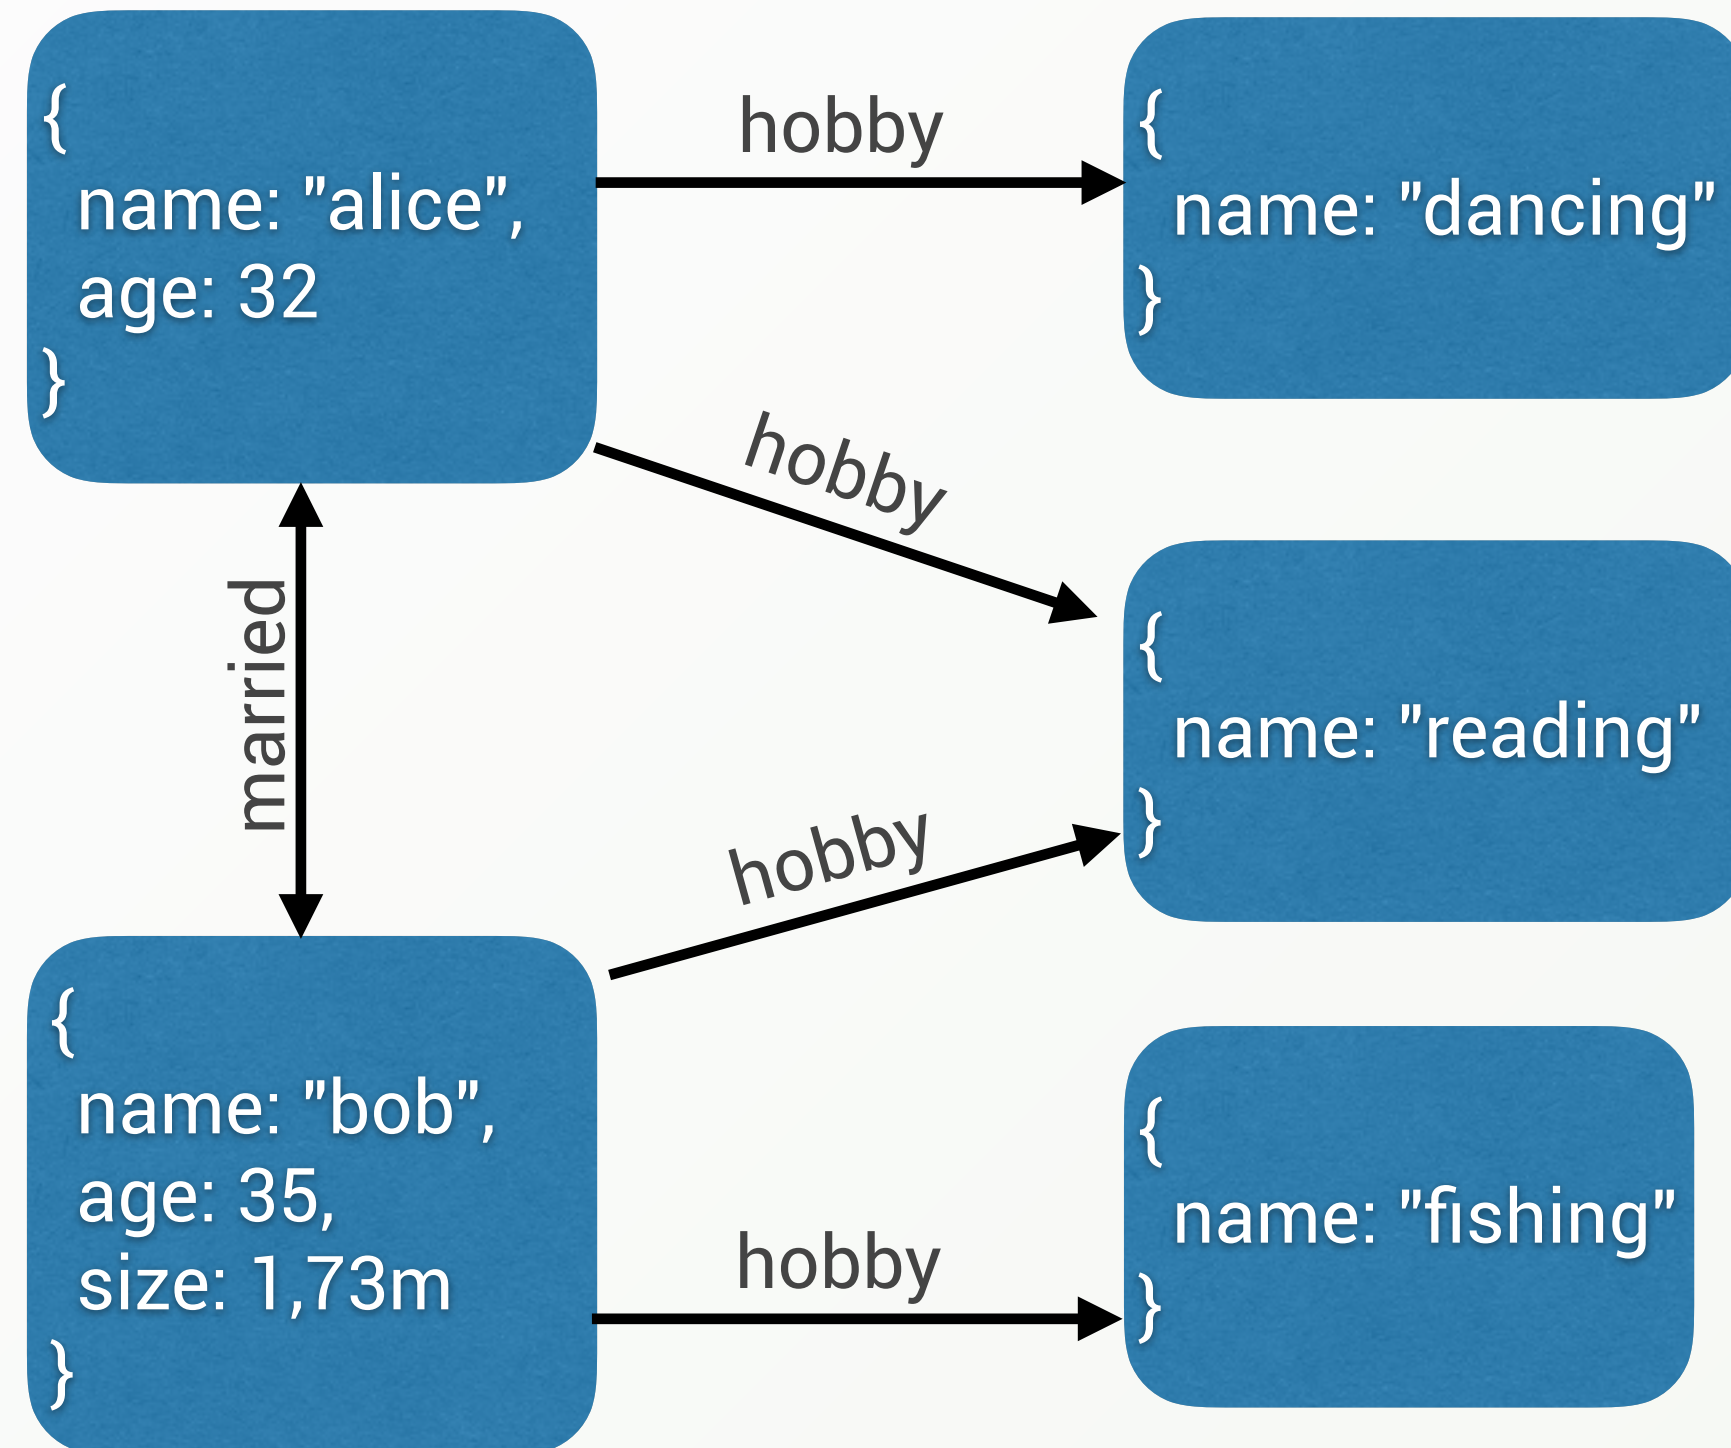


What are Graph Databases

- ▶ Schema-free Objects (Vertices)
- ▶ Relations between them (Edges)
- ▶ Edges have a direction



What are Graph Databases



- ▶ Schema-free Objects (Vertices)
- ▶ Relations between them (Edges)
- ▶ Edges have a direction
- ▶ Edges can be queried in both directions
- ▶ Easily query a range of edges (2 to 5)
- ▶ Undefined number of edges (1 to *)
- ▶ Shortest Path between two vertices

Typical Graph Queries

Typical Graph Queries

- ▶ Give me all friends of **Alice**

Typical Graph Queries

- ▶ Give me all friends of **Alice**
- ▶ Give me all friends-of-friends of **Alice**

Typical Graph Queries

- ▶ Give me all friends of **Alice**
- ▶ Give me all friends-of-friends of **Alice**
- ▶ What is the linking path between **Alice** and **Bob**

Typical Graph Queries

- ▶ Give me all friends of **Alice**
- ▶ Give me all friends-of-friends of **Alice**
- ▶ What is the linking path between **Alice** and **Bob**
- ▶ Which **Trainstations** can I reach if I am allowed to drive **a distance of 6 stations** on my ticket

Typical Graph Queries

- ▶ Give me all friends of **Alice**
- ▶ Give me all friends-of-friends of **Alice**
- ▶ What is the linking path between **Alice** and **Bob**
- ▶ Which **Trainstations** can I reach if I am allowed to drive **a distance of 6 stations** on my ticket
- ▶ Pattern Matching:

Typical Graph Queries

- ▶ Give me all friends of **Alice**
- ▶ Give me all friends-of-friends of **Alice**
- ▶ What is the linking path between **Alice** and **Bob**
- ▶ Which **Trainstations** can I reach if I am allowed to drive **a distance of 6 stations** on my ticket
- ▶ Pattern Matching:
 - ▶ Give me all **users** that share two **hobbies** with **Alice**

Typical Graph Queries

- ▶ Give me all friends of **Alice**
- ▶ Give me all friends-of-friends of **Alice**
- ▶ What is the linking path between **Alice** and **Bob**
- ▶ Which **Trainstations** can I reach if I am allowed to drive **a distance of 6 stations** on my ticket
- ▶ Pattern Matching:
 - ▶ Give me all **users** that share two **hobbies** with **Alice**
 - ▶ Give me all products that at least one of my friends has bought together with the products I already own, ordered by how many friends have bought it and the products rating, but only 20 of them.

Non-Typical Graph Queries

Non-Typical Graph Queries

- ▶ Give me all users which have an **age** attribute between **21** and **35**.

Non-Typical Graph Queries

- ▶ Give me all users which have an **age** attribute between **21** and **35**.
- ▶ Give me the **age** distribution of all **users**

Non-Typical Graph Queries

- ▶ Give me all users which have an **age** attribute between **21** and **35**.
- ▶ Give me the **age** distribution of all **users**
- ▶ Group all **users** by their **name**

Traversal

Iterate down two edges with some filters

Traversal


Iterate down two edges with some filters



- ▶ We first pick a start vertex (S)

Traversal


Iterate down two edges with some filters



- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S

Traversal


Iterate down two edges with some filters



- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges

Traversal


Iterate down two edges with some filters



- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)

Traversal


Iterate down two edges with some filters



- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)
- ▶ We apply filters on edges

Traversal


Iterate down two edges with some filters



- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)
- ▶ We apply filters on edges
- ▶ The next vertex (E) is in desired depth. Return the path S -> A -> E

Traversal


Iterate down two edges with some filters



- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)
- ▶ We apply filters on edges
- ▶ The next vertex (E) is in desired depth. Return the path S -> A -> E
- ▶ Go back to the next unfinished vertex (B)

Traversal


Iterate down two edges with some filters



- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)
- ▶ We apply filters on edges
- ▶ The next vertex (E) is in desired depth. Return the path S -> A -> E
- ▶ Go back to the next unfinished vertex (B)
- ▶ We iterate down on (B)

Traversal


Iterate down two edges with some filters



- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)
- ▶ We apply filters on edges
- ▶ The next vertex (E) is in desired depth. Return the path S -> A -> E
- ▶ Go back to the next unfinished vertex (B)
- ▶ We iterate down on (B)
- ▶ We apply filters on edges

Traversal

Iterate down two edges with some filters



- ▶ We first pick a start vertex (S)
- ▶ We collect all edges on S
- ▶ We apply filters on edges
- ▶ We iterate down one of the new vertices (A)
- ▶ We apply filters on edges
- ▶ The next vertex (E) is in desired depth. Return the path S -> A -> E
- ▶ Go back to the next unfinished vertex (B)
- ▶ We iterate down on (B)
- ▶ We apply filters on edges
- ▶ The next vertex (F) is in desired depth. Return the path S -> B -> F

Traversal - Complexity

| | | |
|--------------------------------|---------------------------|---------|
| ▶ Once: | | 0 |
| ▶ Find the start vertex | Depends on indexes: Hash: | 1 |
| ▶ For every depth: | | |
| ▶ Find all connected edges | Edge-Index or Index-Free: | 1 |
| ▶ Filter non-matching edges | Linear in edges: | n |
| ▶ Find connected vertices | Depends on indexes: Hash: | $n * 1$ |
| ▶ Filter non-matching vertices | Linear in vertices: | n |
| | Only one pass: | 3n |

Traversal - Complexity

- ▶ Linear sounds evil?
 - ▶ NOT linear in All Edges $O(E)$
 - ▶ Only Linear in relevant Edges $n < E$
- ▶ Traversals solely scale with their result size.
- ▶ They are not effected at all by total amount of data
- ▶ BUT: Every depth increases the exponent: $O(3 * n^d)$
- ▶ "7 degrees of separation": $3 * n^6 < E < 3 * n^7$

- ▶ MULTI-MODEL database
 - ▶ Stores Documents and Graphs
- ▶ Query language AQL
 - ▶ Document Queries
 - ▶ Graph Queries
 - ▶ Joins
 - ▶ All can be combined in the same statement
- ▶ ACID support including Multi Collection Transactions

AQL

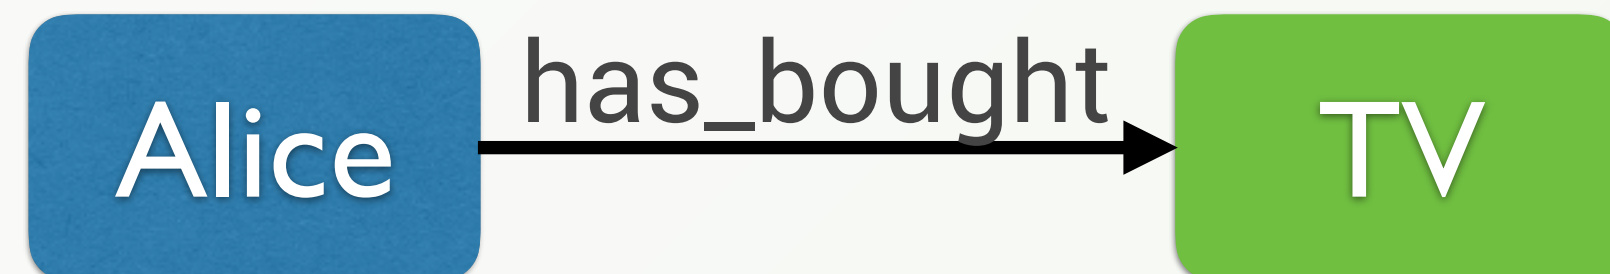
```
FOR user IN users  
  RETURN user
```

AQL

```
FOR user IN users  
  FILTER user.name == "alice"  
  RETURN user
```

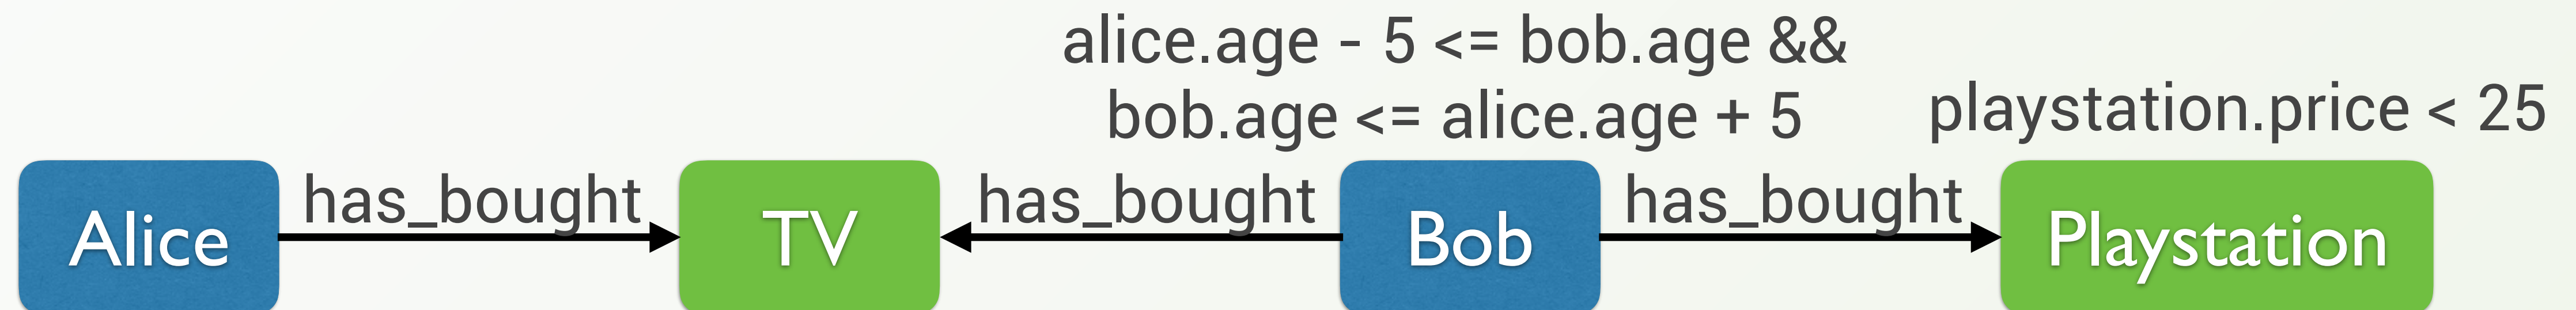
AQL

```
FOR user IN users  
  FILTER user.name == "alice"  
  FOR product IN OUTBOUND user has_bought  
  RETURN product
```



AQL

```
FOR user IN users
  FILTER user.name == "alice"
  FOR recommendation, action, path IN 3 ANY user has_bought
    FILTER path.vertices[2].age <= user.age + 5
    AND path.vertices[2].age >= user.age - 5
    FILTER recommendation.price < 25
  LIMIT 10
  RETURN recommendation
```



Demo Time

Querying basics

First Boost - Vertex Centric Indices

- ▶ Remember Complexity? $O(3 * n^d)$
- ▶ Filtering of non-matching edges is linear for every depth
- ▶ Index all edges based on their vertices and arbitrary other attributes
 - ▶ Find initial set of edges in identical time
 - ▶ Less / No post-filtering required
 - ▶ This decreases the n

Demo Time

Vertex-Centric Indices

Scaling

- ▶ Vertex-Centric Indexes help with super-nodes
- ▶ But what if the graph is too large for one machine?
- ▶ Distribute graph on several machines (sharding)
- ▶ How to query it now?
 - ▶ No global view of the graph possible any more
 - ▶ What about edges between servers?


**First let's do
the cluster thingy**



ArangoDB



MESOSPHERE



Marathon



DC/OS

Demo Time
DC / OS

Is Mesosphere required?

- ▶ ArangoDB can run clusters without it
 - ▶ Setup Requires manual effort (can be scripted):
 - ▶ Configure IP addresses
 - ▶ Correct startup ordering
 - ▶ This works:
 - ▶ Automatic Failover (Follower takes over if leader dies)
 - ▶ Rebalancing of shards
 - ▶ Everything inside of ArangoDB
 - ▶ This is based on Mesos:
 - ▶ Complete self healing
 - ▶ Automatic restart of ArangoDBs (on new machines)
- ➡ We suggest you have someone on call

**Now distribute
the graph**

Dangers of Sharding

- ▶ Only parts of the graph on every machine
- ▶ Neighboring vertices may be on different machines
- ▶ Even edges could be on other machines than their vertices

- ▶ Queries need to be executed in a distributed way
- ▶ Result needs to be merged locally



Random Distribution

▶ Advantages:

- ▶ every server takes an equal portion of data
- ▶ easy to realize
- ▶ no knowledge about data required
- ▶ always works

▶ Disadvantages:

- ▶ Neighbors on different machines
- ▶ Probably edges on other machines than their vertices
- ▶ A lot of network overhead is required for querying




Random Distribution

▶ Advantages:

- ▶ every server takes an equal portion of data
- ▶ easy to realize
- ▶ no knowledge about data required
- ▶ always works



▶ Disadvantages:

- ▶ Neighbors on different machines
- ▶ Probably edges on other machines than their vertices
- ▶ A lot of network overhead is required for querying




Index-Free Adjacency

- ▶ Used by most other graph databases
- ▶ Every vertex maintains two lists of its edges (IN and OUT)
 - ▶ Do not use an index to find edges
 - ▶ How to shard this?



Index-Free Adjacency


- ▶ Used by most other graph databases
- ▶ Every vertex maintains two lists of its edges (IN and OUT)
 - ▶ Do not use an index to find edges
 - ▶ How to shard this?



Index-Free Adjacency

- ▶ Used by most other graph databases
- ▶ Every vertex maintains two lists of its edges (IN and OUT)
 - ▶ Do not use an index to find edges
 - ▶ How to shard this?


————



Index-Free Adjacency

- ▶ Used by most other graph databases
- ▶ Every vertex maintains two lists of its edges (IN and OUT)
 - ▶ Do not use an index to find edges
 - ▶ How to shard this?


????



Index-Free Adjacency

- ▶ Used by most other graph databases
- ▶ Every vertex maintains two lists of its edges (IN and OUT)
 - ▶ Do not use an index to find edges
 - ▶ How to shard this?



????



- ▶ ArangoDB uses an hash-based EdgeIndex ($O(1)$ - lookup)
 - ▶ The vertex is independent of its edges
 - ▶ It can be stored on a different machine


Domain Based Distribution

- ▶ Many Graphs have a natural distribution
 - ▶ By country/region for People
 - ▶ By tags for Blogs
 - ▶ By category for Products
- ▶ Most edges in same group
- ▶ Rare edges between groups



Domain Based Distribution

- ▶ Many Graphs have a natural distribution
 - ▶ By country/region for People
 - ▶ By tags for Blogs
 - ▶ By category for Products
- ▶ Most edges in same group
- ▶ Rare edges between groups



Domain Based Distribution

- ▶ Many Graphs have a natural distribution
 - ▶ By country/region for People
 - ▶ By tags for Blogs
 - ▶ By category for Products
- ▶ Most edges in same group
- ▶ Rare edges between groups



**uses Domain Knowledge
for short-cuts**

Sneak Preview

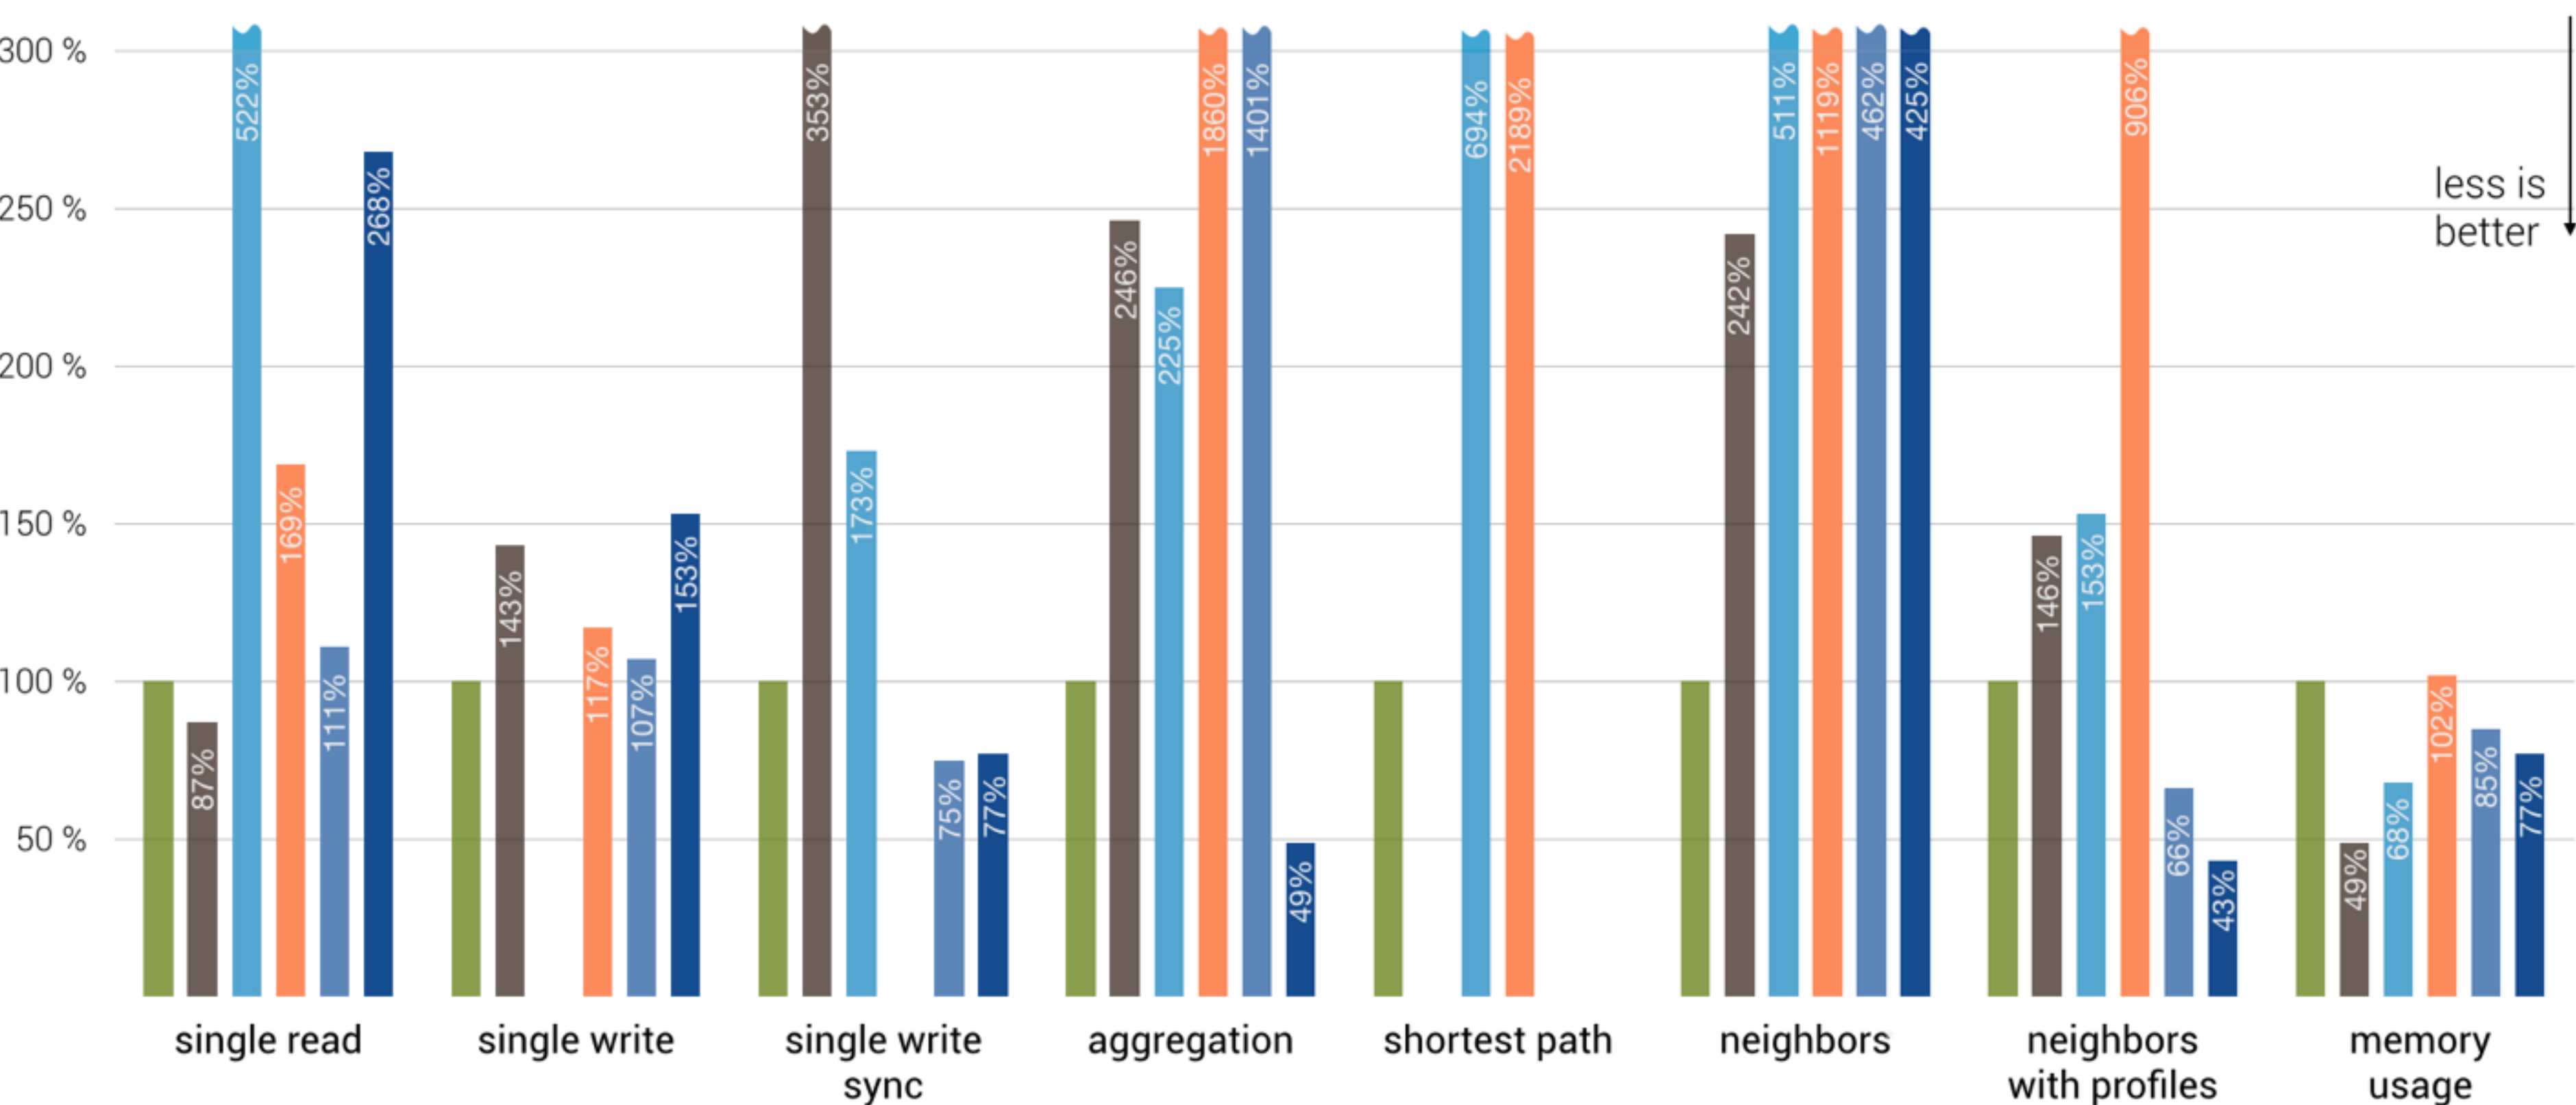



SmartGraphs

Benchmark Comparison

NoSQL Performance Test

ArangoDB, Postgres, MongoDB, Neo4j and OrientDB



*) neighbors and neighbors of neighbors (distinct)

Database versions: ArangoDB 2.7 RC2, OrientDB 2.2 alpha, MongoDB 3.0.6, Neo4J 2.3 M3, PostgreSQL 9.4.4

Weinberger 2015-10-13 (r207)

Source: <https://www.arangodb.com/2015/10/benchmark-postgresql-mongodb-arangodb/>

Thank you

- ▶ Further questions?
 - ▶ Follow us on twitter: @arangodb
 - ▶ Join our slack: slack.arangodb.com
- ▶ Follow me on twitter/github: @mchacki
- ▶ Write me a mail: michael@arangodb.com