



Polyglot Persistence & Multi-Model Databases

Java User Group Hessen

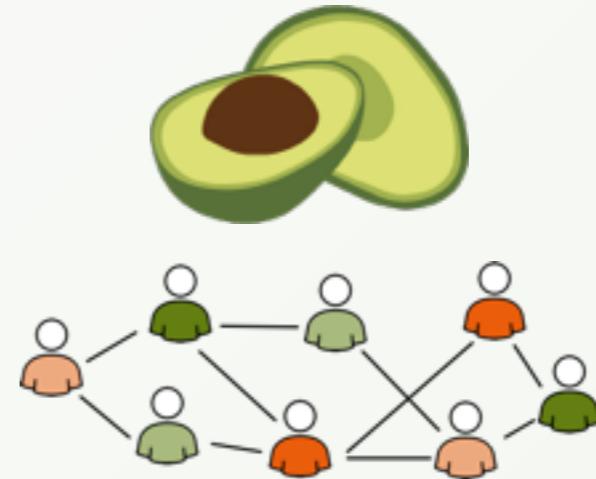
Michael Hackstein
@mchacki

www.arangodb.com

Michael Hackstein

- ▶ ArangoDB Core Team

- ▶ Web Frontend
- ▶ Graph visualisation
- ▶ Graph features



- ▶ Host of cologne.js

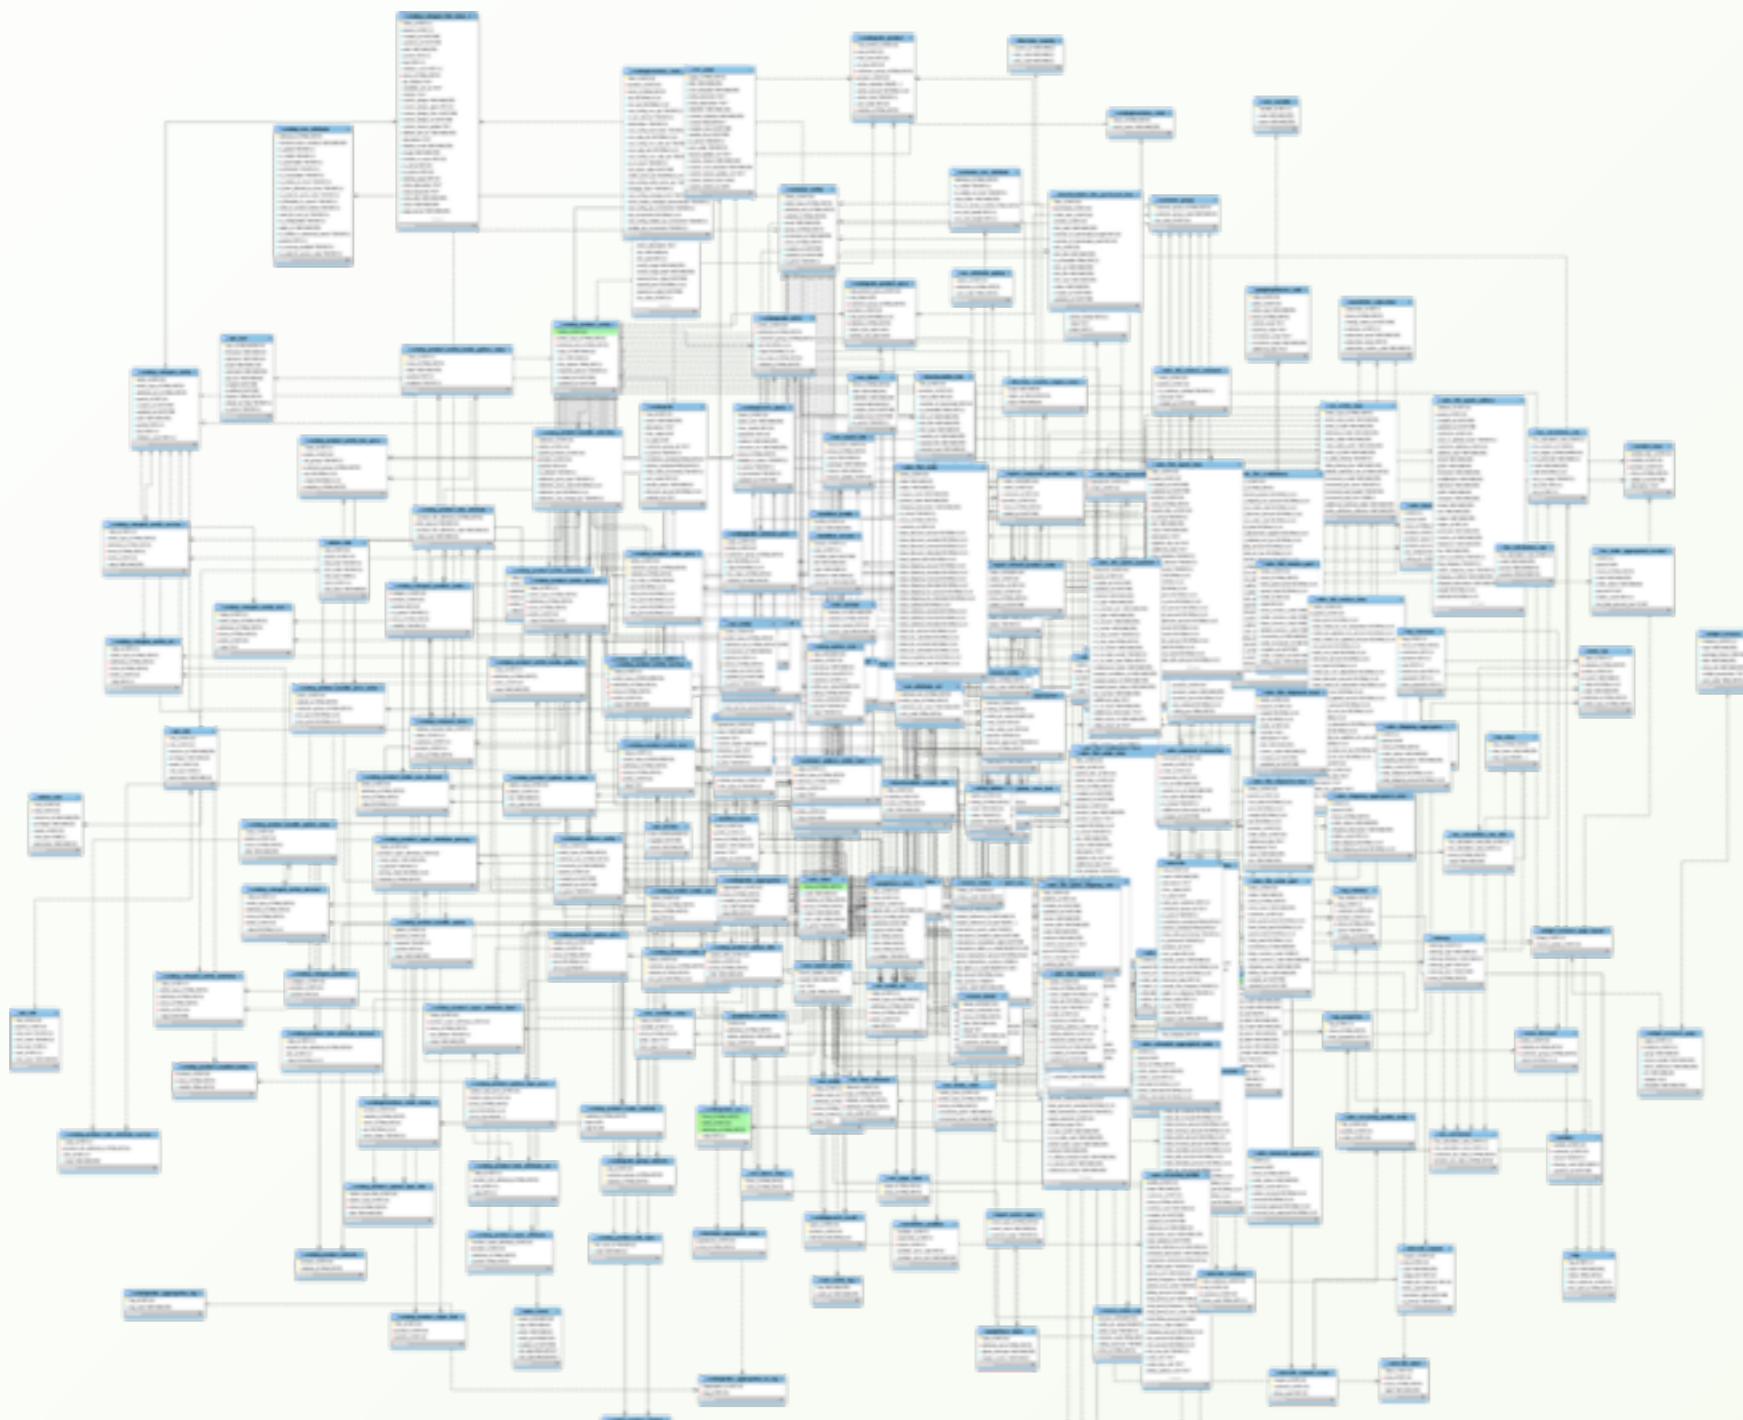


- ▶ Master's Degree
(spec. Databases and
Information Systems)



The Single Model era is over

Relational World



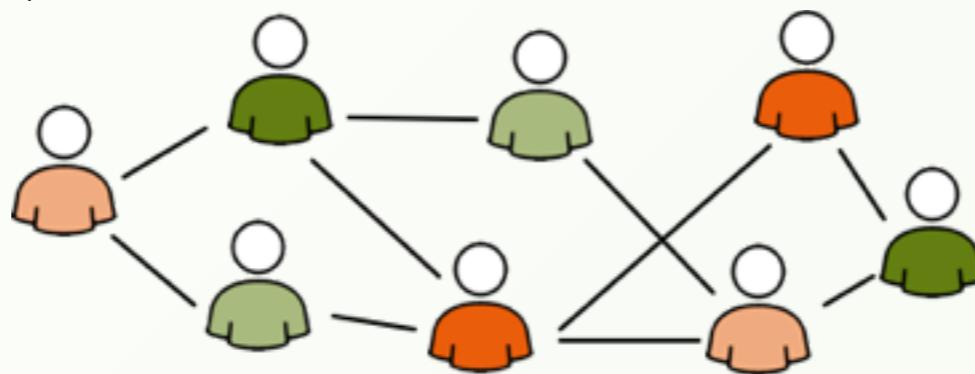
The Multi Model era begins

NoSQL World

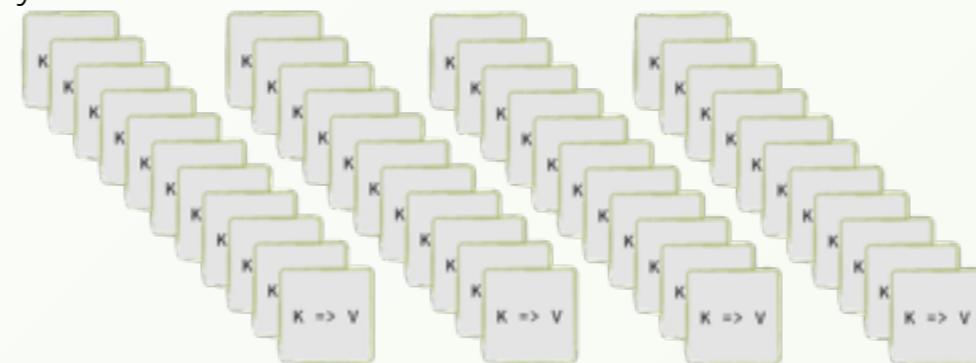
Documents - JSON

```
{  
  "type": "sweater",  
  "color": "blue",  
  "size": "M",  
  "material": "wool",  
  "form": "turtleneck"  
}  
  
{  
  "type": "pants",  
  "waist": 32,  
  "length": 34,  
  "color": "blue",  
  "material": "cotton"  
}  
  
{  
  "type": "sweater",  
  "color": "blue",  
  "size": "M",  
  "material": "wool",  
  "form": "turtleneck"  
}  
  
{  
  "type": "television",  
  "diagonal screen size": 46,  
  "hdmi inputs": 3,  
  "wall mountable": true,  
  "built-in digital tuner": true,  
  "dynamic contrast ratio": "50,000:1",  
  "Resolution": "1920x1080"  
}
```

Graphs



Key Value



The Multi Model era begins

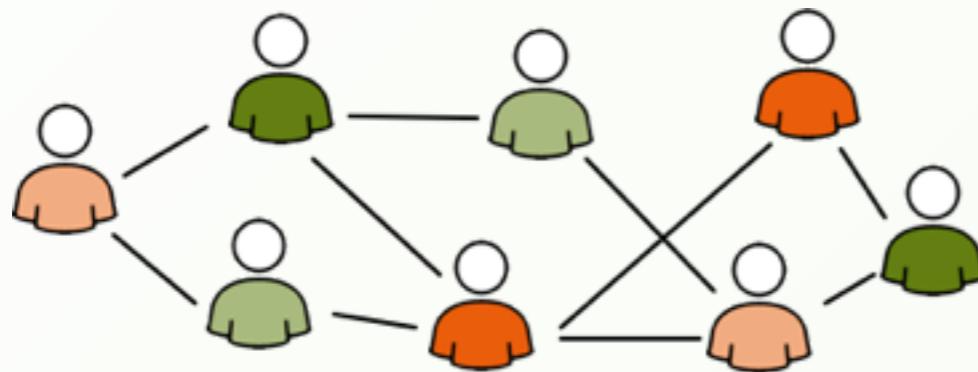
NoSQL World

Documents - JSON

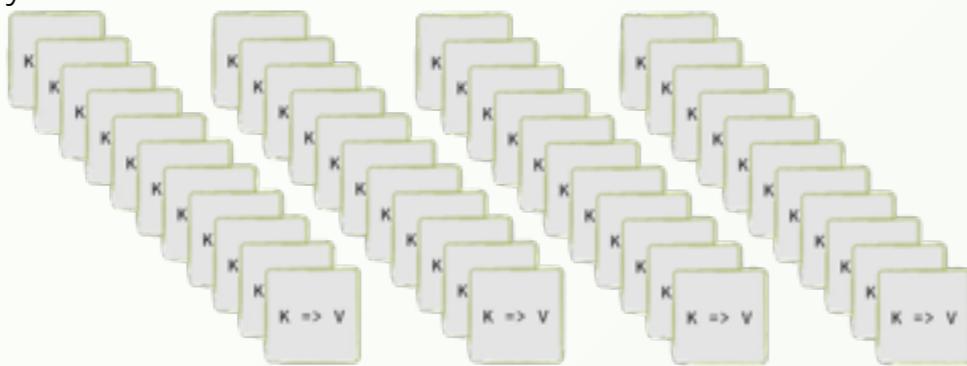
```
{  
  "type": "sweater",  
  "color": "blue",  
  "size": "M",  
  "material": "wool",  
  "form": "turtleneck"  
}  
  
{  
  "type": "pants",  
  "waist": 32,  
  "length": 34,  
  "color": "blue",  
  "material": "cotton"  
}  
  
{  
  "type": "sweater",  
  "color": "blue",  
  "size": "M",  
  "material": "wool",  
  "form": "turtleneck"  
}  
  
{  
  "type": "television",  
  "diagonal screen size": 46,  
  "hdmi inputs": 3,  
  "wall mountable": true,  
  "built-in digital tuner": true,  
  "dynamic contrast ratio": "50,000:1",  
  "Resolution": "1920x1080"  
}
```

- ▶ Normally based on key-value stores (each document still has a unique key)
- ▶ Allow to save documents with logical similarity in "collections"
- ▶ Treat data records as attribute-structured documents (data is no more opaque)
- ▶ Often allow querying and indexing document attributes

Graphs



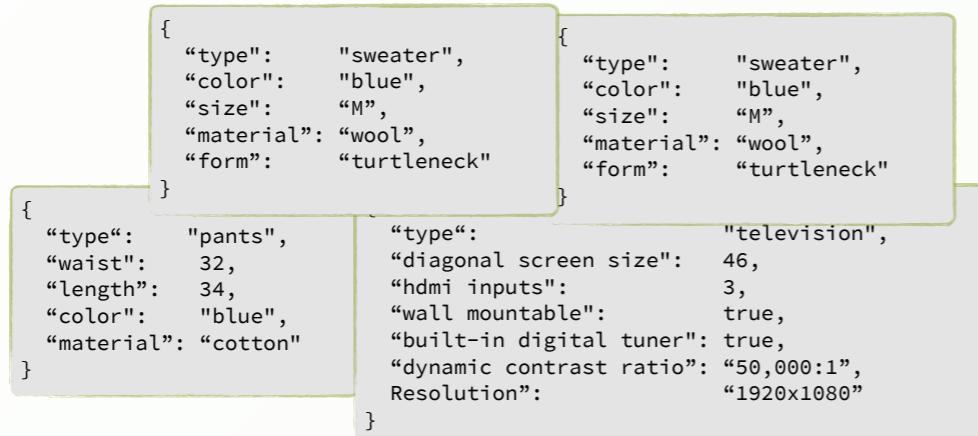
Key Value



The Multi Model era begins

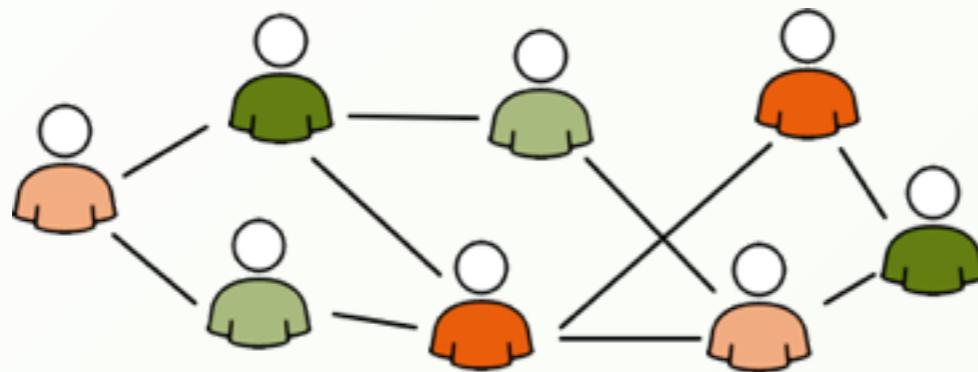
NoSQL World

Documents - JSON



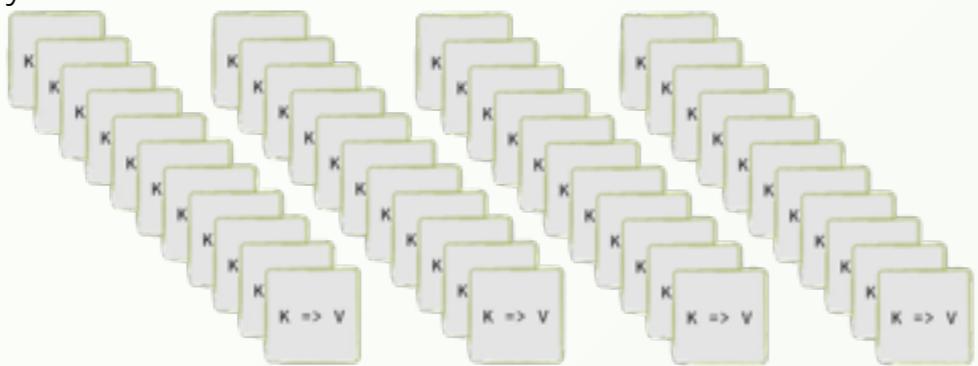
- ▶ Normally based on key-value stores (each document still has a unique key)
- ▶ Allow to save documents with logical similarity in "collections"
- ▶ Treat data records as attribute-structured documents (data is no more opaque)
- ▶ Often allow querying and indexing document attributes

Graphs



- ▶ Focussed on m-to-n relations between entities
- ▶ Stores property graphs: entities and edges can have attributes
- ▶ Easily query paths of variable length

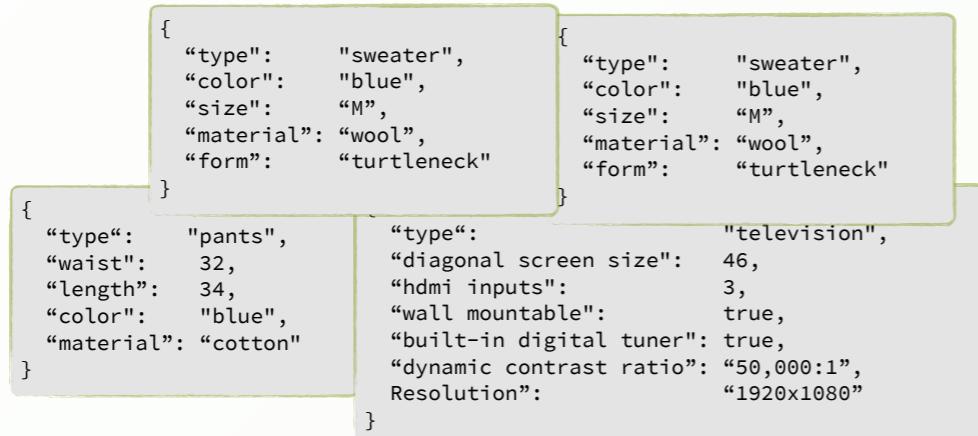
Key Value



The Multi Model era begins

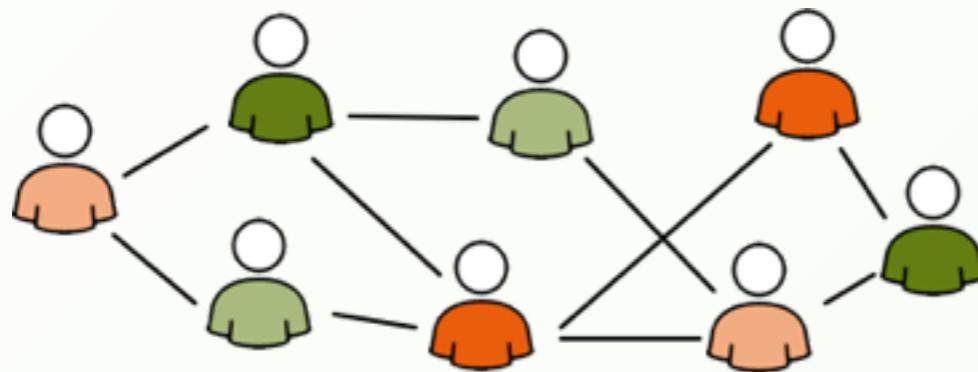
NoSQL World

Documents - JSON



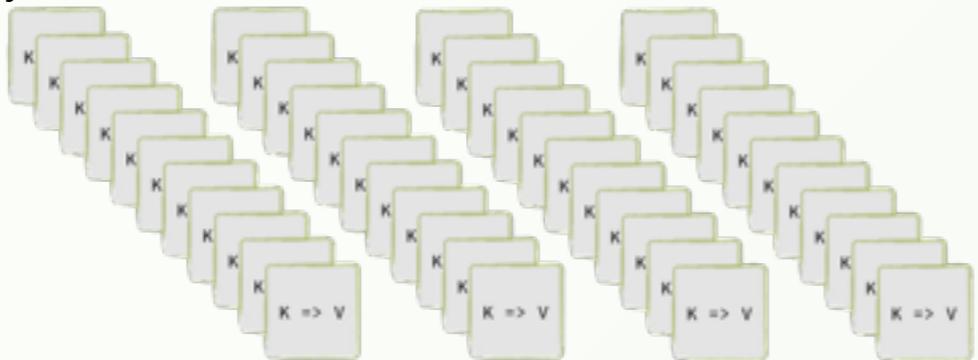
- ▶ Normally based on key-value stores (each document still has a unique key)
- ▶ Allow to save documents with logical similarity in "collections"
- ▶ Treat data records as attribute-structured documents (data is no more opaque)
- ▶ Often allow querying and indexing document attributes

Graphs



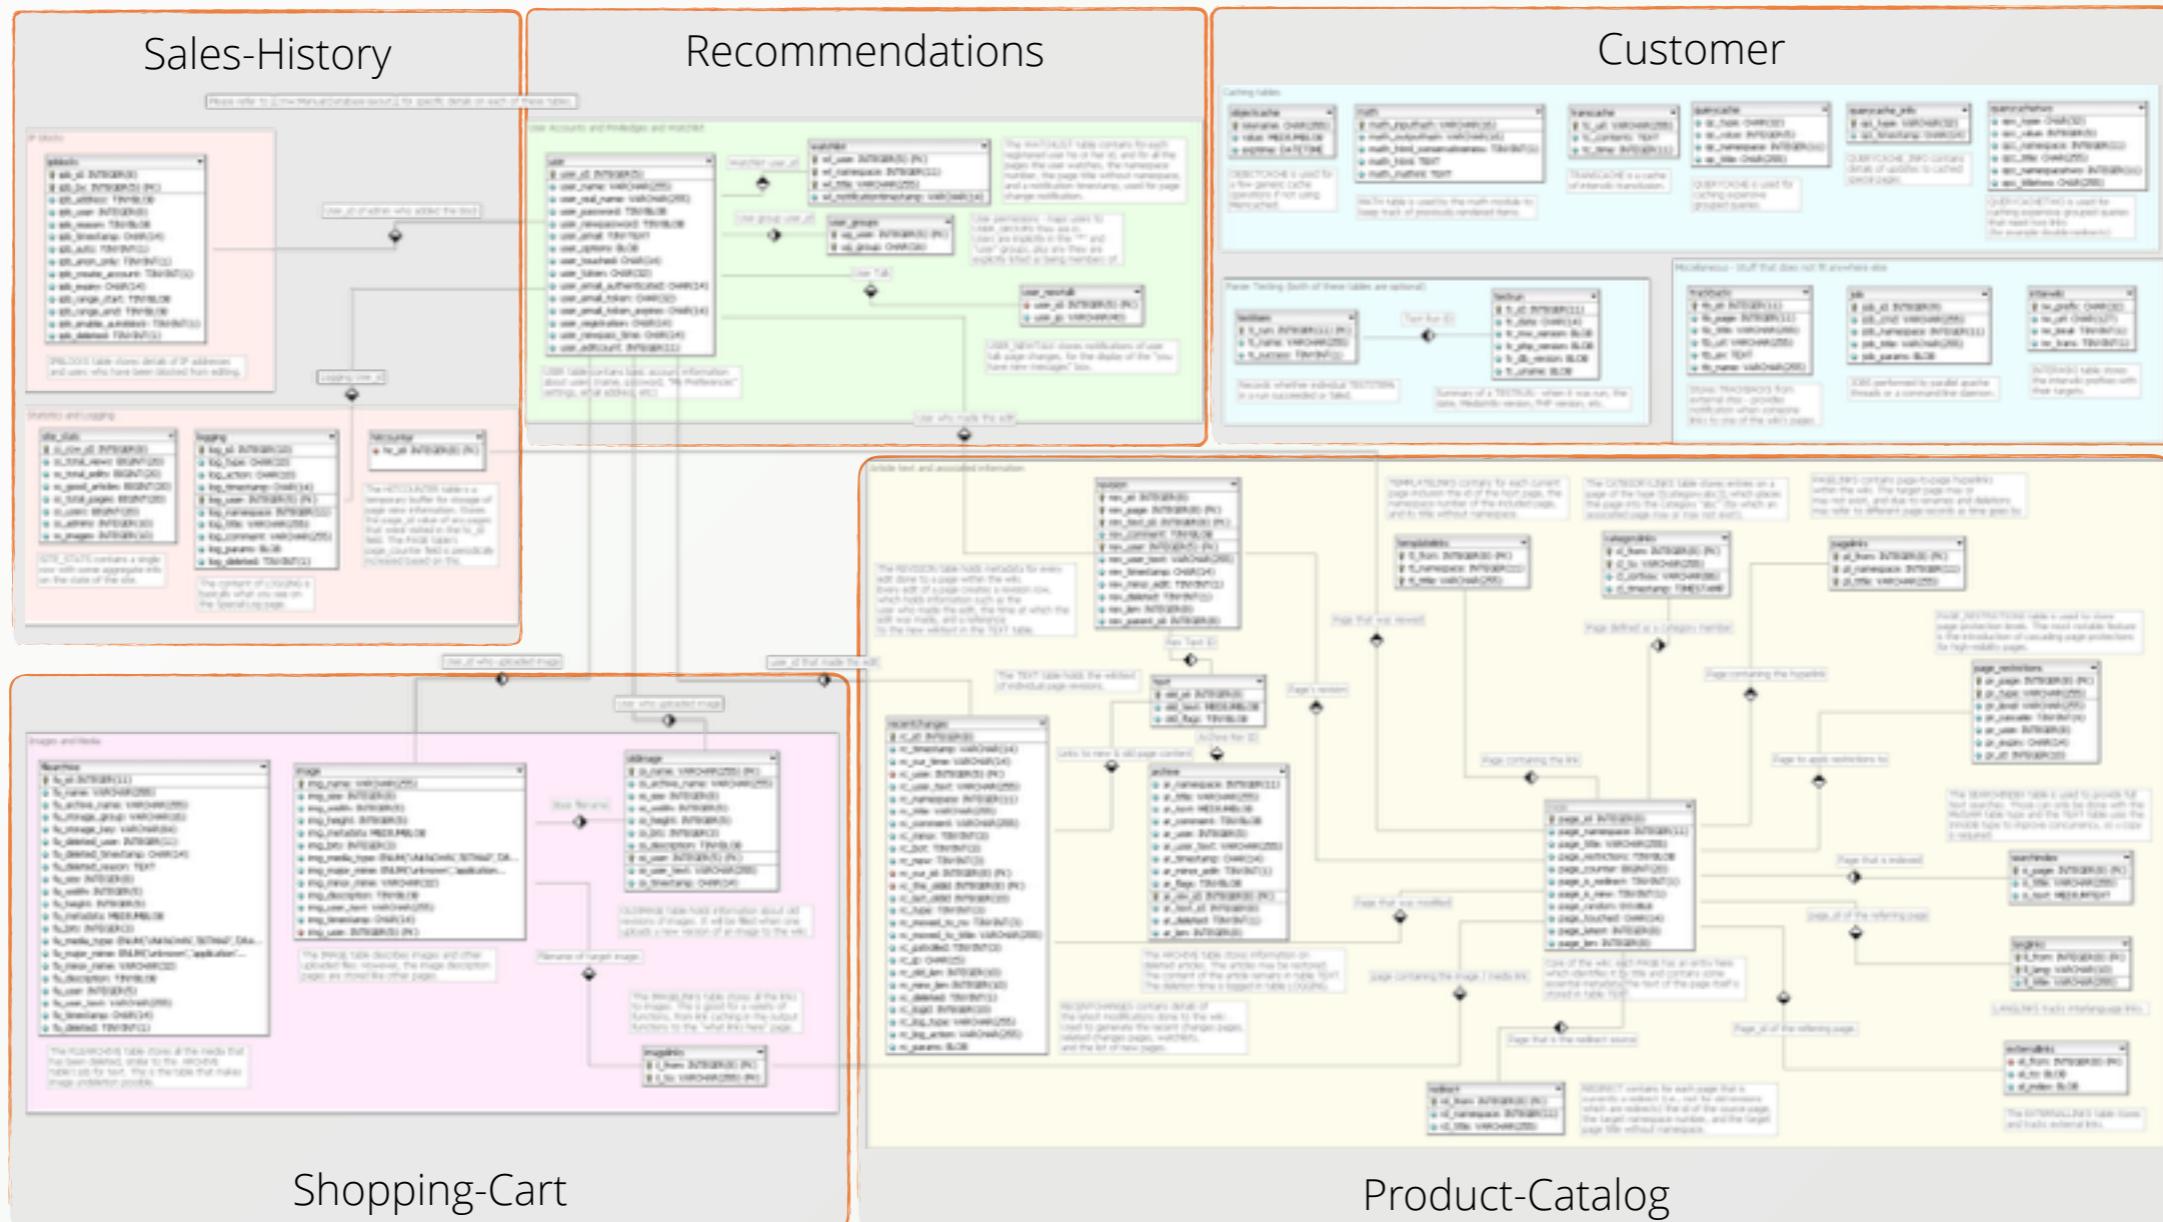
- ▶ Focussed on m-to-n relations between entities
- ▶ Stores property graphs: entities and edges can have attributes
- ▶ Easily query paths of variable length

Key Value

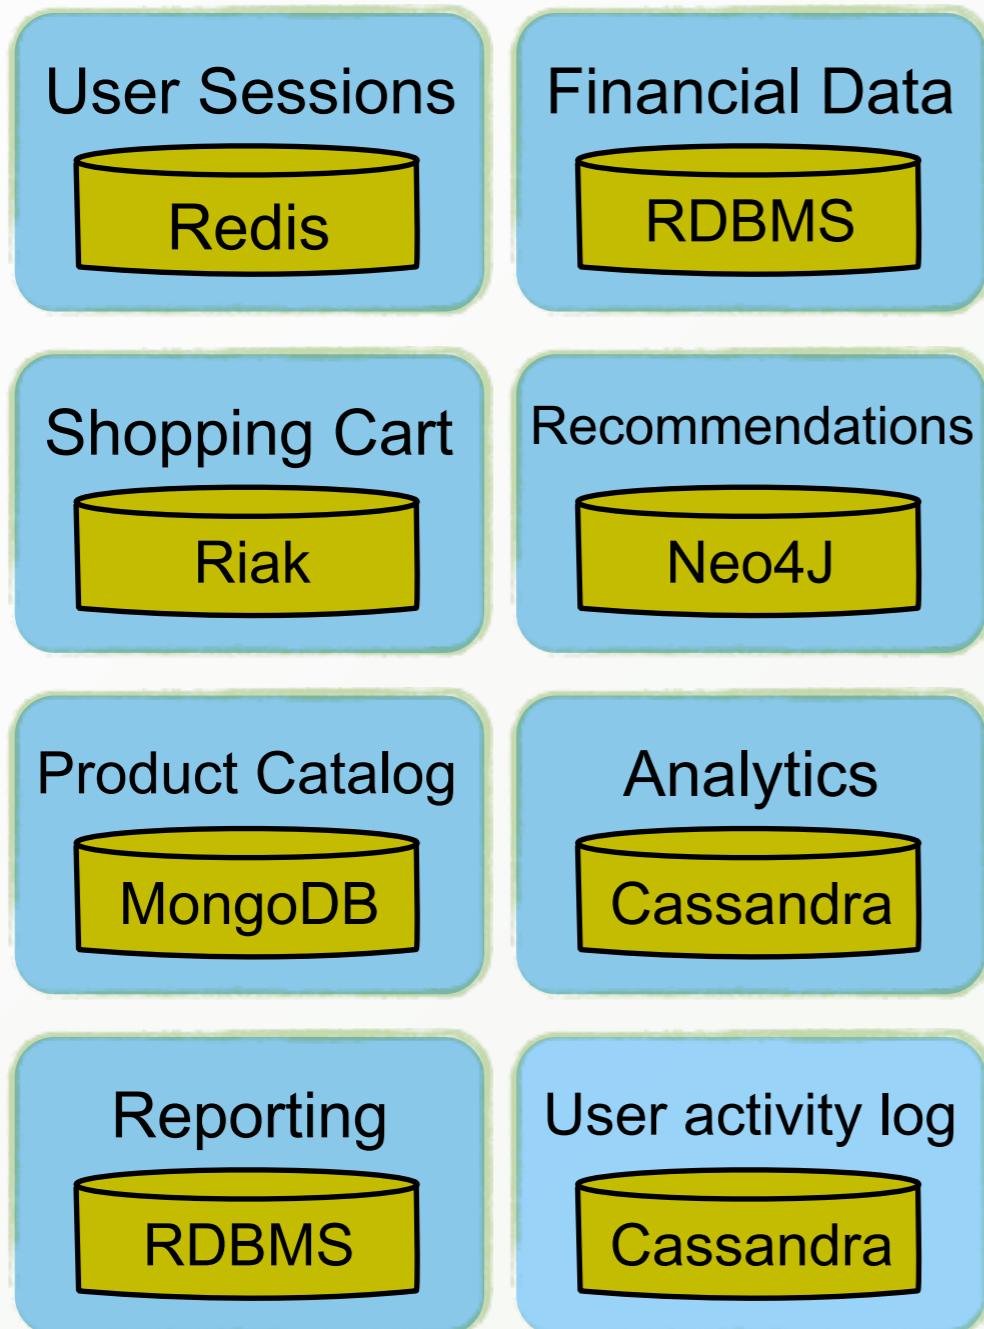


- ▶ Map value data to unique string keys (identifiers)
- ▶ Treat data as opaque (data has no schema)
- ▶ Can implement scaling and partitioning easily

e-commerce system in Relational World

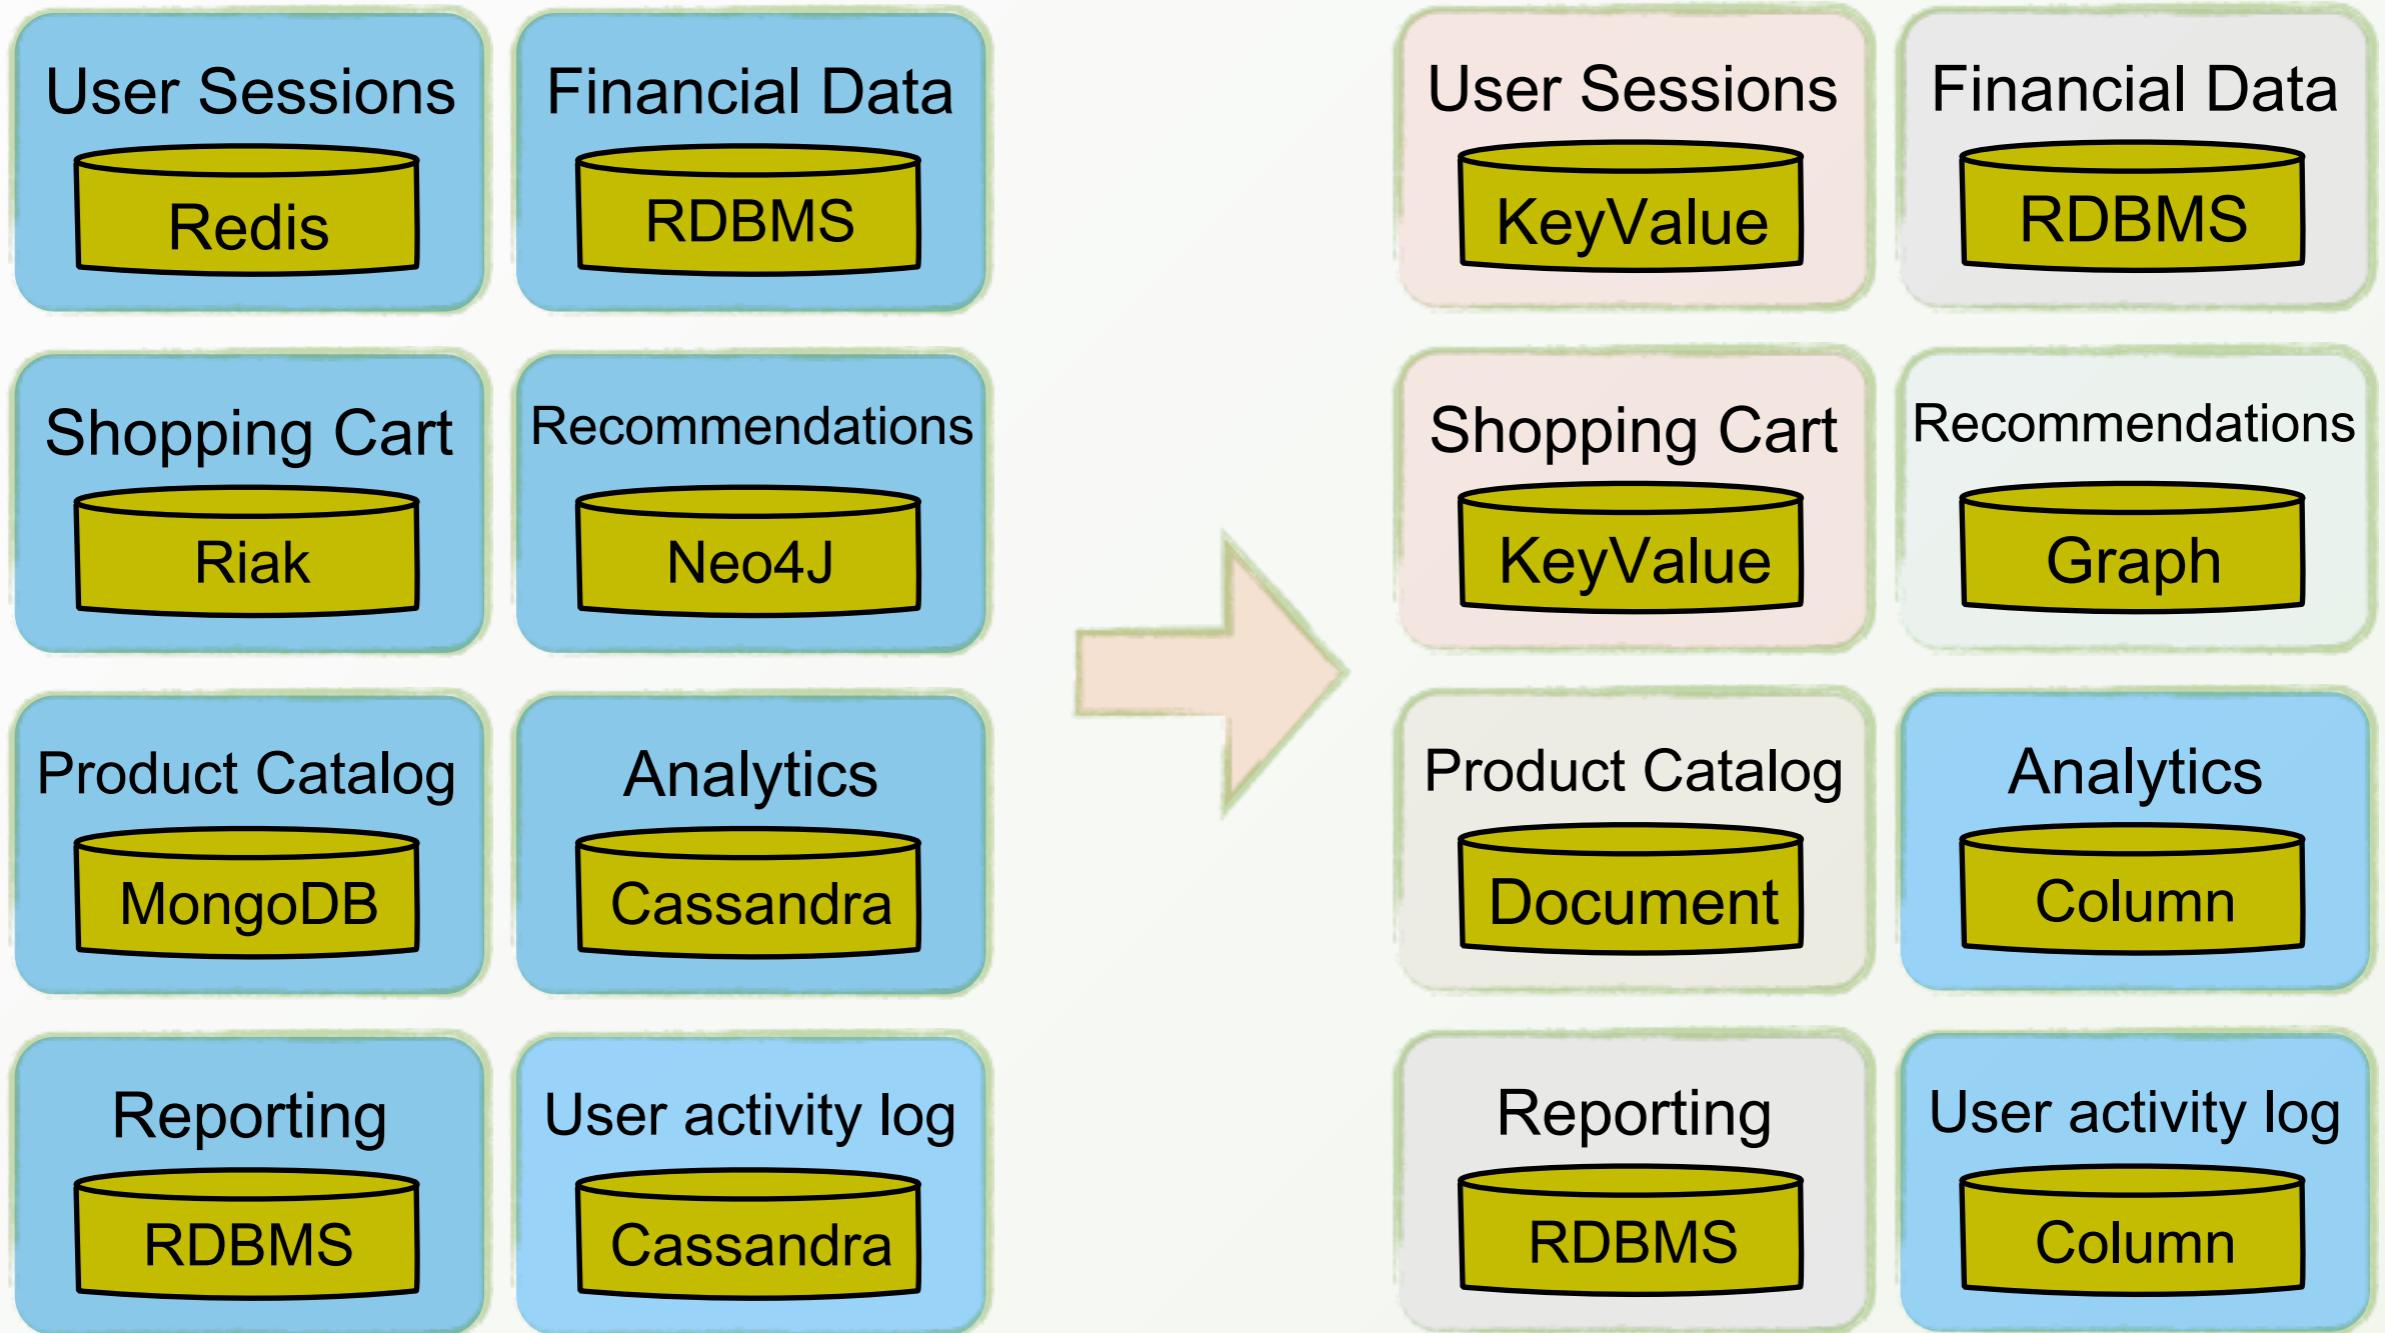


Polyglot Persistence



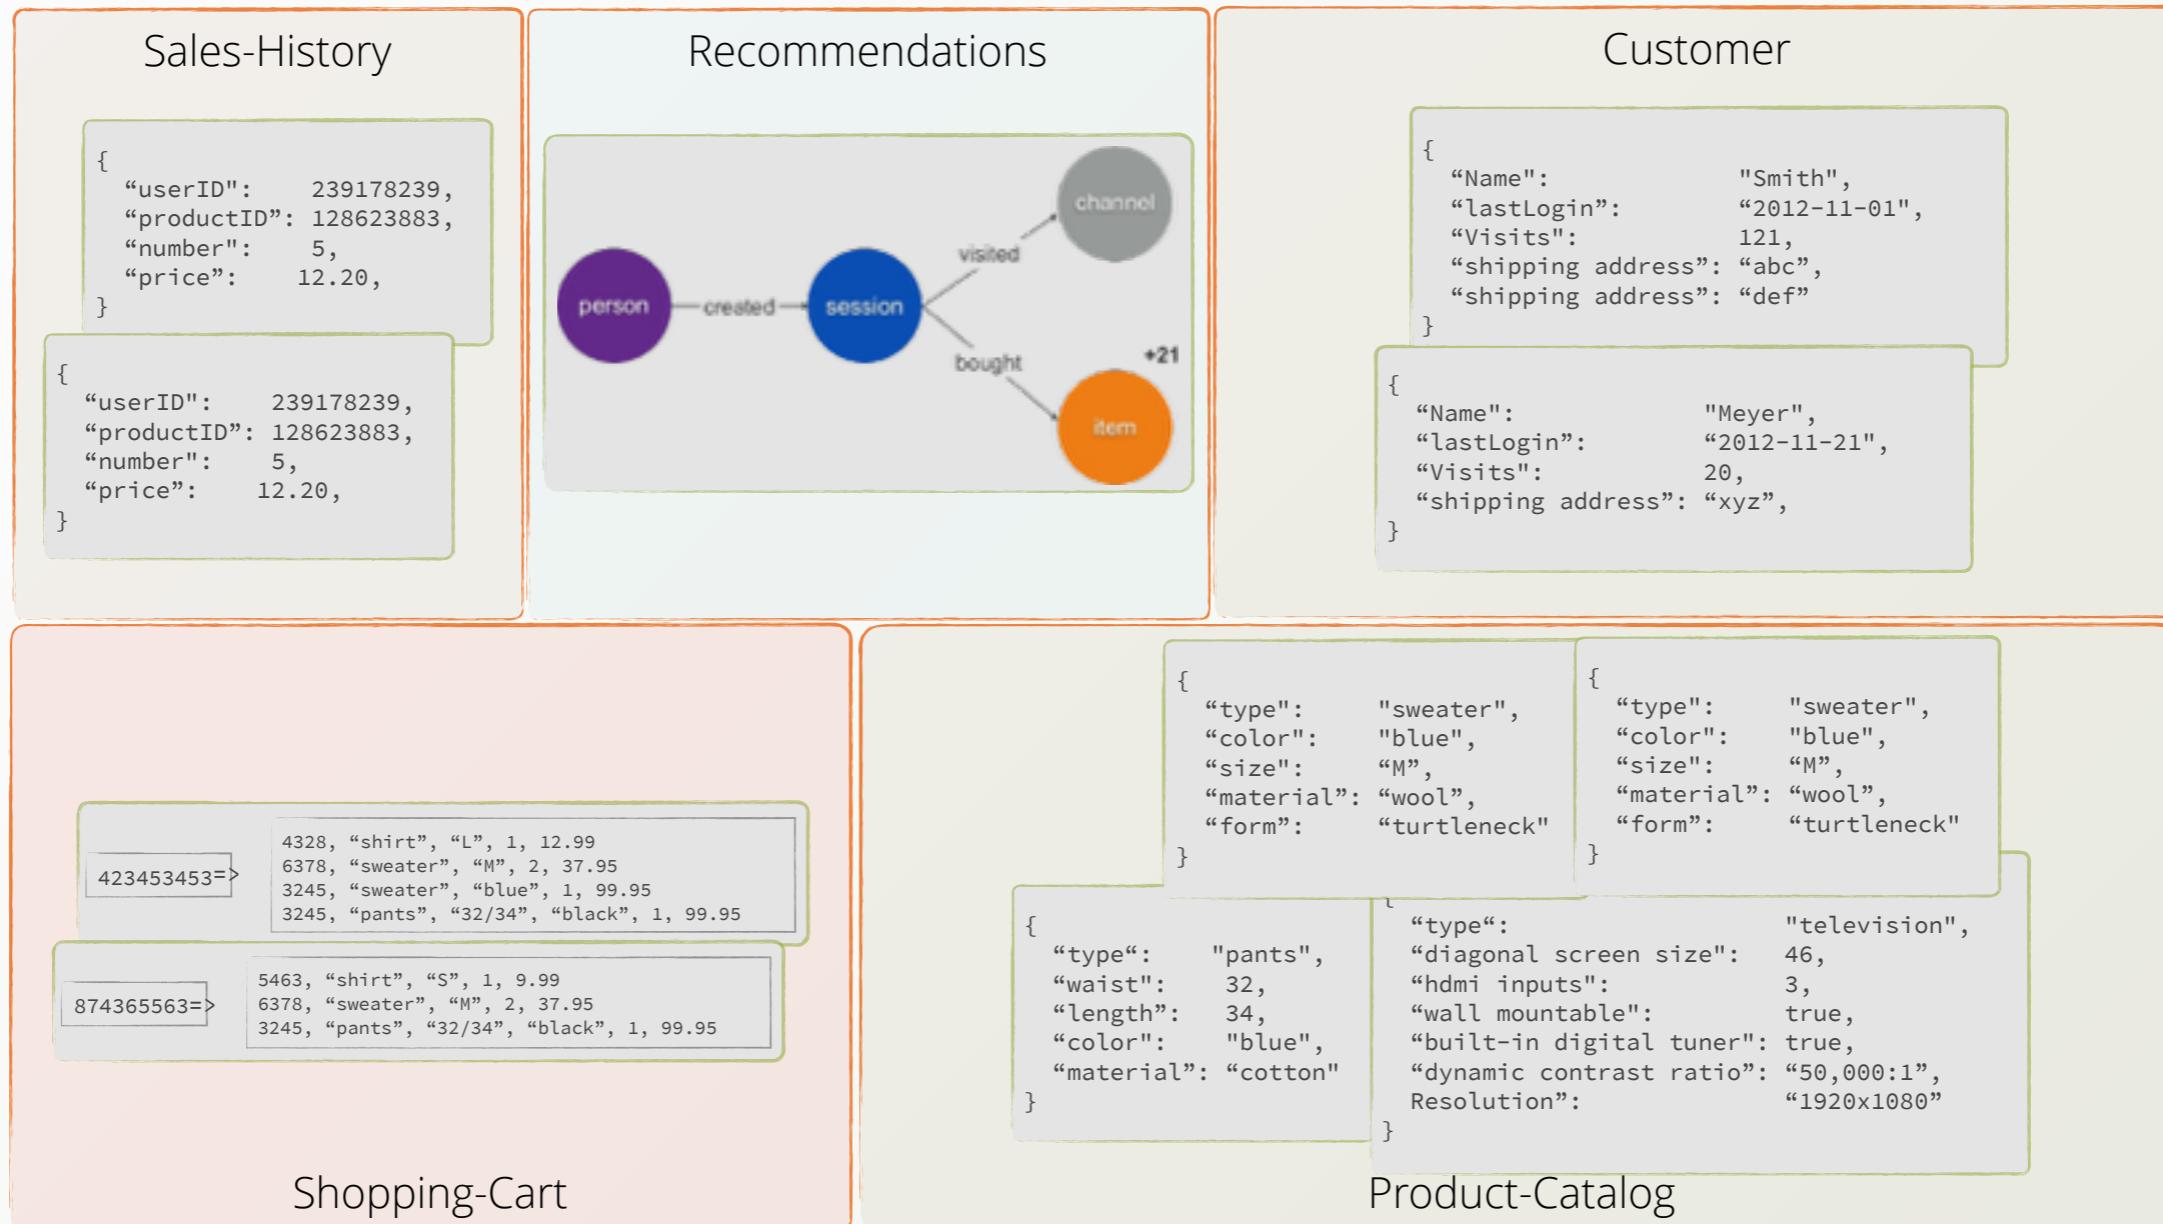
Source: Martin Fowler, <http://martinfowler.com/articles/nosql-intro.pdf>

Polyglot Persistence

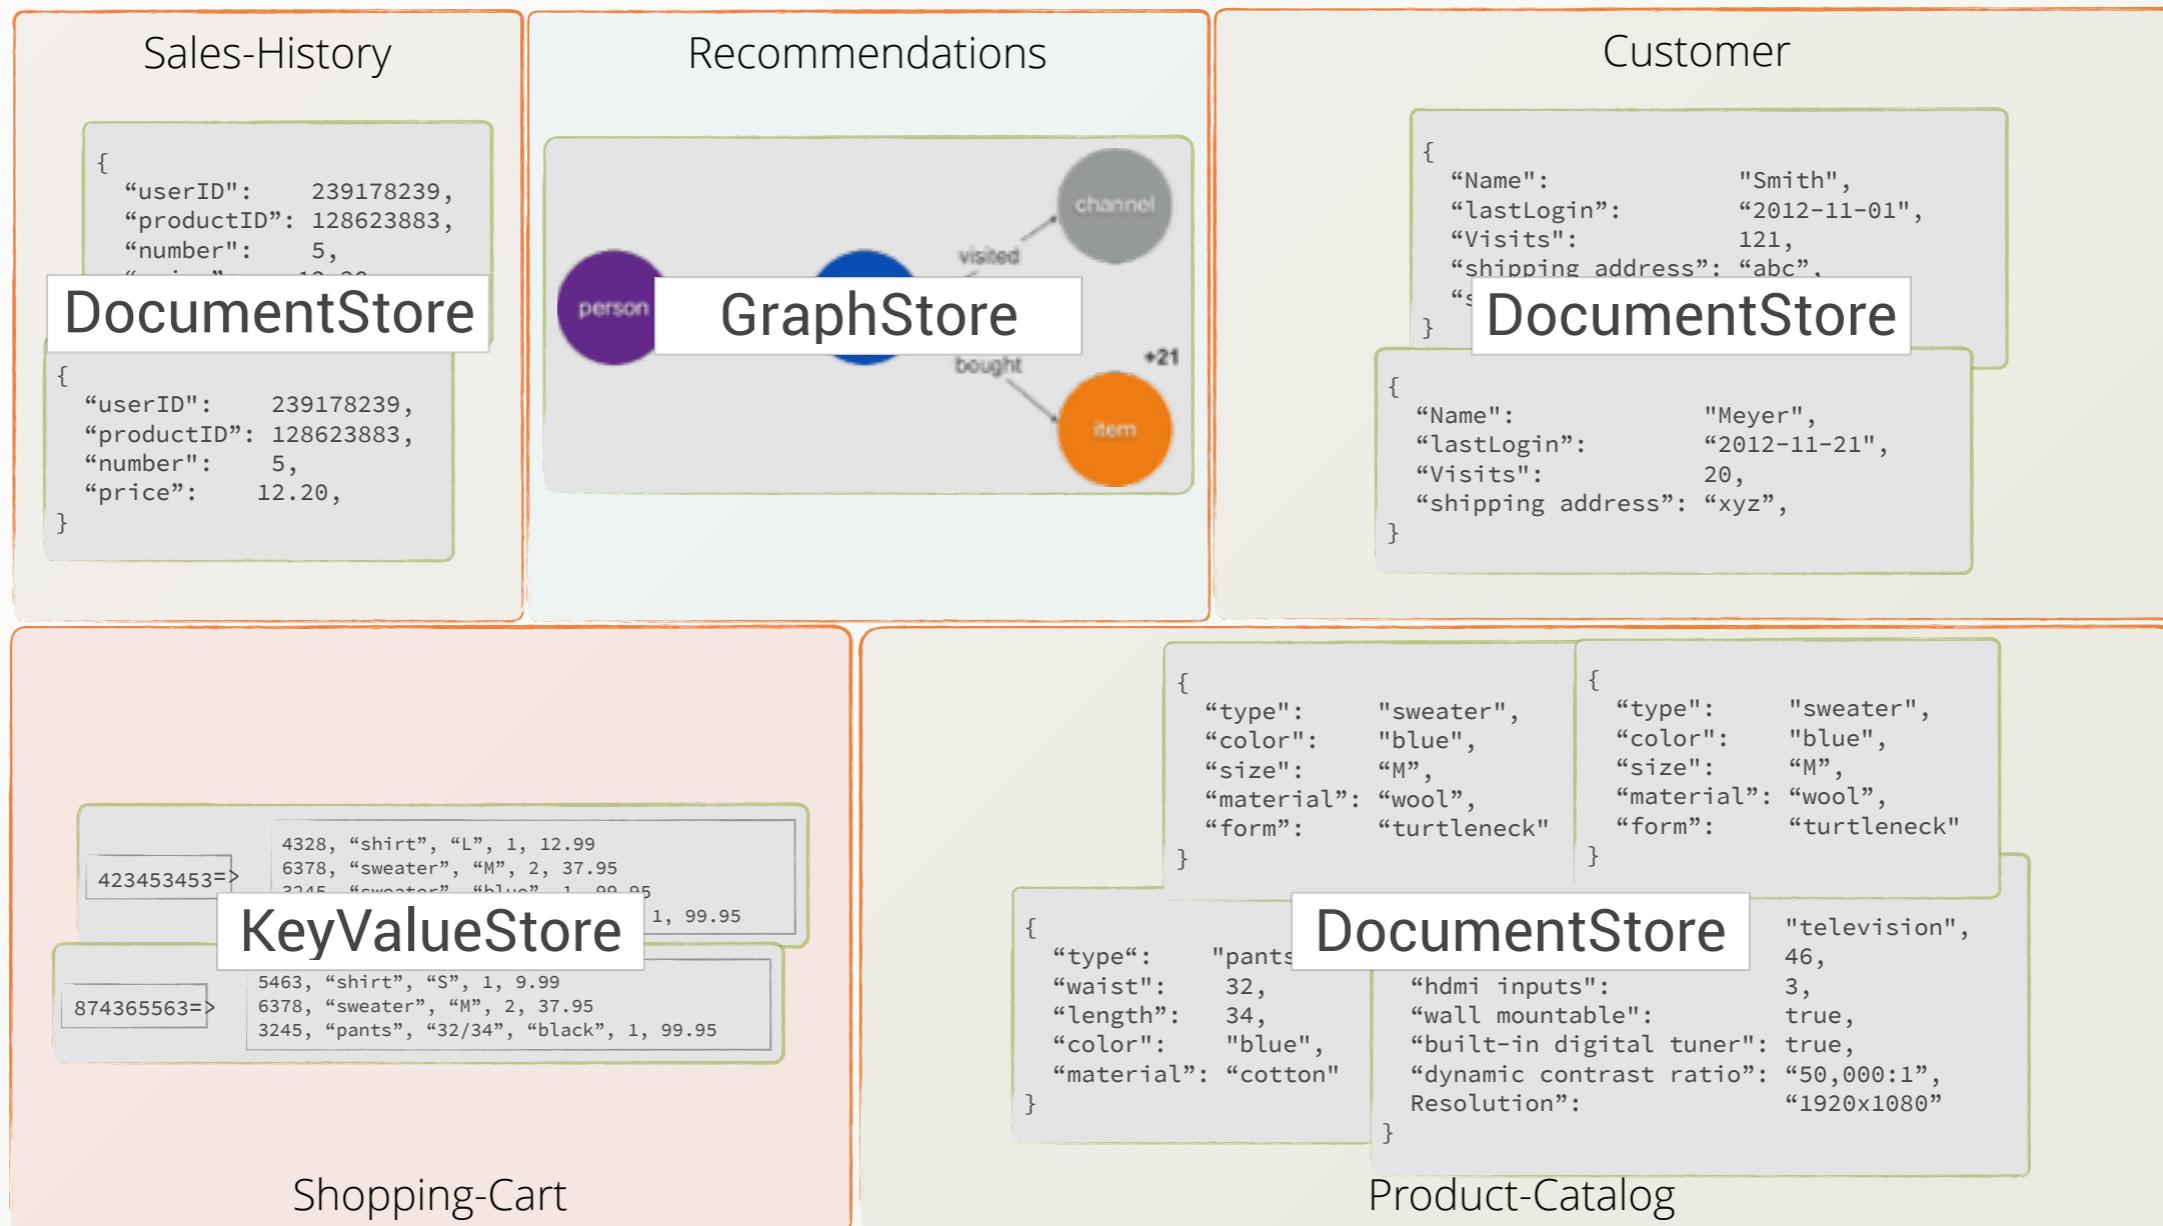


Source: Martin Fowler, <http://martinfowler.com/articles/nosql-intro.pdf>

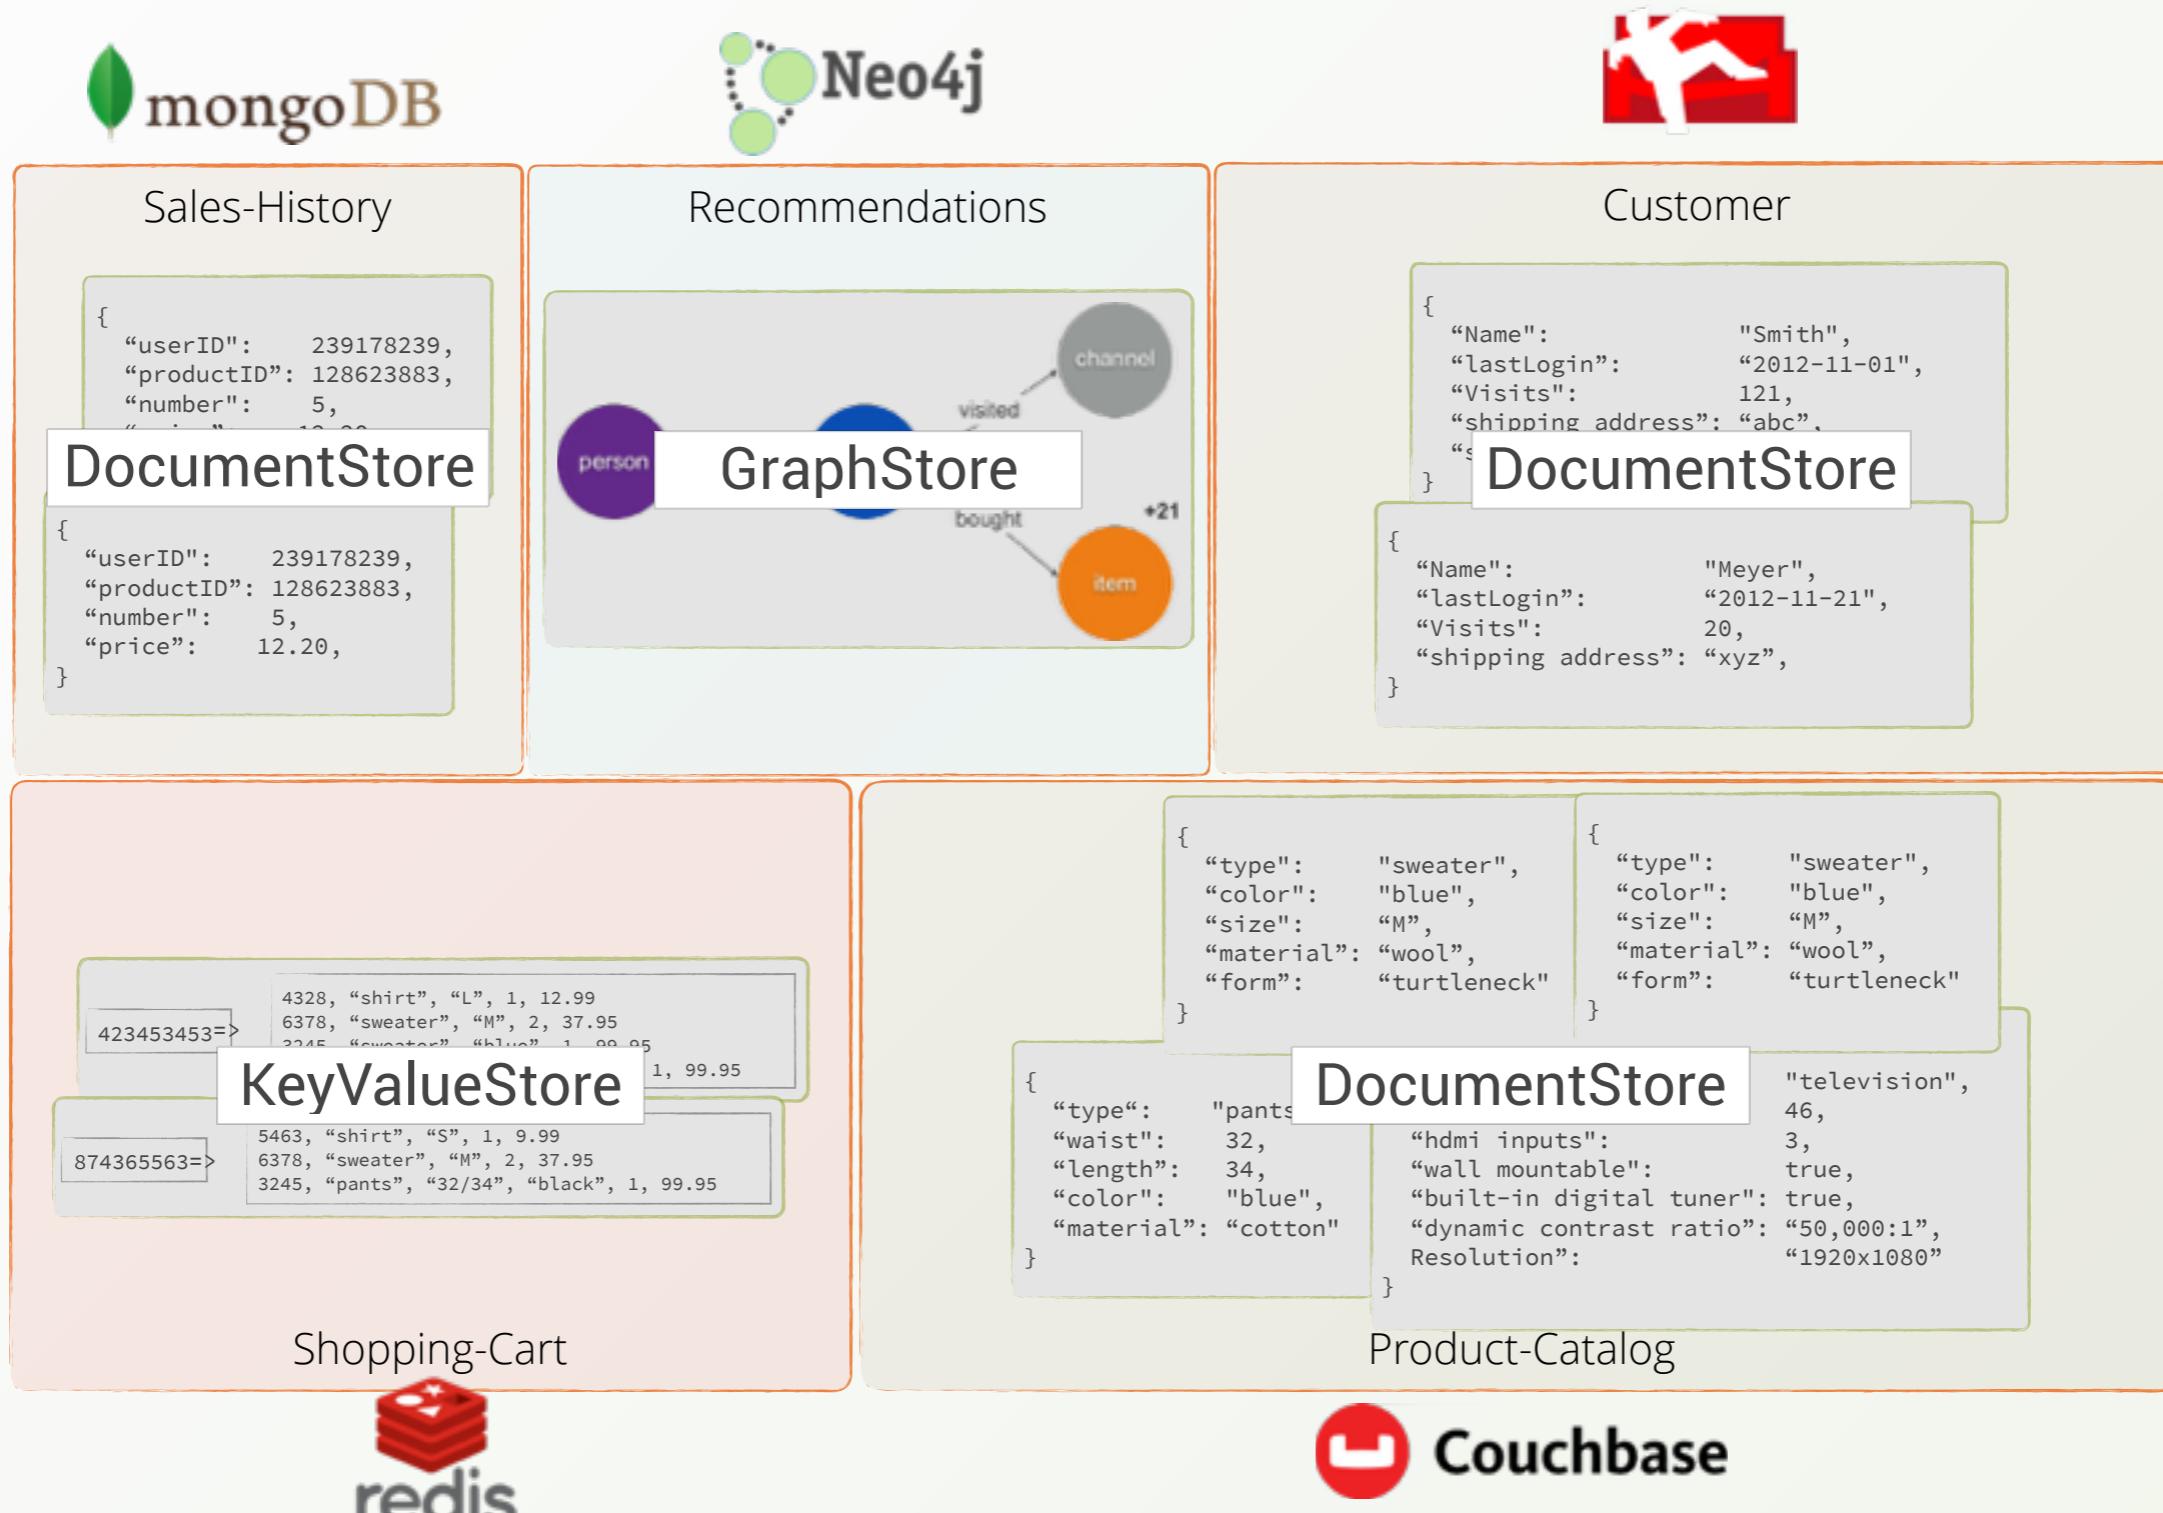
Single Model Databases



Single Model Databases



Single Model Databases



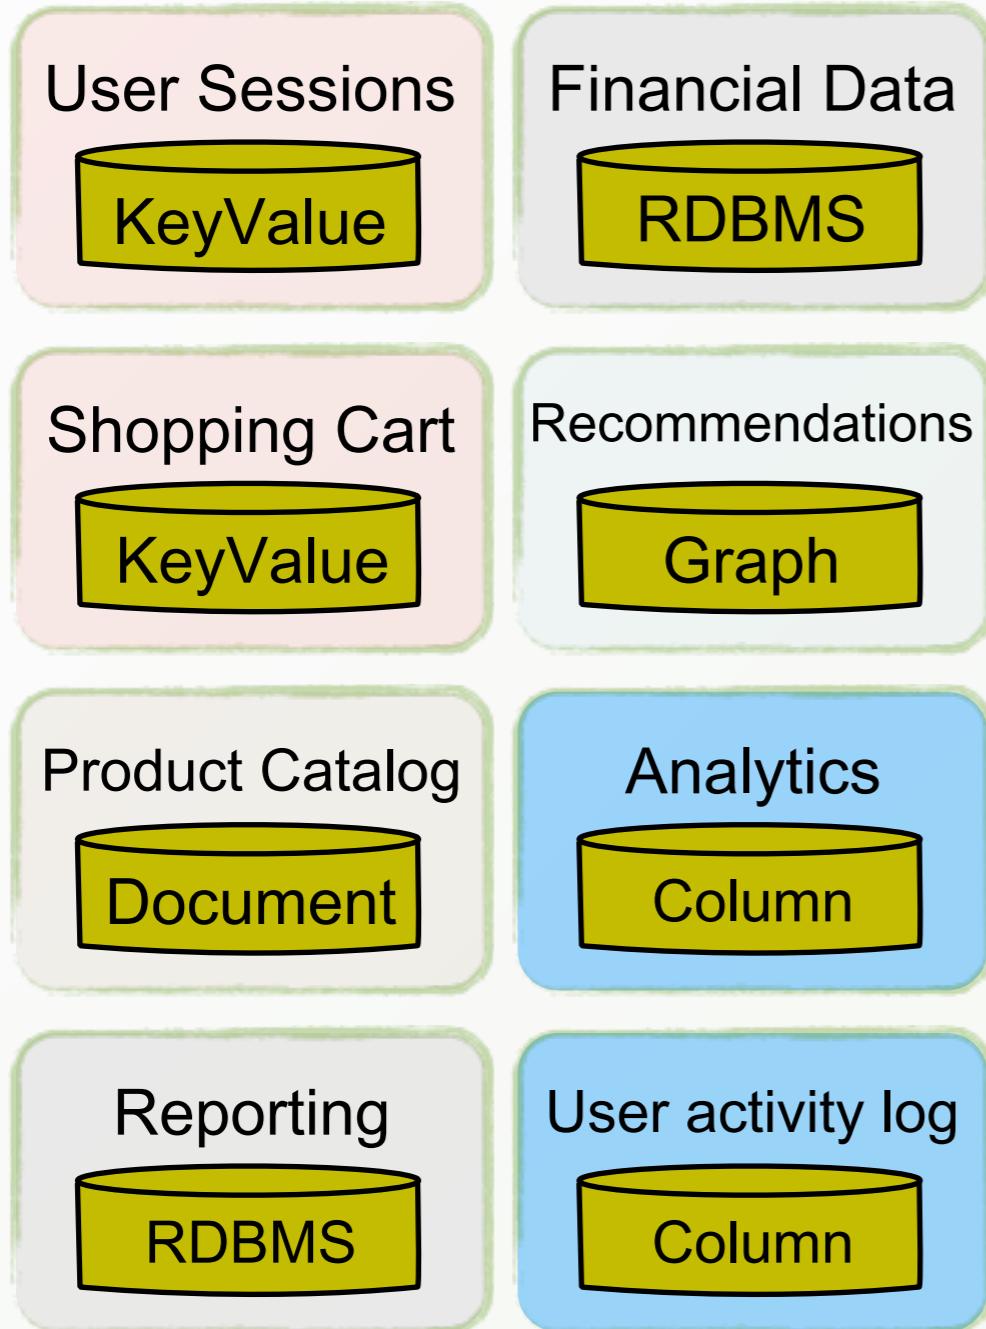
Benefits & Overhead

- ▶ Natural mapping of data into DB
- ▶ DB optimized for the data format
- ▶ Queries are tailored for your data format
- ▶ Focus on writing business logic
- ▶ Data has to be stored redundantly and has to be kept in sync
- ▶ Several technologies involved
- ▶ Administration effort is huge

Solution: Multi Model Database

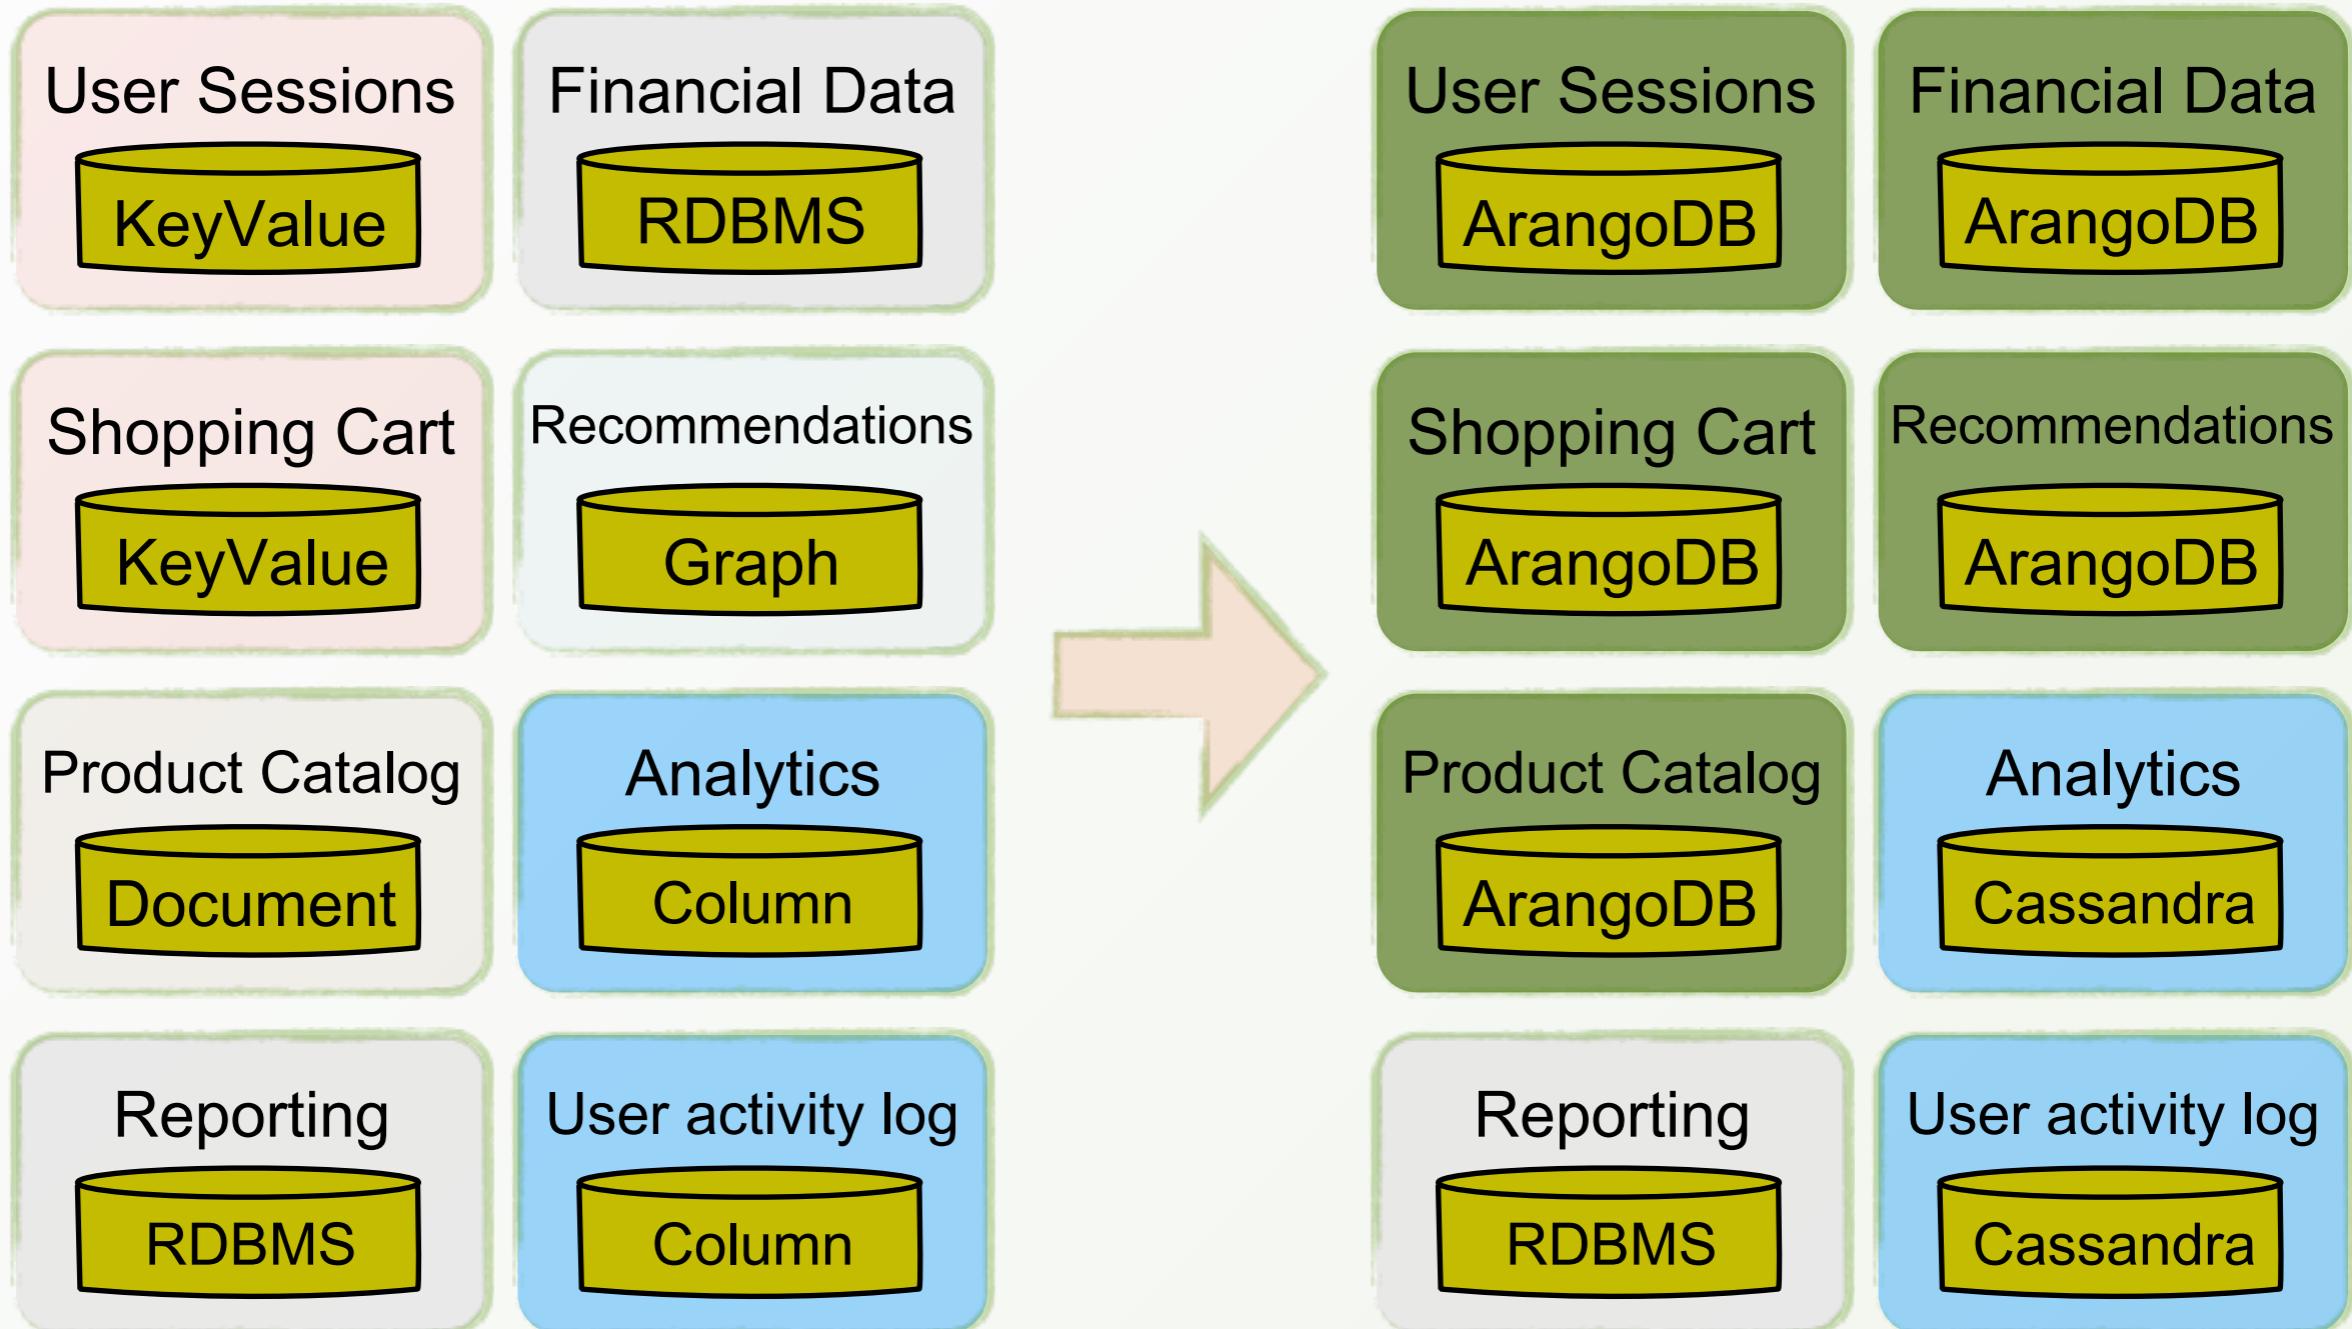
- ▶ Can natively store several kinds of data models:
 - ▶ Key-value pairs
 - ▶ Documents
 - ▶ Graphs
- ▶ Delivers query mechanisms for all data models

Polyglot Persistence Revisited



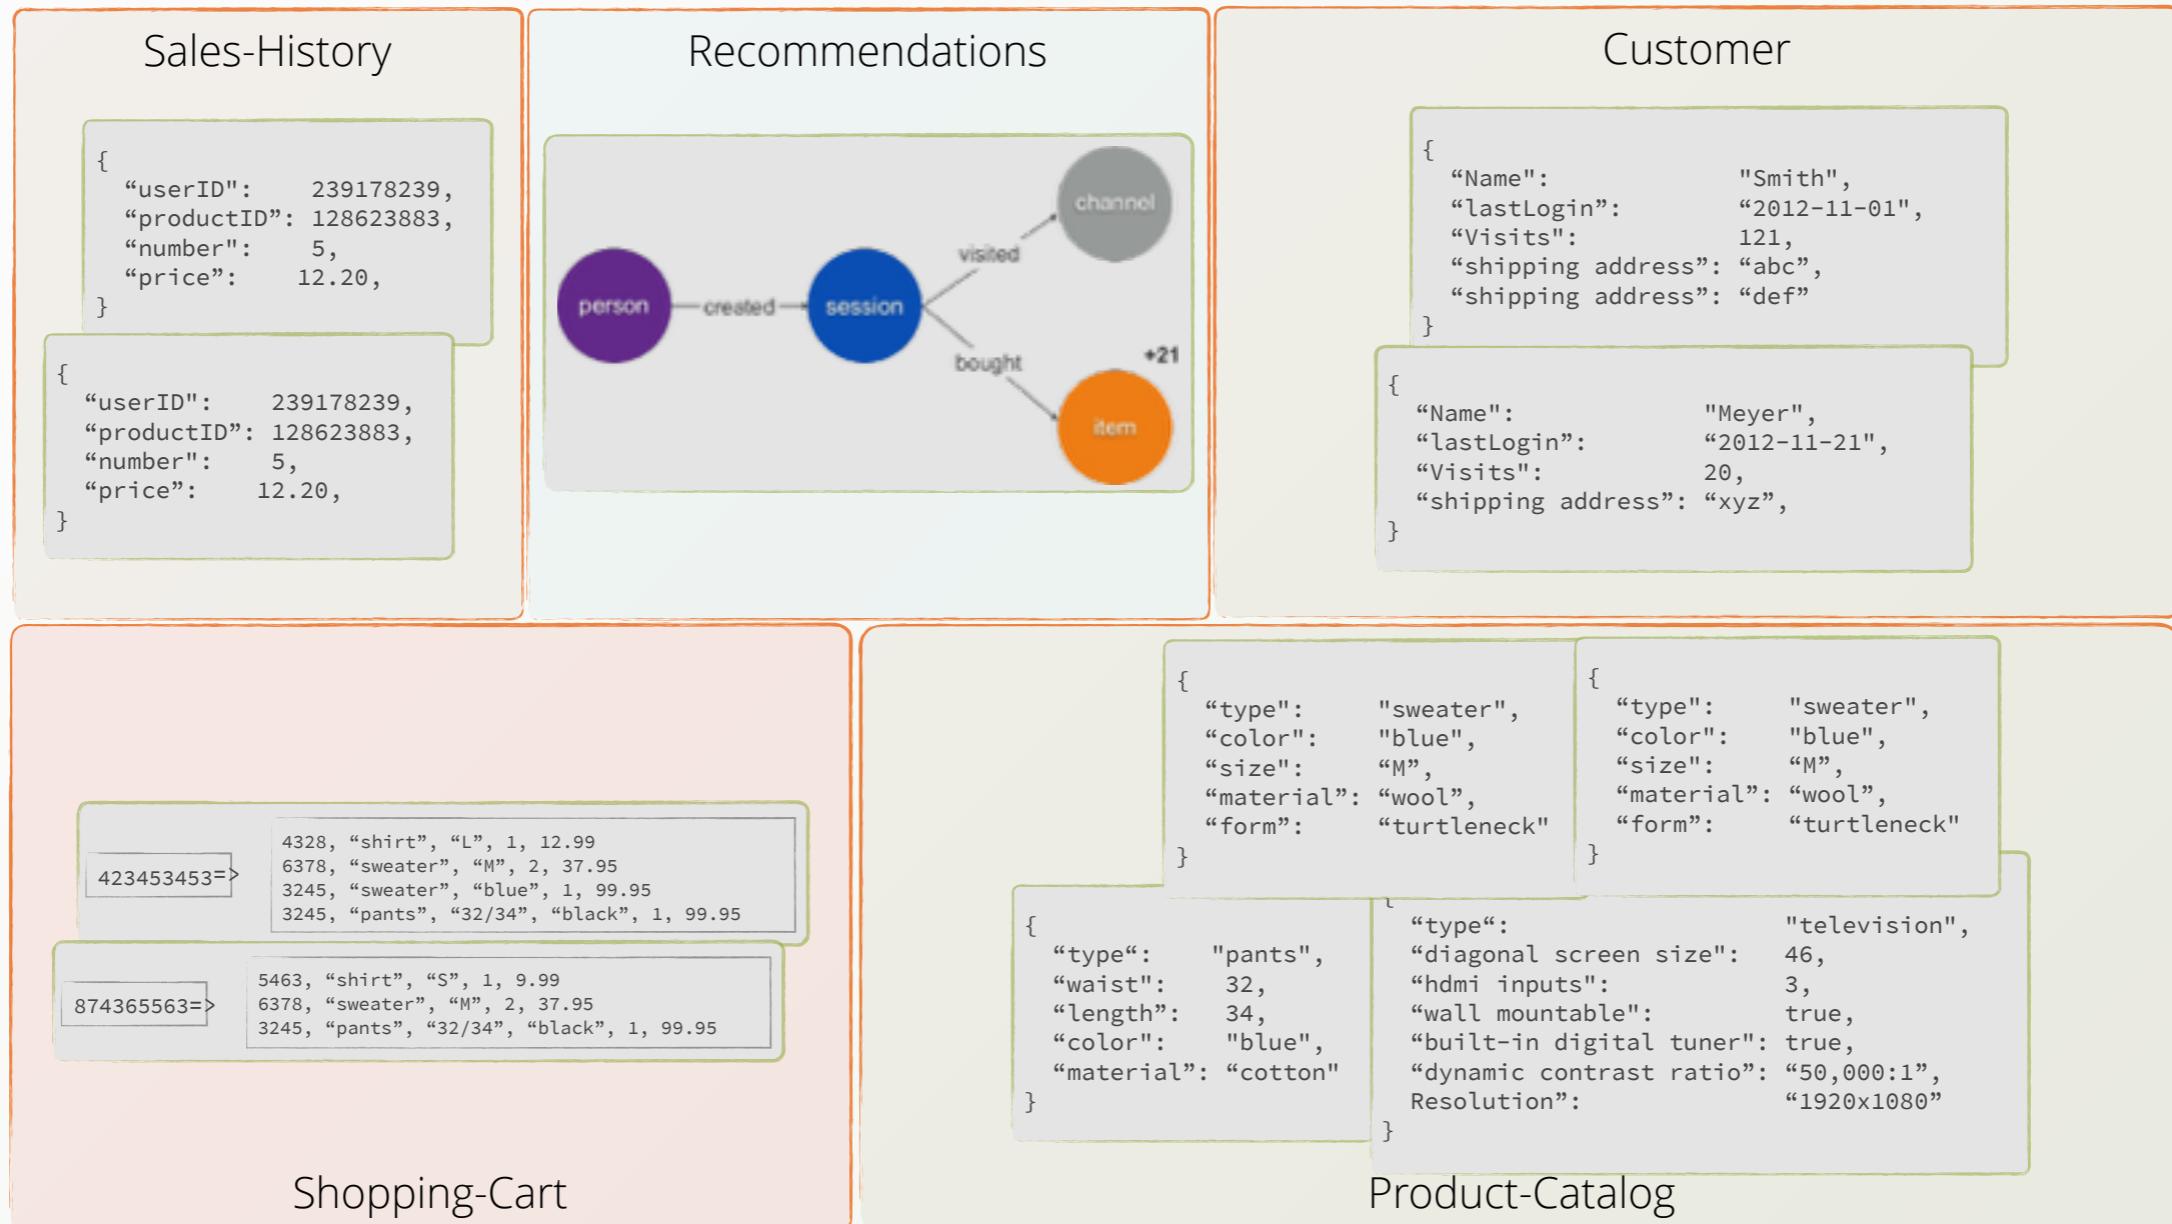
Source: Martin Fowler, <http://martinfowler.com/articles/nosql-intro.pdf>

Polyglot Persistence Revisited

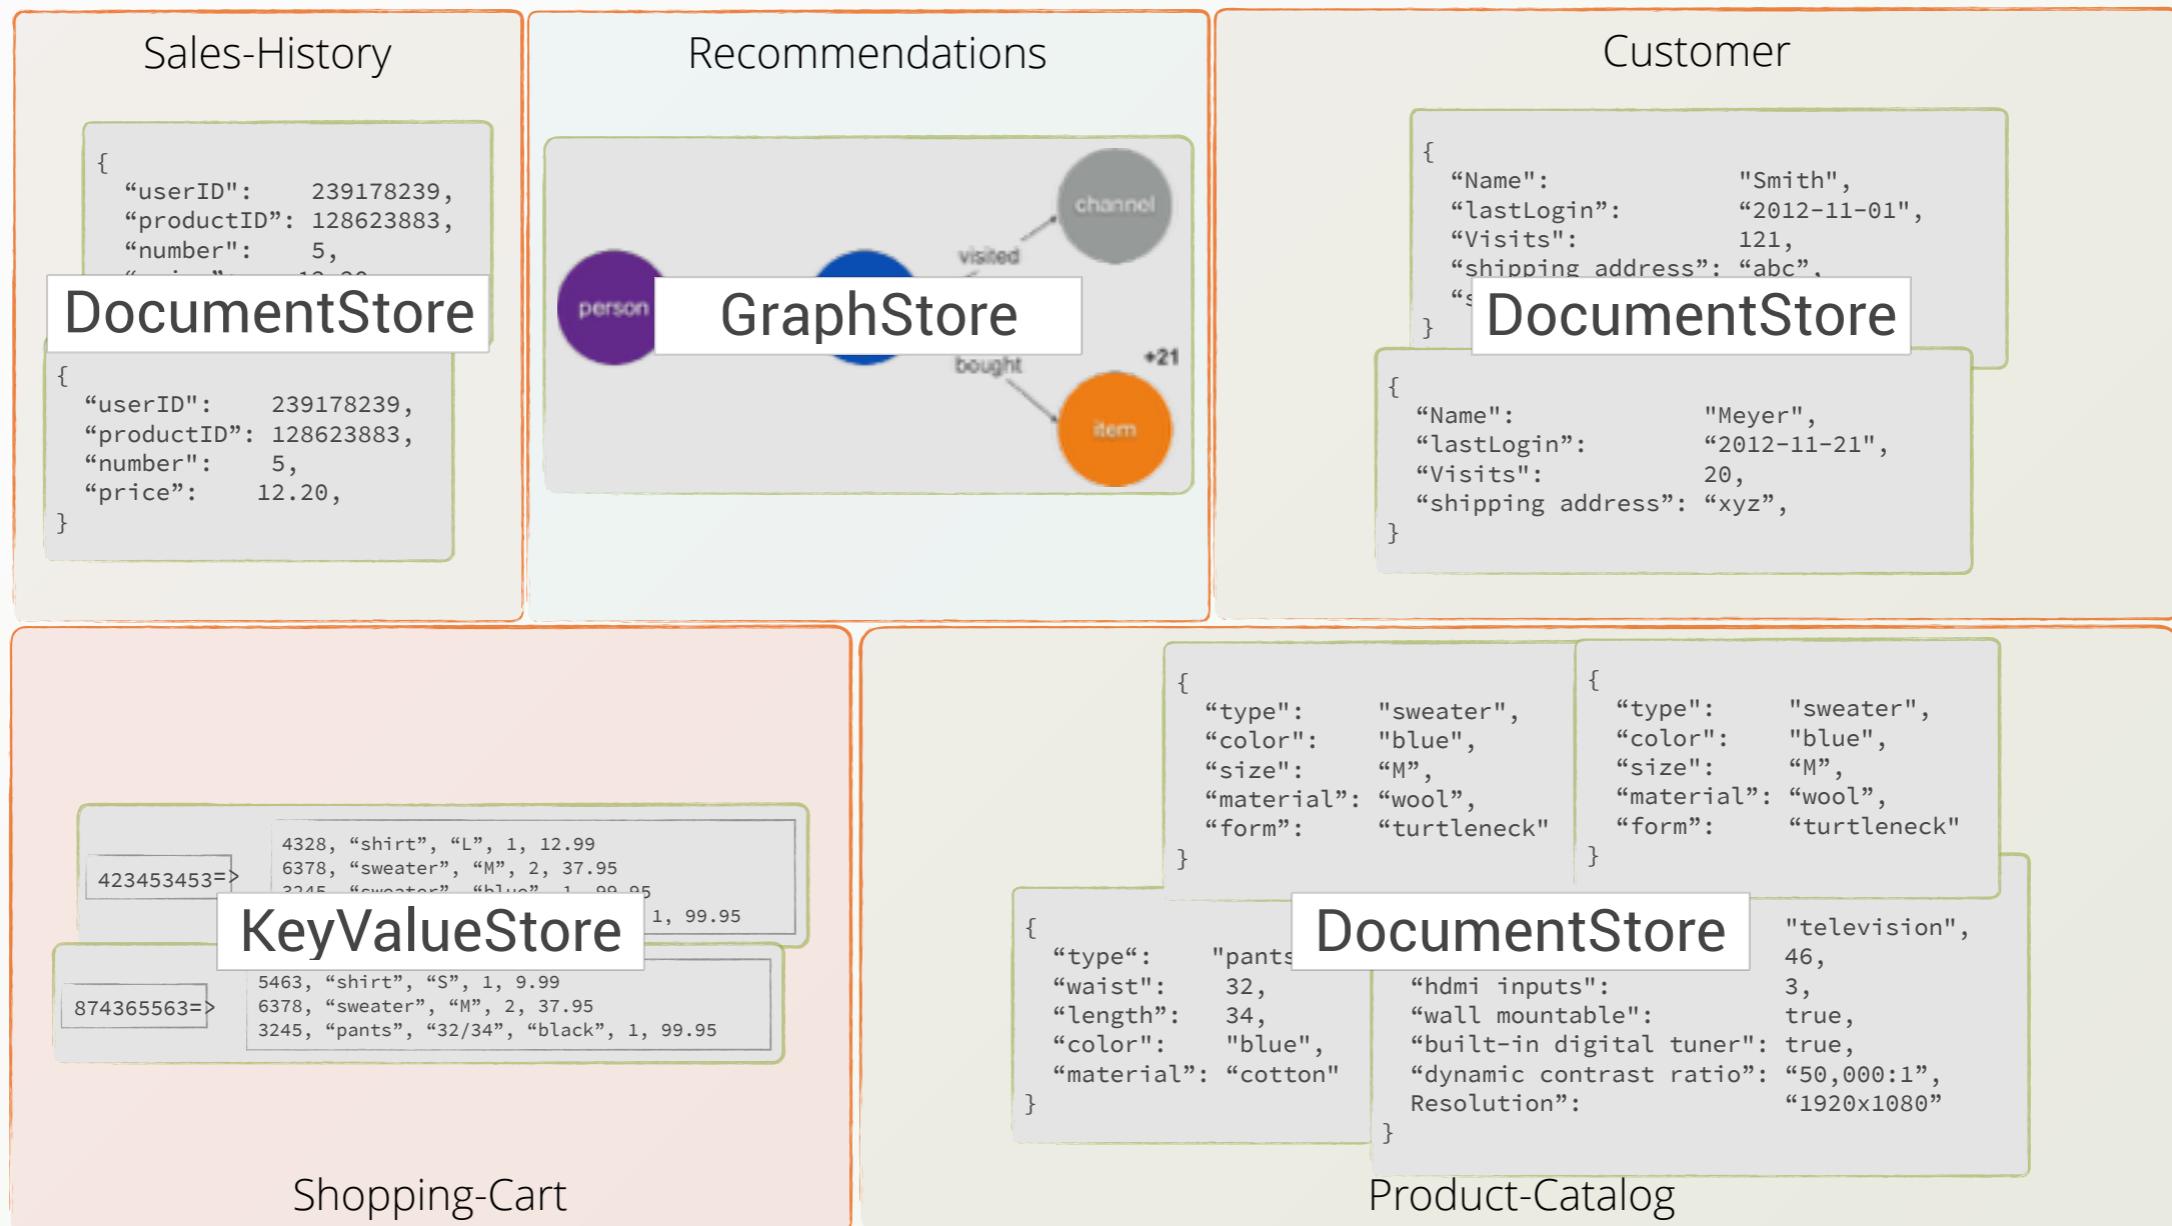


Source: Martin Fowler, <http://martinfowler.com/articles/nosql-intro.pdf>

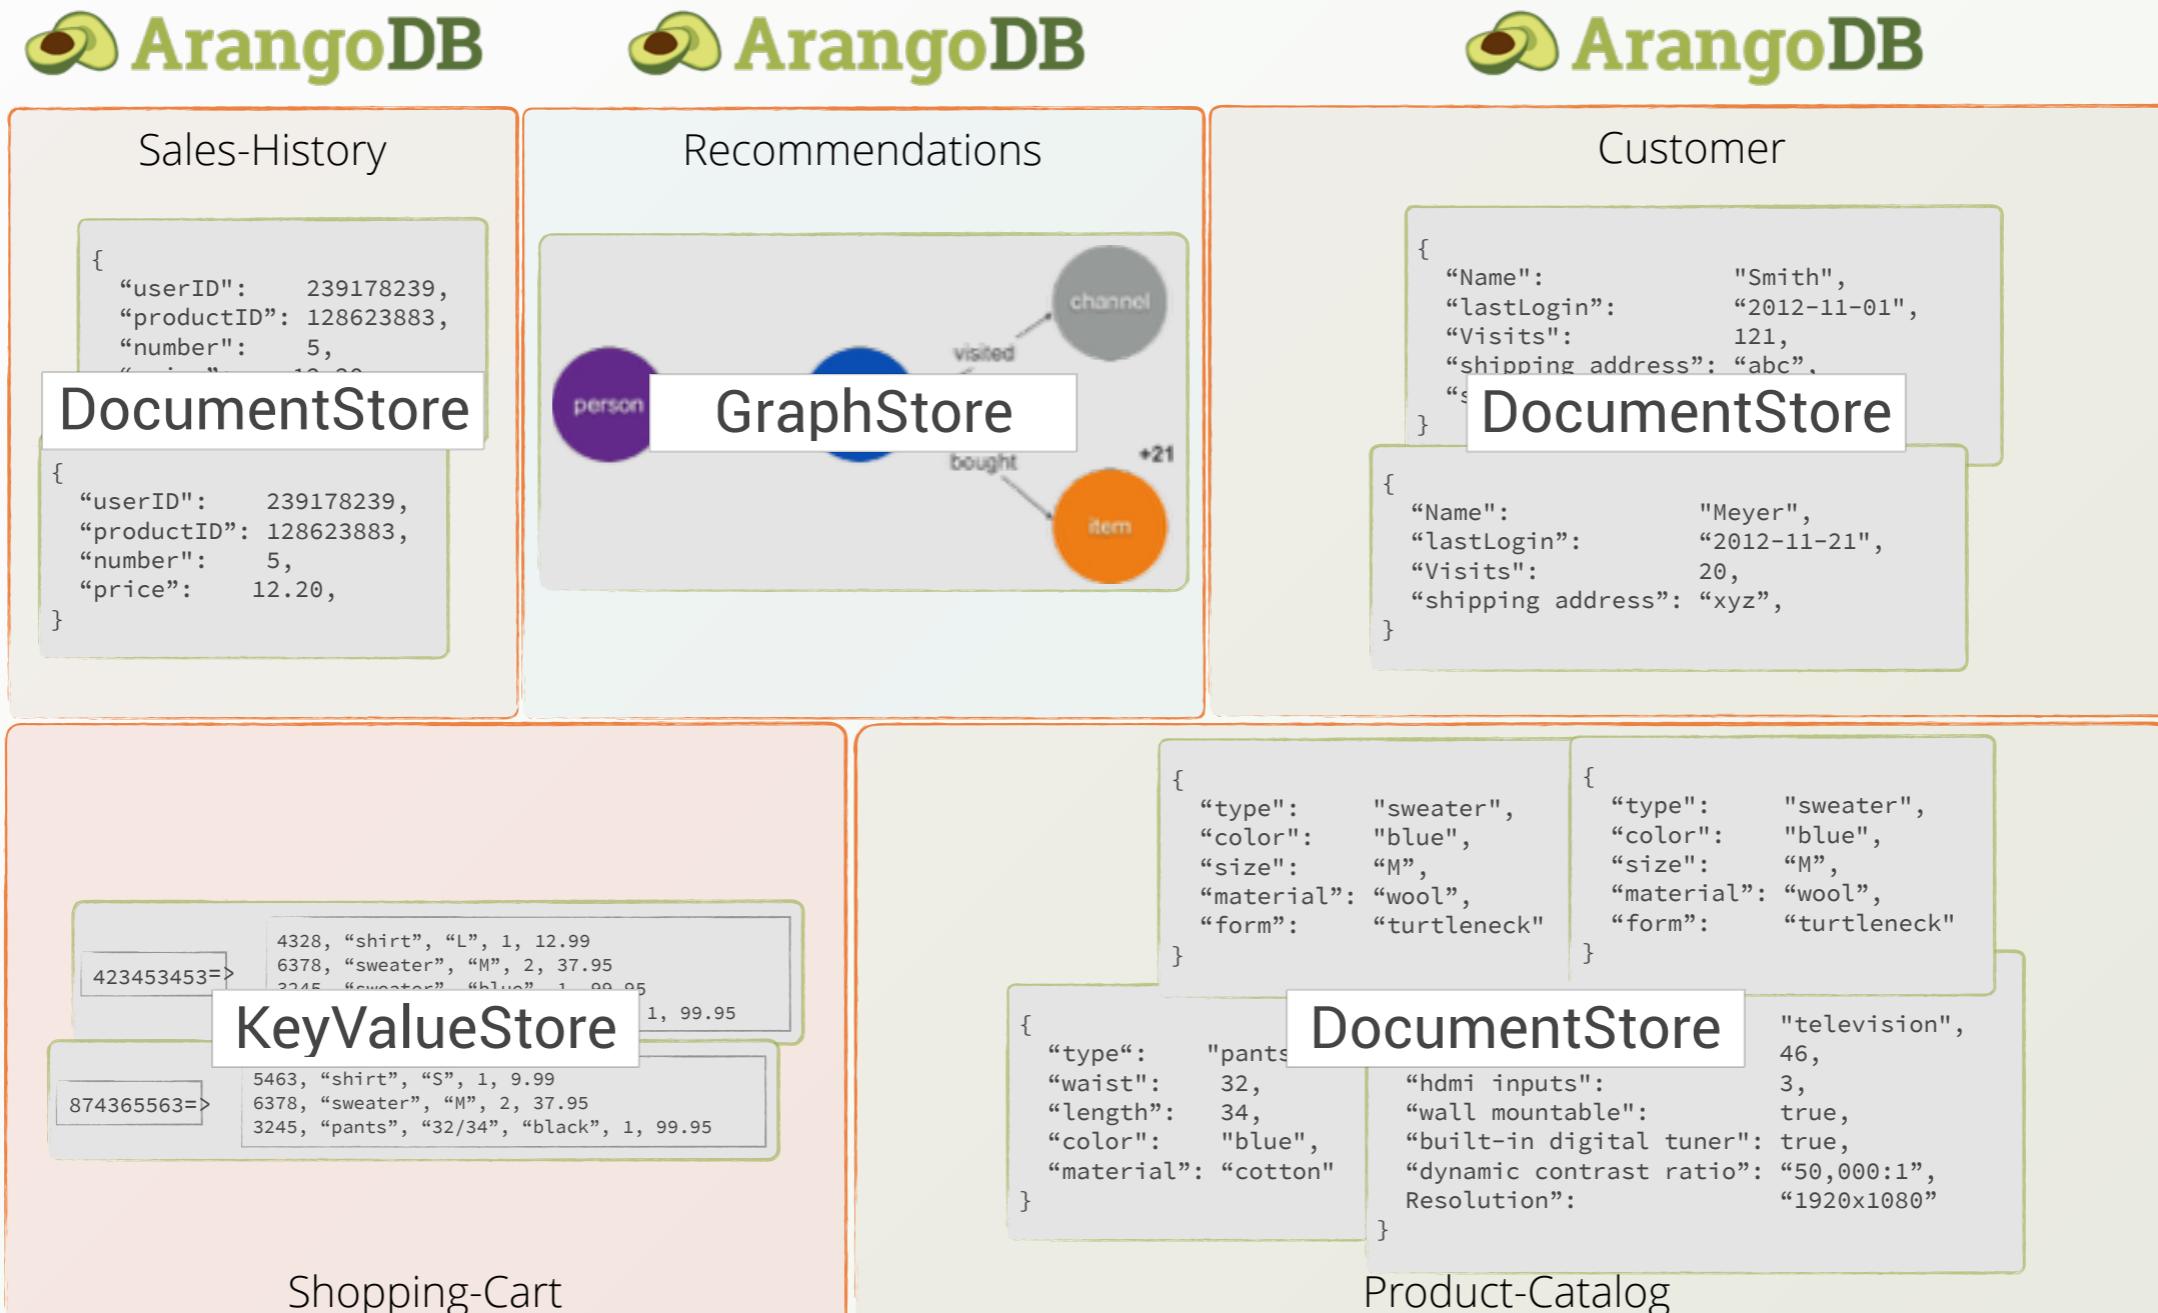
Use Case: Multi-Model-Databases



Use Case: Multi-Model-Databases



Use Case: Multi-Model-Databases



ArangoDB

ArangoDB

My four favorite features of ArangoDB

- ▶ AQL offering joins & traversals
- ▶ ACID including Multi Collection Transactions
- ▶ MULTI-MODEL stores graphs and documents

My four favorite features of ArangoDB

- ▶ AQL offering joins & traversals
- ▶ ACID including Multi Collection Transactions
- ▶ MULTI-MODEL stores graphs and documents
- ▶ FOXX extend the API and adapt it to your needs

AQL

```
FOR user IN users  
  RETURN user
```

AQL

FOR user IN users

 FILTER user.name == "alice"

 RETURN user

AQL

```
FOR user IN users
  FILTER user.name == "alice"
  FOR invoice IN invoices
    FILTER user._key == invoice.customer
    RETURN {
      user: user,
      invoice: invoice
    }
```

AQL

```
FOR user IN users
  FILTER user.name == "alice"
  RETURN {
    user: user,
    invoices: (
      FOR invoice IN invoices
        FILTER user._key == invoice.customer
        RETURN invoice
    )
  }
```

AQL

```
FOR user IN users
  FILTER user.name == "alice"
  RETURN {
    user: user,
    hasToPay: SUM(
      FOR invoice IN invoices
        FILTER user._key == invoice.customer
        FILTER !invoice.payed
        RETURN invoice.price
    )
  }
```

AQL - Alternative

FOR user IN users

 FILTER user.name == "alice"

FOR invoice IN invoices

 FILTER user._key == invoice.customer

 FILTER !invoice.payed

COLLECT u = user AGGREGATE toPay = SUM(invoice.price)

RETURN {

 user: u,

 hasToPay: toPay

}

AQL

FOR user IN users

 FILTER user.name == "alice"

 FOR product IN OUTBOUND user has_bought

 RETURN product

AQL

```
FOR user IN users
  FILTER user.name == "alice"
  FOR product IN OUTBOUND user has_bought
    LIMIT 10
    RETURN product
```

AQL

```
FOR user IN users
  FILTER user.name == "alice"
  FOR product, action IN OUTBOUND user has_bought
    SORT action.timestamp DESC
    LIMIT 10
    RETURN product
```

AQL

```
FOR user IN users
```

```
  FILTER user.name == "alice"
```

```
  FOR recommendation IN 3 ANY user has_bought
```

```
    OPTIONS { bfs: true, uniqueVertices: "global" }
```

```
  LIMIT 10
```

```
  RETURN recommendation
```

AQL

```
FOR user IN users
  FILTER user.name == "alice"
  FOR recommendation, action, path IN 3 ANY user has_bought
    FILTER path.vertices[2].age <= user.age + 5
    AND path.vertices[2].age >= user.age - 5
  LIMIT 10
  RETURN recommendation
```

AQL

```
FOR user IN users
  FILTER user.name == "alice"
  FOR recommendation, action, path IN 3 ANY user has_bought
    FILTER path.vertices[2].age <= user.age + 5
    AND path.vertices[2].age >= user.age - 5
    FILTER recommendation.price < 25
  LIMIT 10
  RETURN recommendation
```

ACID - Transactions

- ▶ Invoke a transaction:

```
db._executeTransaction({  
  collections: {  
    write: ["users", "products"],  
    read: "has_bought"  
  },  
  action: function() {  
    // all operations go here  
  
  }  
});
```

ACID - Transactions

► Invoke a transaction:

```
db._executeTransaction({  
  collections: {  
    write: ["users", "products"],  
    read: "has_bought"  
  },  
  action: function() {  
    // all operations go here  
    throw "failure"; // Triggers rollback  
  }  
});
```

Benefits & Overhead

- ▶ Native mapping of data into DB
- ▶ DB optimized
- ▶ Queries are tailored for your data format
- ▶ Focus on writing business logic
- ▶ Data has to be stored redundantly and has to be kept in sync
- ▶ Several technologies involved
- ▶ Administration effort is huge

Benefits & Overhead

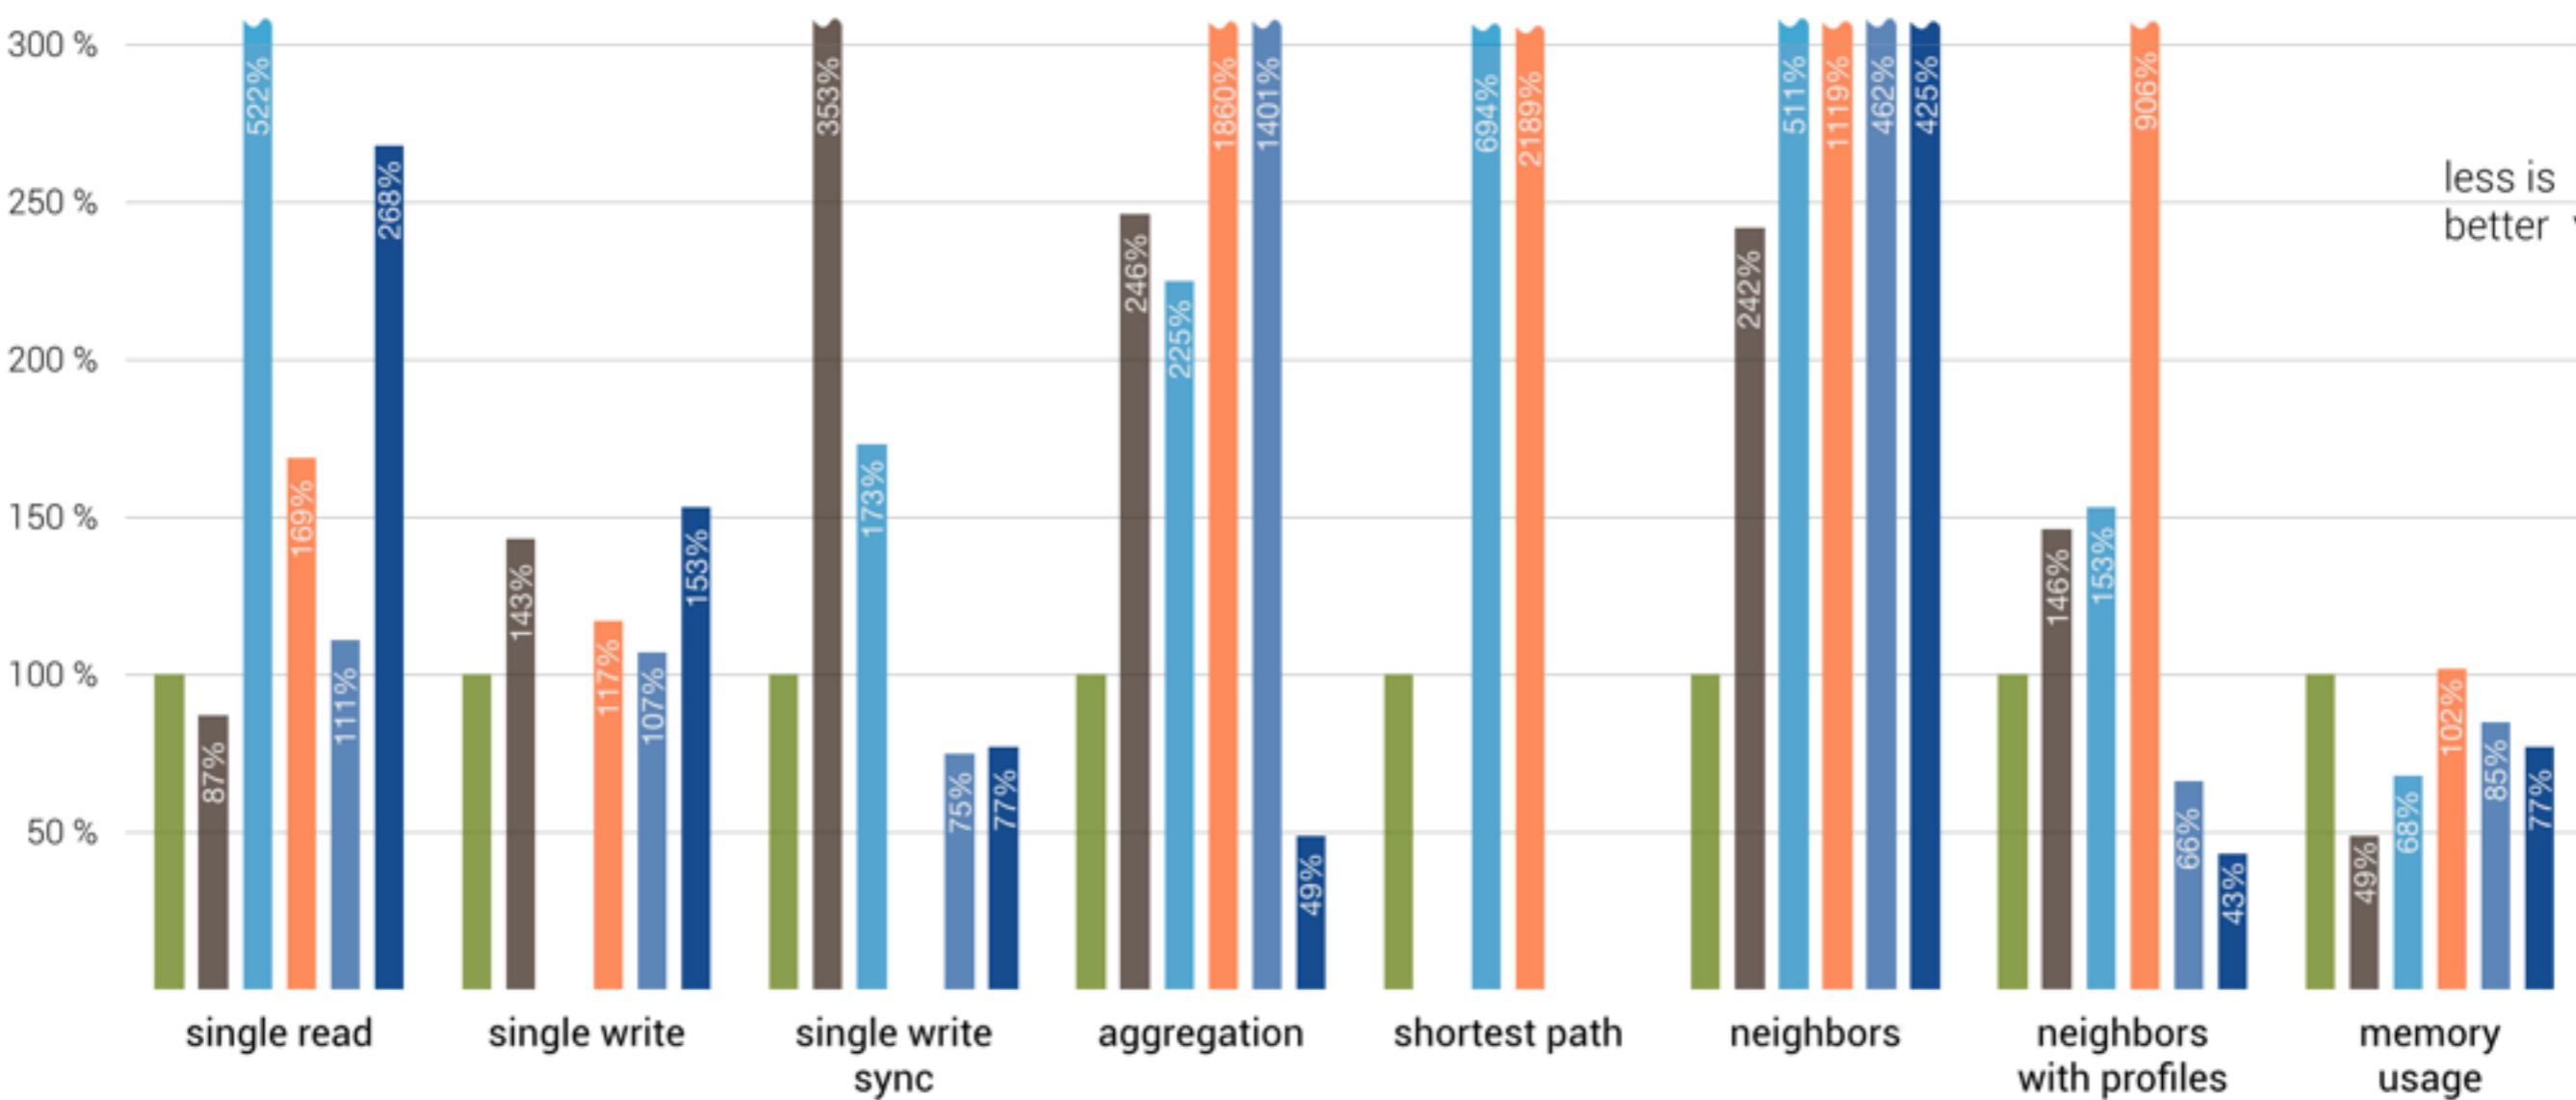
- ▶ Native mapping of data into DB
- ▶ DB optimized
- ▶ Queries are tailored for your data format
- ▶ Focus on writing business logic
- ▶ One technology involved

Benchmark Comparison

NoSQL Performance Test

ArangoDB, Postgres, MongoDB, Neo4j and OrientDB

ArangoDB Neo4j Postgres (json)
MongoDB OrientDB Postgres (tab)



*) neighbors and neighbors of neighbors (distinct)

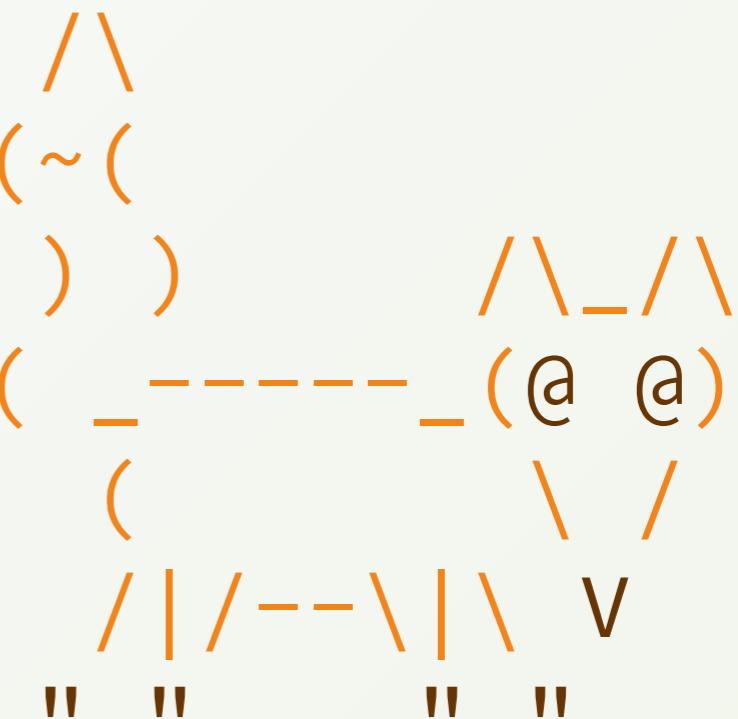
Database versions: ArangoDB 2.7 RC2, OrientDB 2.2 alpha, MongoDB 3.0.6, Neo4J 2.3 M3, PostgreSQL 9.4.4

Weinberger 2015-10-13 (r207)

Source: <https://www.arangodb.com/2015/10/benchmark-postgresql-mongodb-arangodb/>

FoXX

- ▶ Add your own customized and versioned REST-API on top of ArangoDB in JavaScript
 - ▶ Include as a microservice in Rails, Node.js etc.
 - ▶ Ship an administration fronted with it
- ▶ Built-in authentication using OAuth2.0 or HTTP-Basic Auth
- ▶ Operations are encapsulated in the database
 - ▶ low network traffic, direct data access
 - ▶ increases data privacy
- ➔ Multi-device setups
- ➔ Microservice Architectures



Foxx Example

```
router.get("/recommend/:price", function(req, res) {
  res.json(db._query(
    `FOR user IN users
      FILTER user.name == @name
      FOR recommendation, action, path IN 3 ANY user has_bought
        OPTIONS { bfs: true, uniqueVertices: "global" }
        FILTER path.vertices[2].age <= user.age + 5
        AND path.vertices[2].age >= user.age - 5
        FILTER recommendation.price < @maxPrice
        LIMIT 10
        RETURN recommendation`,
    {
      maxPrice: req.params("price"),
      name: req.session.username
    }).toArray());
});
```

Does it Scale?

- ▶ Sharding (Huge Dataset, Write-Scaling)
 - ▶ Collection: distributed across several servers
 - ▶ Distributed by: List of Attributes (default: _key)
 - ▶ Number of Shards immutable
 - ▶ Has to be defined on creation
 - ▶ Can be different for all collections
 - ▶ Suggested: #Servers²

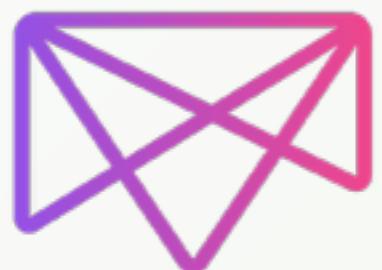
Does it Scale?

- ▶ Sharding (Huge Dataset, Write-Scaling)
 - ▶ Collection: distributed across several servers
 - ▶ Distributed by: List of Attributes (default: _key)
 - ▶ Number of Shards immutable
 - ▶ Has to be defined on creation
 - ▶ Can be different for all collections
 - ▶ Suggested: #Servers²
- ▶ Replication (Failover, Read-Scaling)
 - ▶ Shard: can have **n** Followers (Replicas)
 - ▶ Followers:
 - ▶ do not accept writes
 - ▶ synchronous (identical data)
 - ▶ placed on different machine (if possible)

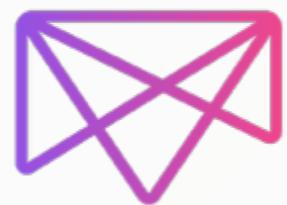
Speaking of Clusters ...



ArangoDB



MESOSPHERE



MESOSPHERE



Apache
MESOSTM



Marathon



DC/OS



- ▶ Connects the cluster
- ▶ Ressource management
 - ▶ Knows which servers are online
 - ▶ Knows which resources they have



Marathon

- ▶ Builds on top of Mesos
- ▶ Application Monitor
 - ▶ Can start new applications
 - ▶ "Monitors" running applications
 - ▶ If a machine/process dies Marathon restarts it
 - ▶ Can stop applications
 - ▶ Can scale up / down running applications within a cluster
- ▶ Offers available resources to new applications
- ▶ If Marathon dies it is restarted somewhere else



- ▶ Builds on top of Marathon
 - ▶ "Data-Center-Operating-System"
 - ▶ Offers one-click installers for applications (DC/OS Universe)
 - ▶ Includes Service Discovery
 - ▶ Includes Reverse Proxying
-
- ➡ You just run an Application through DC/OS and you can communicate with it without ever knowing any IP address

Frameworks

- ▶ Marathons little helpers
 - ▶ Know details about complete Applications
 - ▶ Communicate with Marathon to order sub-processes
 - ▶ Can describe additional conditions for processes
 - ▶ Like: "Do not run all DB instances on the same Machine"
 - ▶ Can trigger marathon to scale up/down processes
- ArangoDB has a Marathon Framework

Live Demo

ArangoDB on DC/OS

Is Mesosphere required?

- ▶ ArangoDB can run clusters without it
 - ▶ Setup Requires manual effort (can be scripted):
 - ▶ Configure IP addresses
 - ▶ Correct startup ordering
 - ▶ This works:
 - ▶ Automatic Failover (Follower takes over if leader dies)
 - ▶ Rebalancing of shards
 - ▶ Everything inside of ArangoDB
 - ▶ This is based on Mesos:
 - ▶ Complete self healing
 - ▶ Automatic restart of ArangoDBs (on new machines)
- We suggest you have someone on call

Does your data-model scale?

- ▶ Rule of thumb: "The more complex the query, the less it scales".
- ▶ Indexed-Attribute lookups:
 - ▶ Scale "infinitely"
 - ▶ Examples
 - ▶ Key/Value
 - ▶ FOR doc IN sharded FILTER doc.name == 'Michael'
- ▶ Joins:
 - ▶ Each join operation requires one additional round-trip (find left side, than find right side)
 - ▶ Scale "okayish"
- ▶ Graphs:
 - ▶ Each search depth potentially requires 2 join-like operations
 - ▶ A single-server needs to hold a large intermediate result
 - ▶ Does not scale too good

An overview of other ArangoDB features

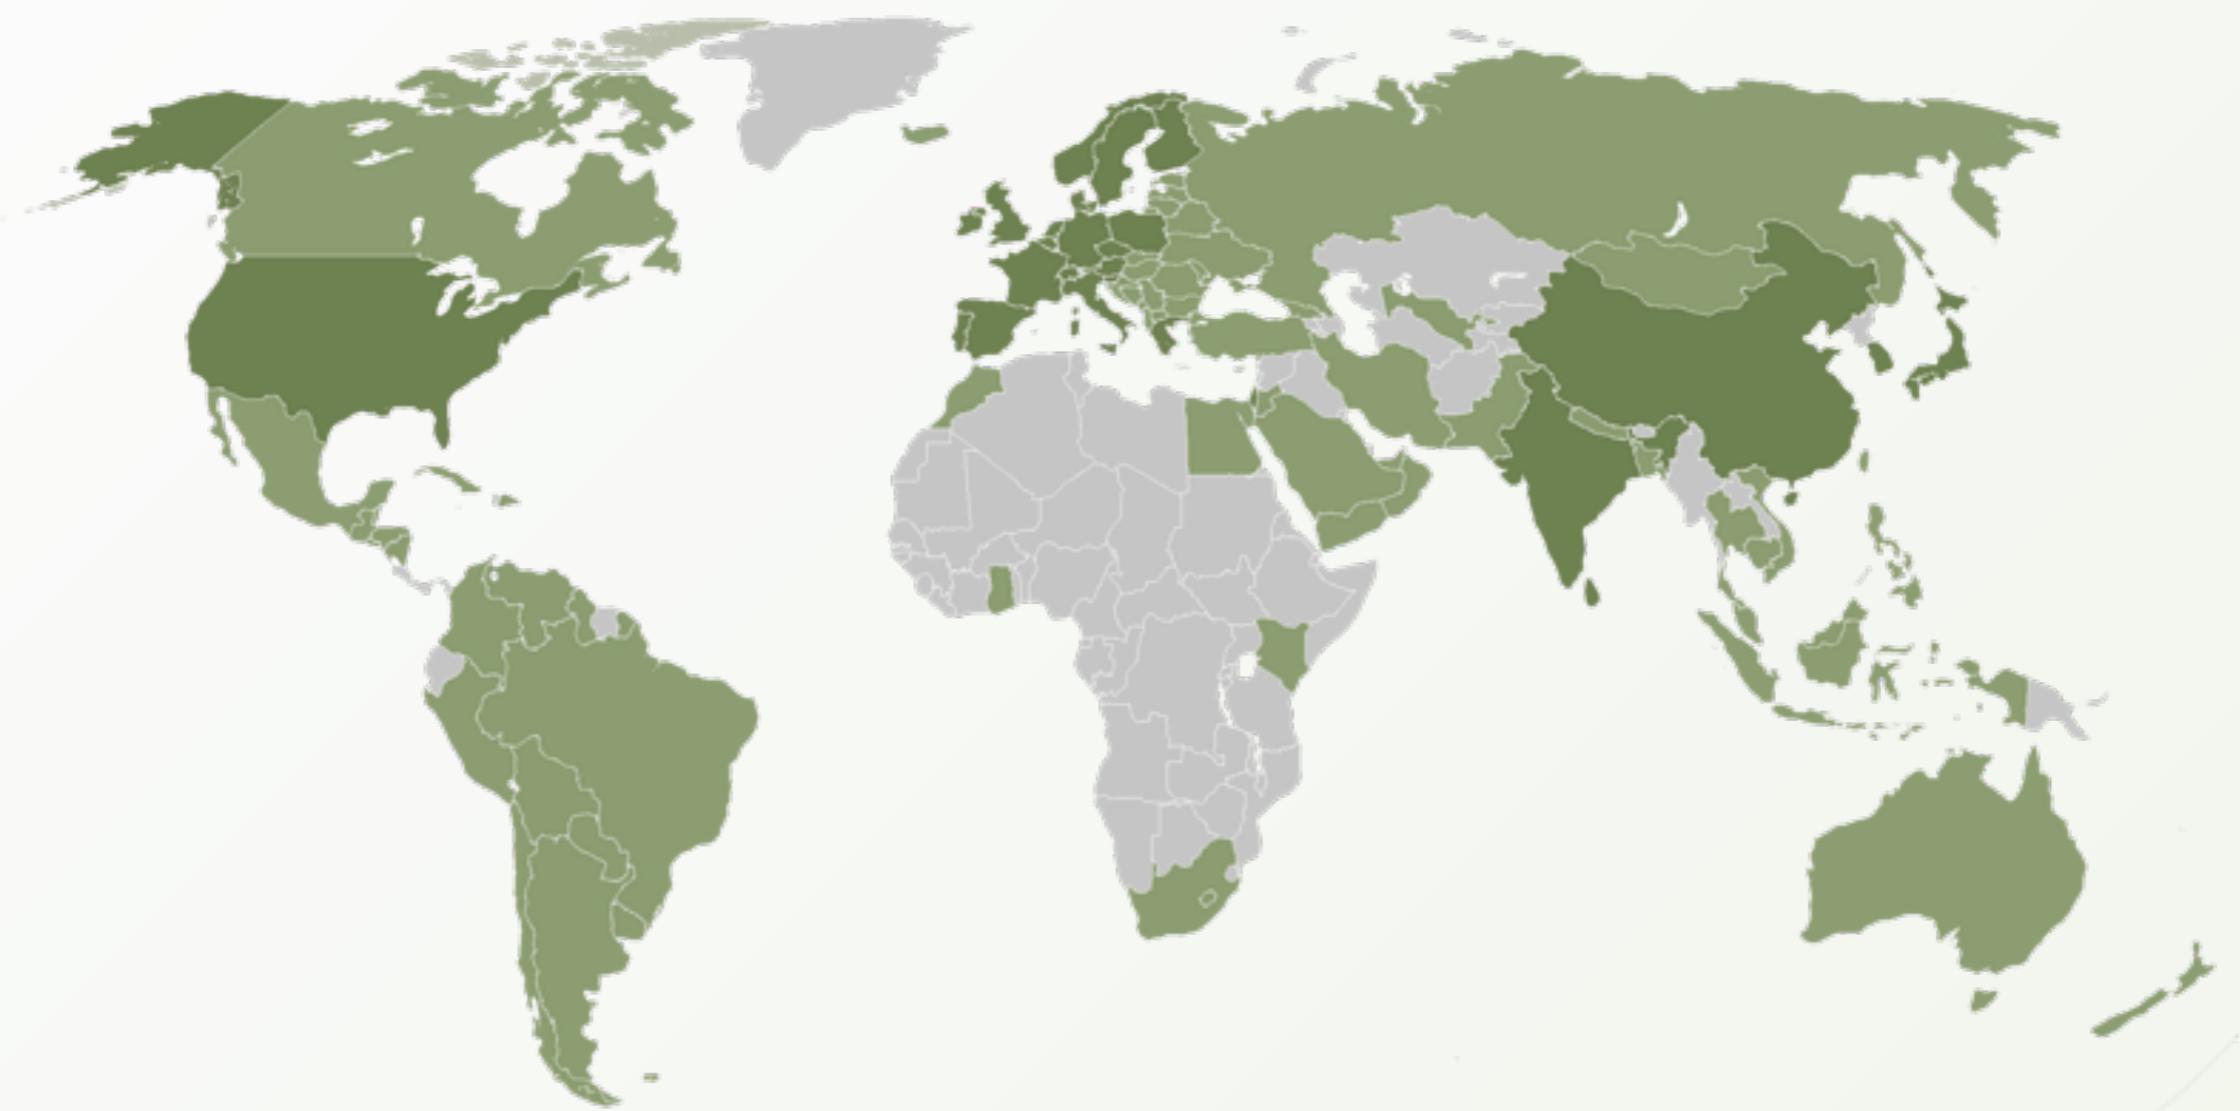
- ▶ open source and free (Apache 2 license)
- ▶ sharding & replication
- ▶ JavaScript throughout (V8 built into server)
- ▶ drivers for a wide range of languages
- ▶ web frontend
- ▶ good & complete documentation
- ▶ professional as well as community support

What about Java?

- ▶ Basic Driver developed in-house
 - ▶ Tutorial: <https://www.arangodb.com/tutorials/tutorial-java/>
- ▶ We are working on a Spring-Data integration
- ▶ We are very happy to get input on this
 - ▶ Feedback
 - ▶ Feature Requests
 - ▶ Personal Opinion on other Drivers/Integration you worked with
 - ▶ Development

Join our growing community

.. working on the geo index, the full text search and many APIs: Ruby, Python, PHP, Java, D, Javascript, ...



Thank you

► Further questions?

- Follow us on twitter: @arangodb
- Join our community slack: arangodb-community.slack.com
- Join our google group: <https://groups.google.com/forum/#!forum/arangodb>

- Follow me on twitter/github: @mchacki
- Write me a mail: michael@arangodb.com